

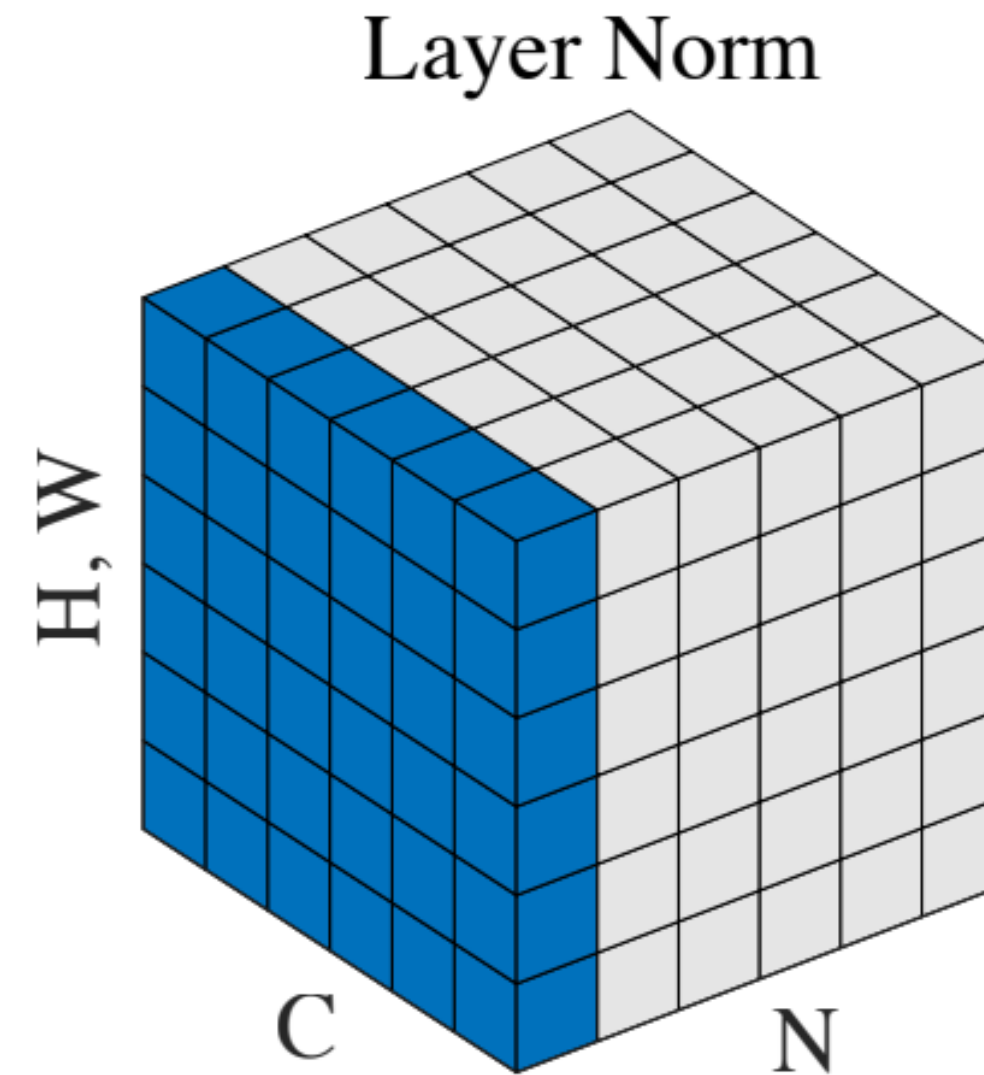
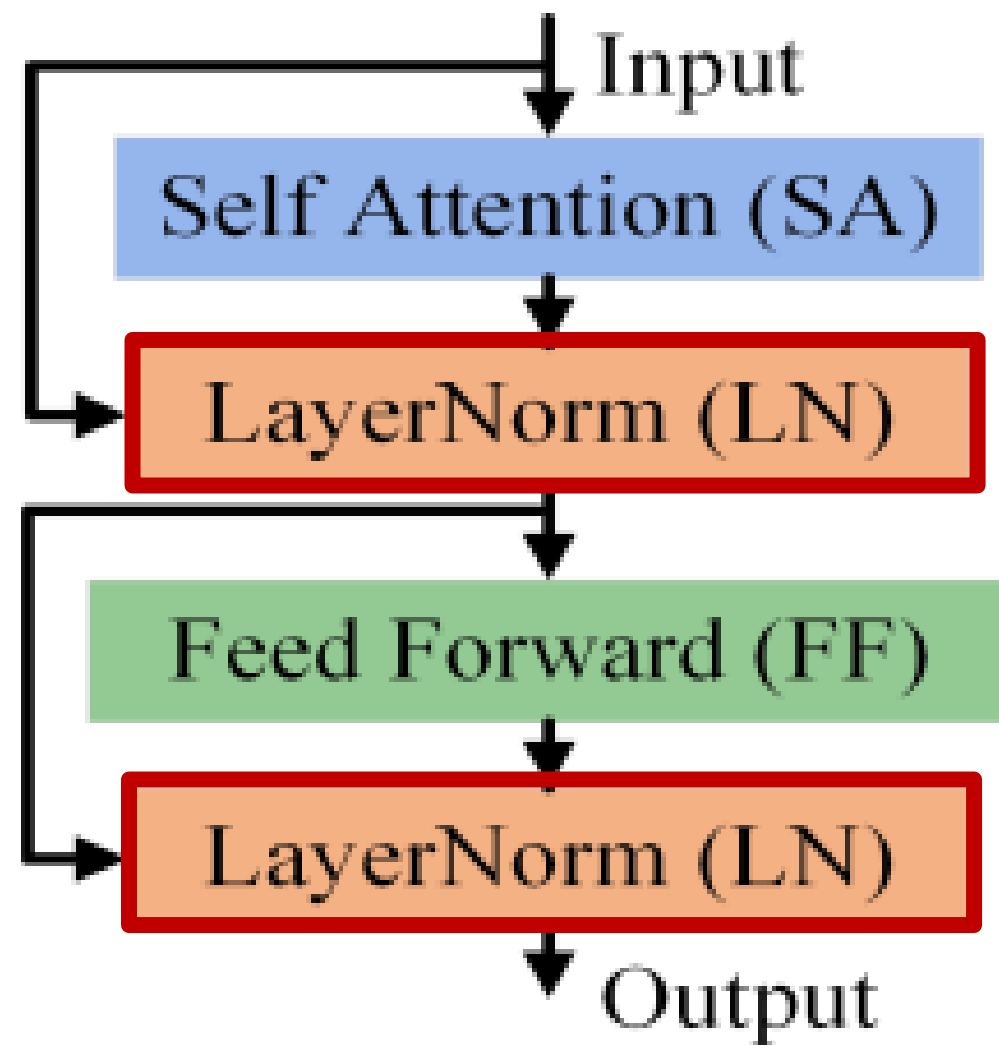
Accuracy-preserving Layer Normalization Approximations

CONTENTS

- 1. What is Layer Normalization?**
- 2. Why Optimize Layer Normalization?**
- 3. What This Paper Proposes?**
 - 1. Fixed-point quantization & Parallelism exploration**
 - 2. Standard Layer Normalization(PWL)**
 - 3. Pairwise Variance Algorithm**
- 4. Conclusion**

What is Layer Normalization?

LN location in Transformer

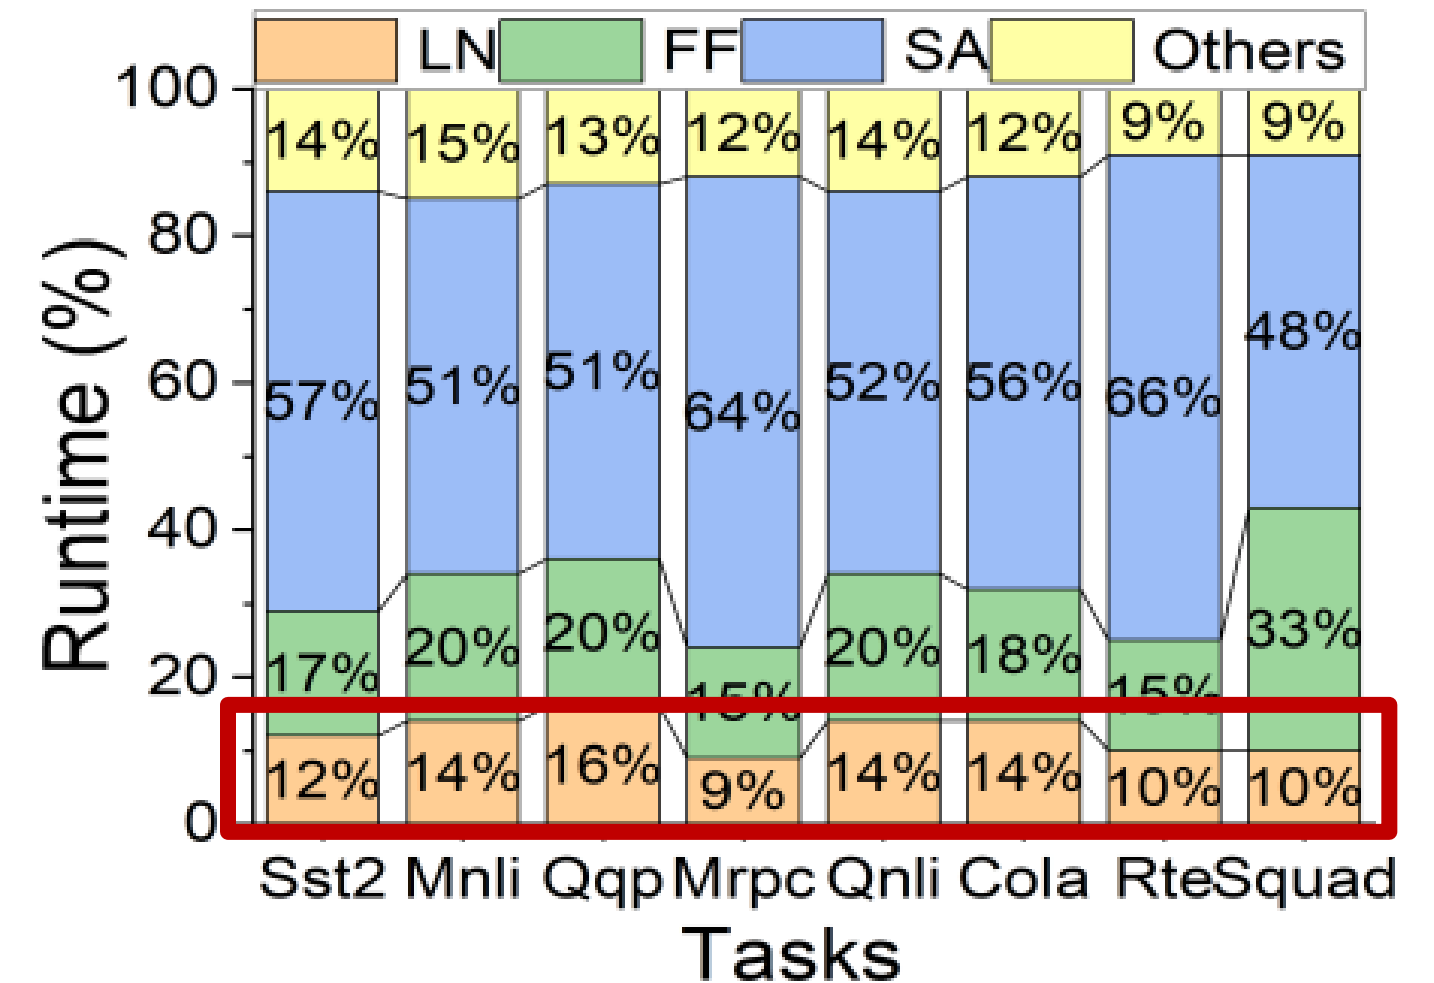


- Transformer have emerged as the choice learning model for NLP.
- LN contributes to the overall latency of the system

Why Optimize Layer Normalization?

$$LN(x_{i,j}) = \widehat{x}_{i,j} \cdot \gamma_j + \beta_j, \quad \text{for } j = 1, 2, \dots, d_{model}$$

- **First iteration** $\mu_i = \frac{1}{d_{model}} \cdot \sum_{j=1}^{d_{model}} x_{i,j}$
- **Second iteration** $\sigma_i^2 = \frac{1}{d_{model}} \cdot \sum_{j=1}^{d_{model}} (x_{i,j} - \mu_i)^2$
- **Third iteration** $\widehat{x}_{i,j} = \frac{x_{i,j} - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}, \quad \text{for } i = 1, 2, \dots, d_{token}$



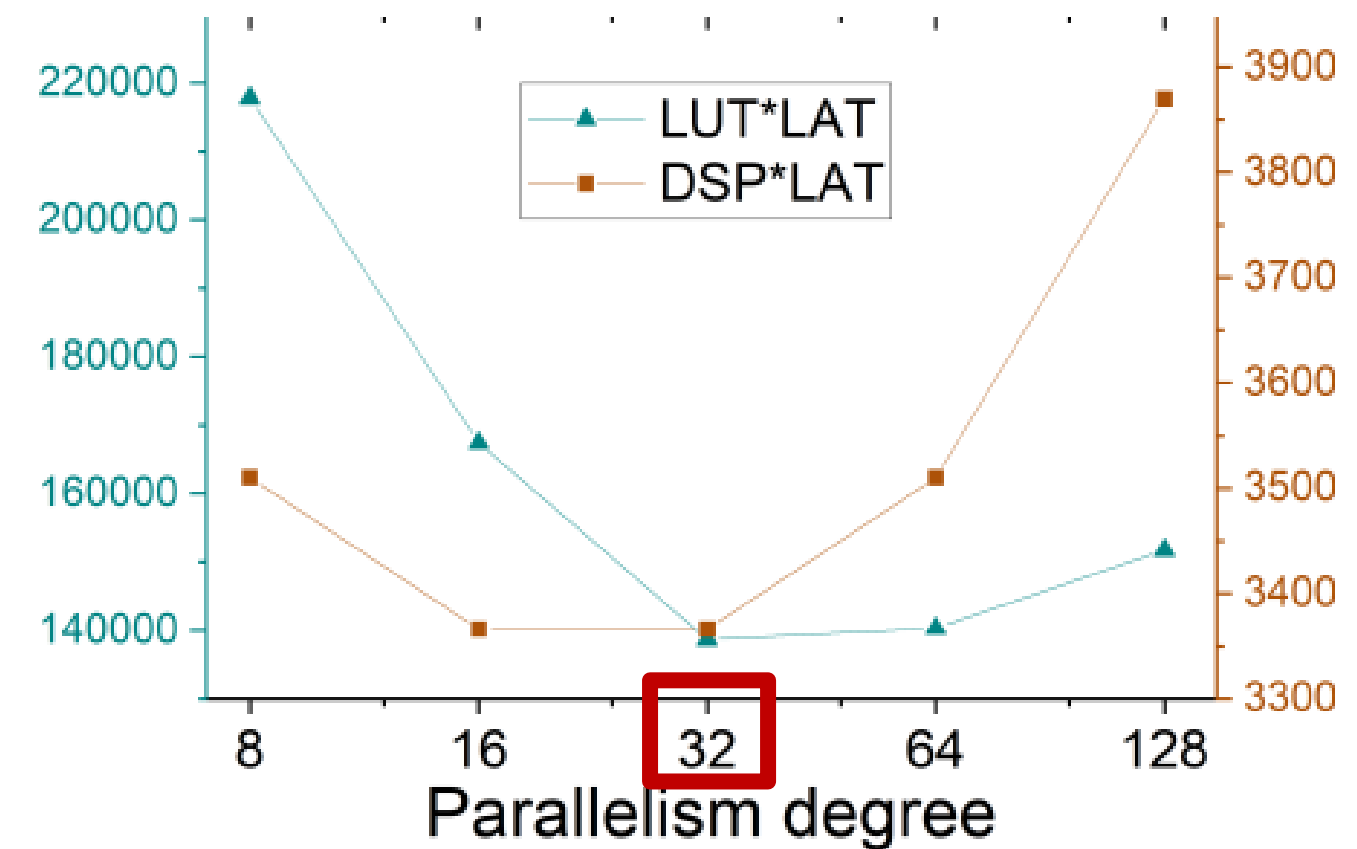
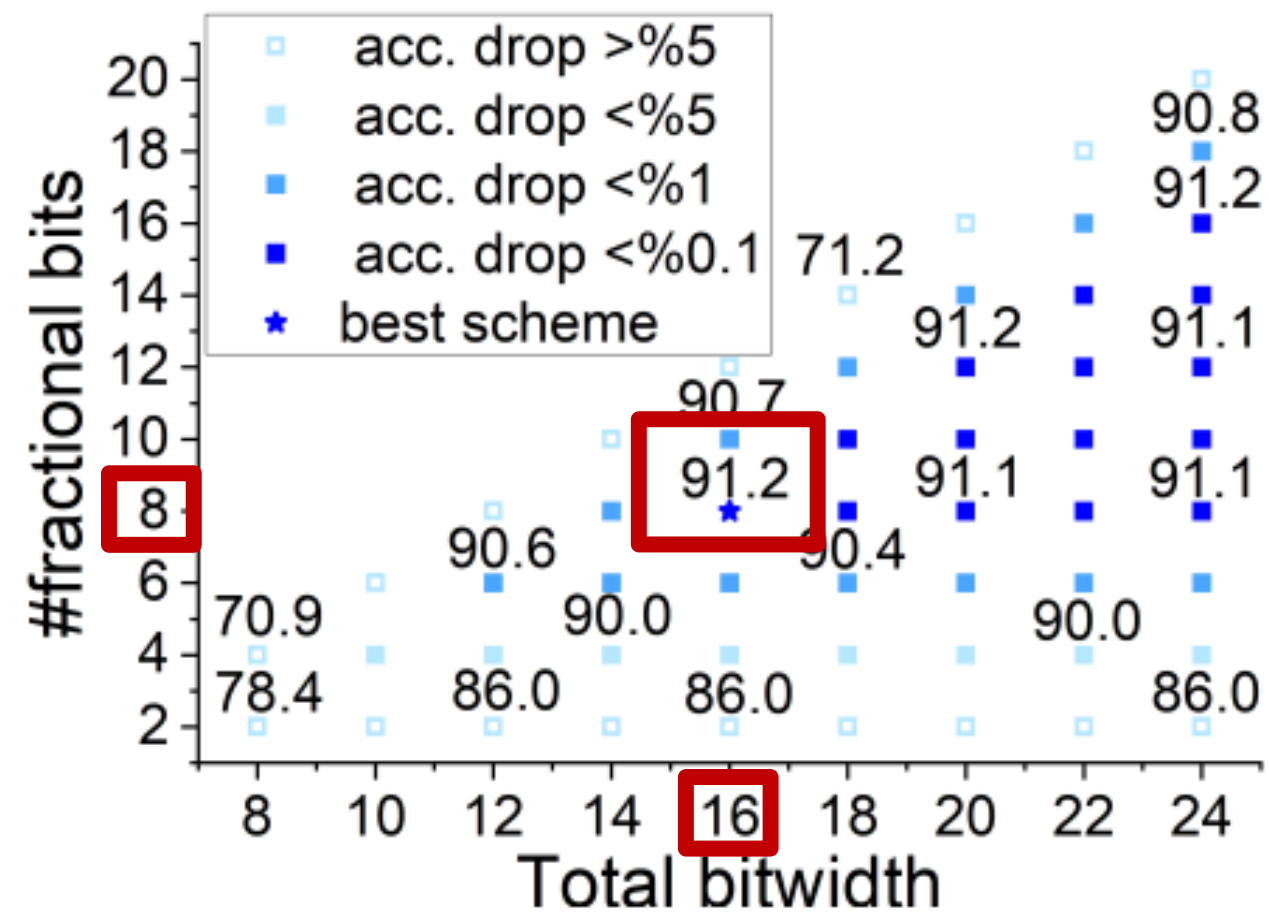
→ **hardware inefficient** due to the **square root** and **division operations**

→ **time consuming** since it needs **three iterations** on the input data

What This Paper Proposes?

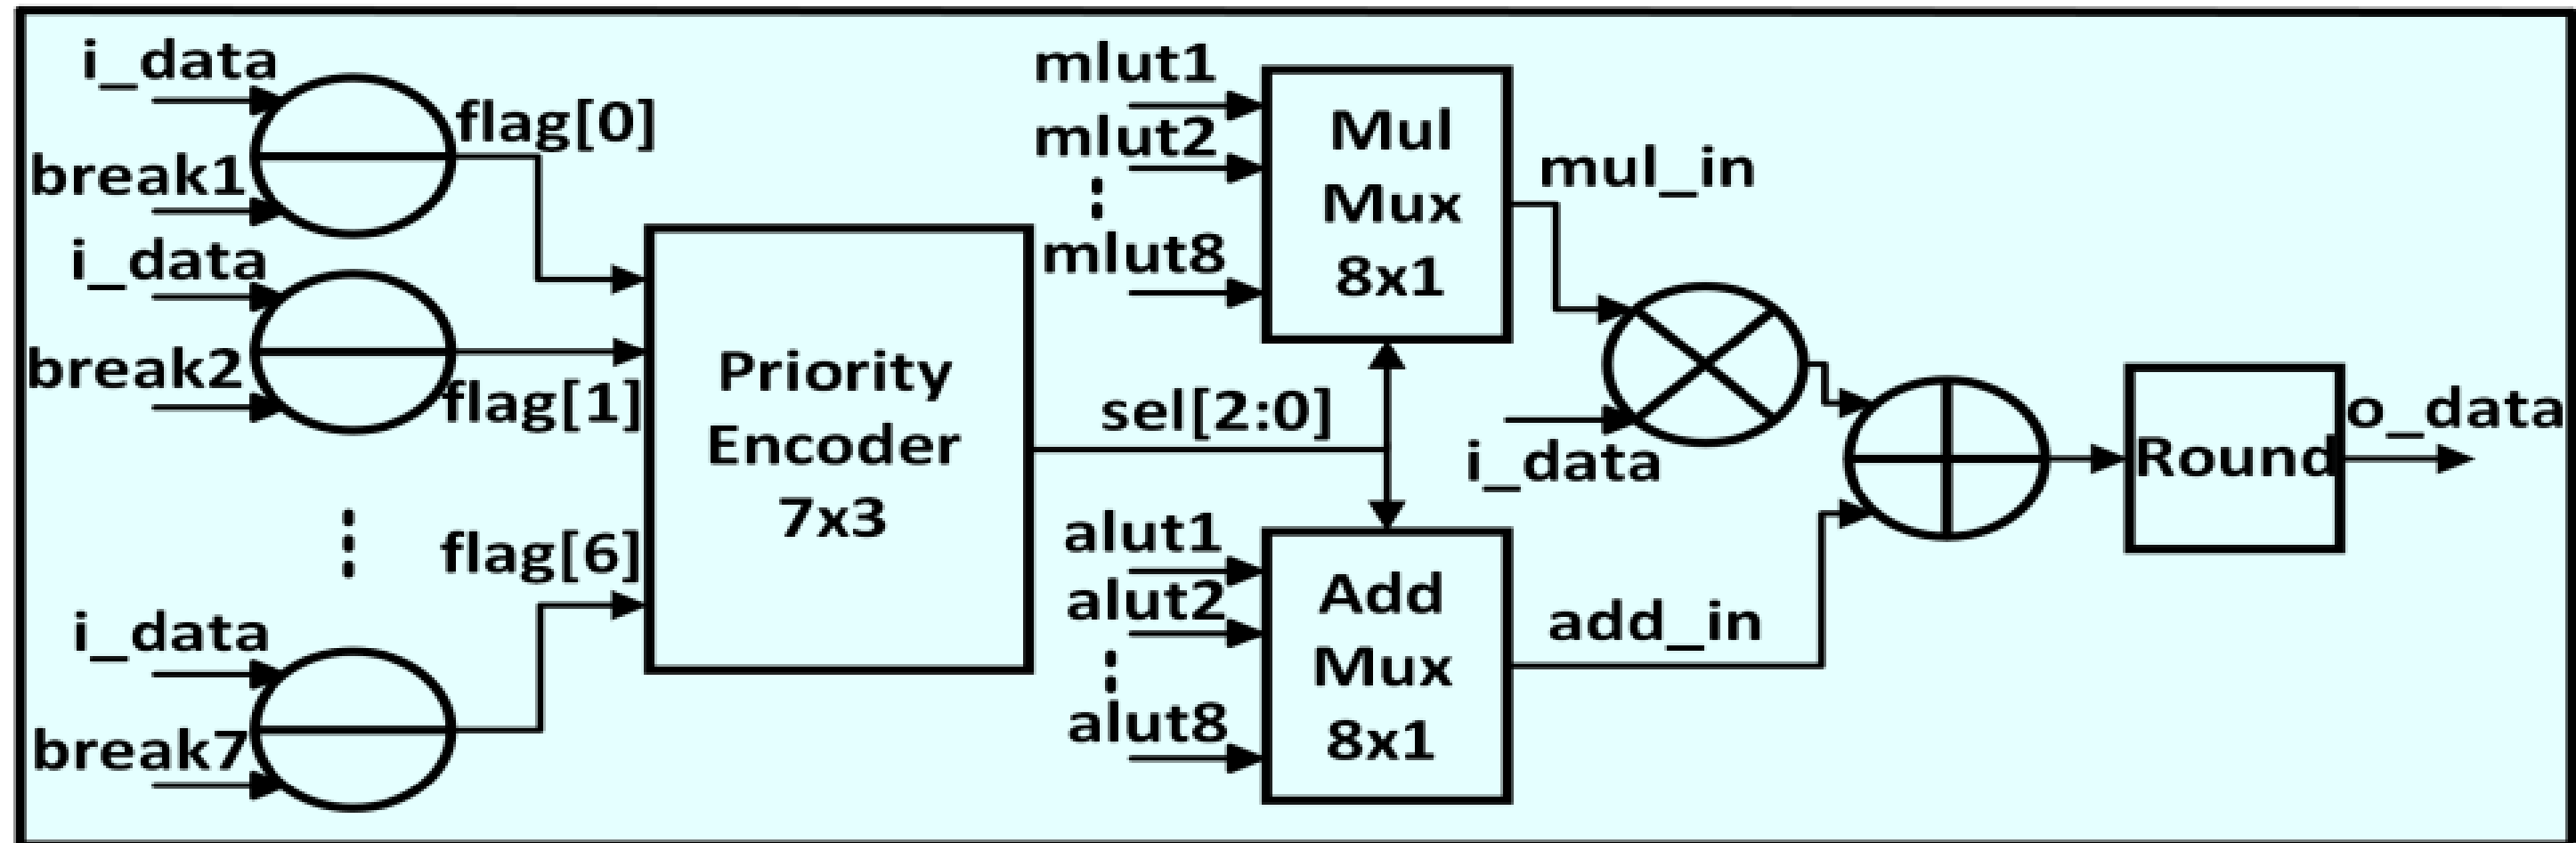
- Hardware-efficient LN approximation **without fine-tuning**
- **Key contributions:**
 - Fixed-point quantization & Parallelism exploration
 - Standard Layer Normalization(PWL)
 - Pairwise Variance Algorithm

Fixed-point quantization & Parallelism exploration



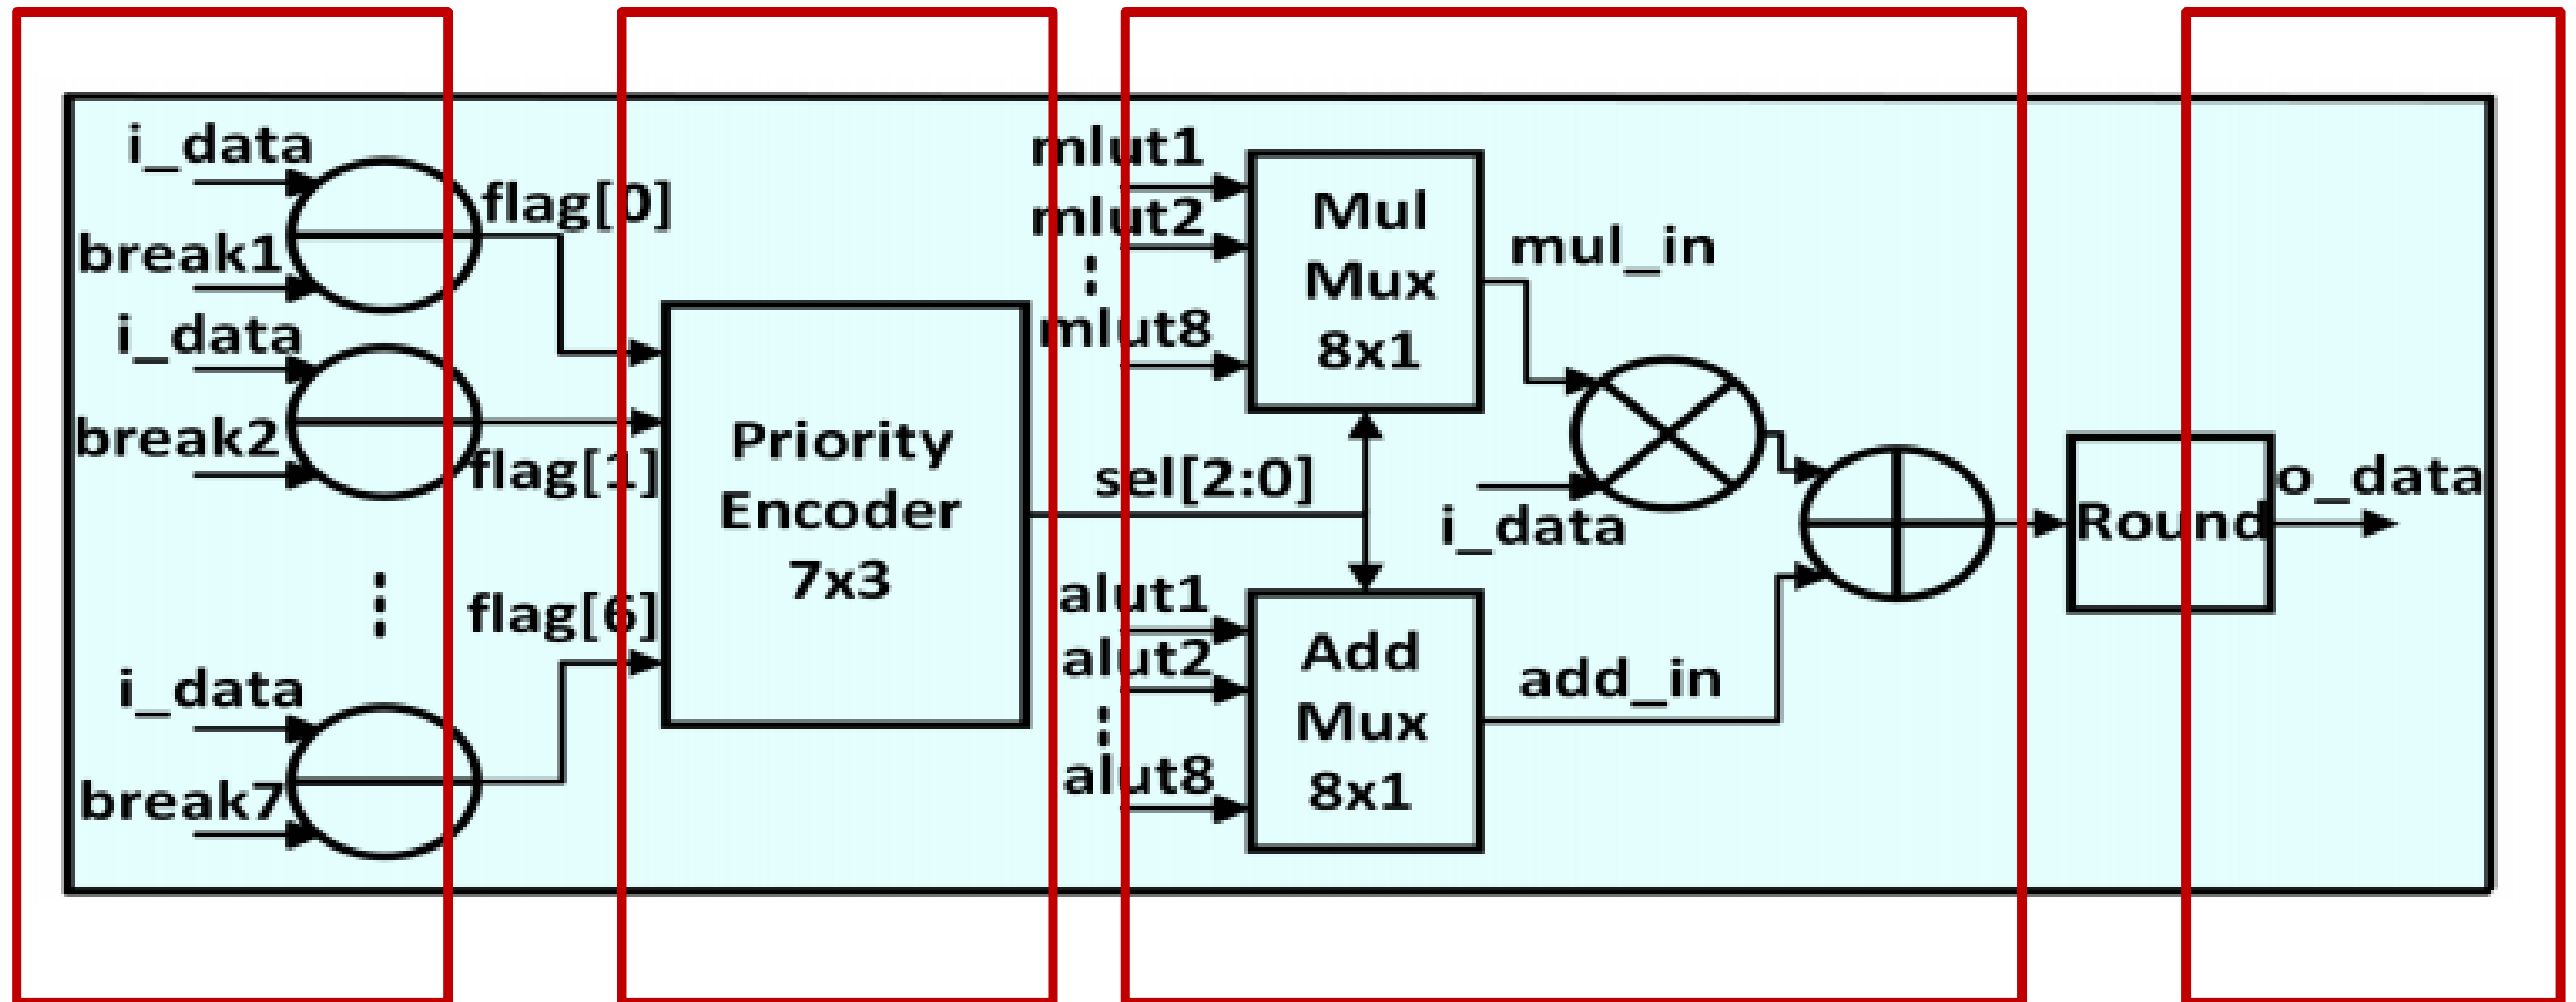
- Replaces 32-bit float with **16-bit fixed point** (8 integer + 8 fractional bits)
- Maintains accuracy ($\leq 0.28\%$ drop)
- The **best area-time trade-off** is found to be the **32-input** design

Standard Layer Normalization(PWL)



- Sqrt and reciprocal approximated using **7 non-uniform breakpoints (8 line segments)**
- Executes in **1 clock cycle**

Standard Layer Normalization(PWL)



Breakpoints
(→ 8 segments)

Input data x
→ Select one of
8 segments using
a 3-bit index [2:0]

PWL Approximation
using the function `pwl_approx()`

Output $y = a_3x + b_3$

Standard Layer Normalization(PWL)

- using pwlf python lib
- **pwlf library**: Automatically finds optimal breakpoints for piecewise linear fitting

```
# PWL fit
x_vals = np.linspace(0.01, 128, 1000)
sqrt_vals = np.sqrt(x_vals)
recip_vals = 1 / sqrt_vals

sqrt_model = pwlf.PiecewiseLinFit(x_vals, sqrt_vals)
sqrt_breaks = sqrt_model.fit(8)
sqrt_slopes = sqrt_model.slopes a
sqrt_intercepts = sqrt_model.intercepts b

recip_model = pwlf.PiecewiseLinFit(x_vals, recip_vals)
recip_breaks = recip_model.fit(8)
recip_slopes = recip_model.slopes
recip_intercepts = recip_model.intercepts
```

Input : 0.01 ~ 128

Precomputing sqrt and reciprocal

Fits a piecewise linear function with 8 segments

Standard Layer Normalization(PWL)

```
# PWL approximation
def pwl_approx(x, breakpoints, slopes, intercepts):
    x = np.clip(x, breakpoints[0], breakpoints[-1])
    out = np.zeros_like(x)
    for i in range(len(slopes)):
        mask = (x >= breakpoints[i]) & (x < breakpoints[i + 1])
        out[mask] = slopes[i] * x[mask] + intercepts[i]
    out[x >= breakpoints[-1]] = slopes[-1] * x[x >= breakpoints[-1]] + intercepts[-1]
    return out
```

If input the data x in breackpoint $[x_3, x_4)$

Then output the result $y = a_3x + b_3$

→ Output can be calculated in a **linear approximation**

→ Reduced latency

Standard Layer Normalization(PWL)

Result Evaluation

Latency

```
===== Timing (ms) =====  
[Sqrt Exact]    0.0268 ms  
[Sqrt PWL]      0.2562 ms  
[Recip Exact]   0.0106 ms  
[Recip PWL]     0.1365 ms
```

PWL > True

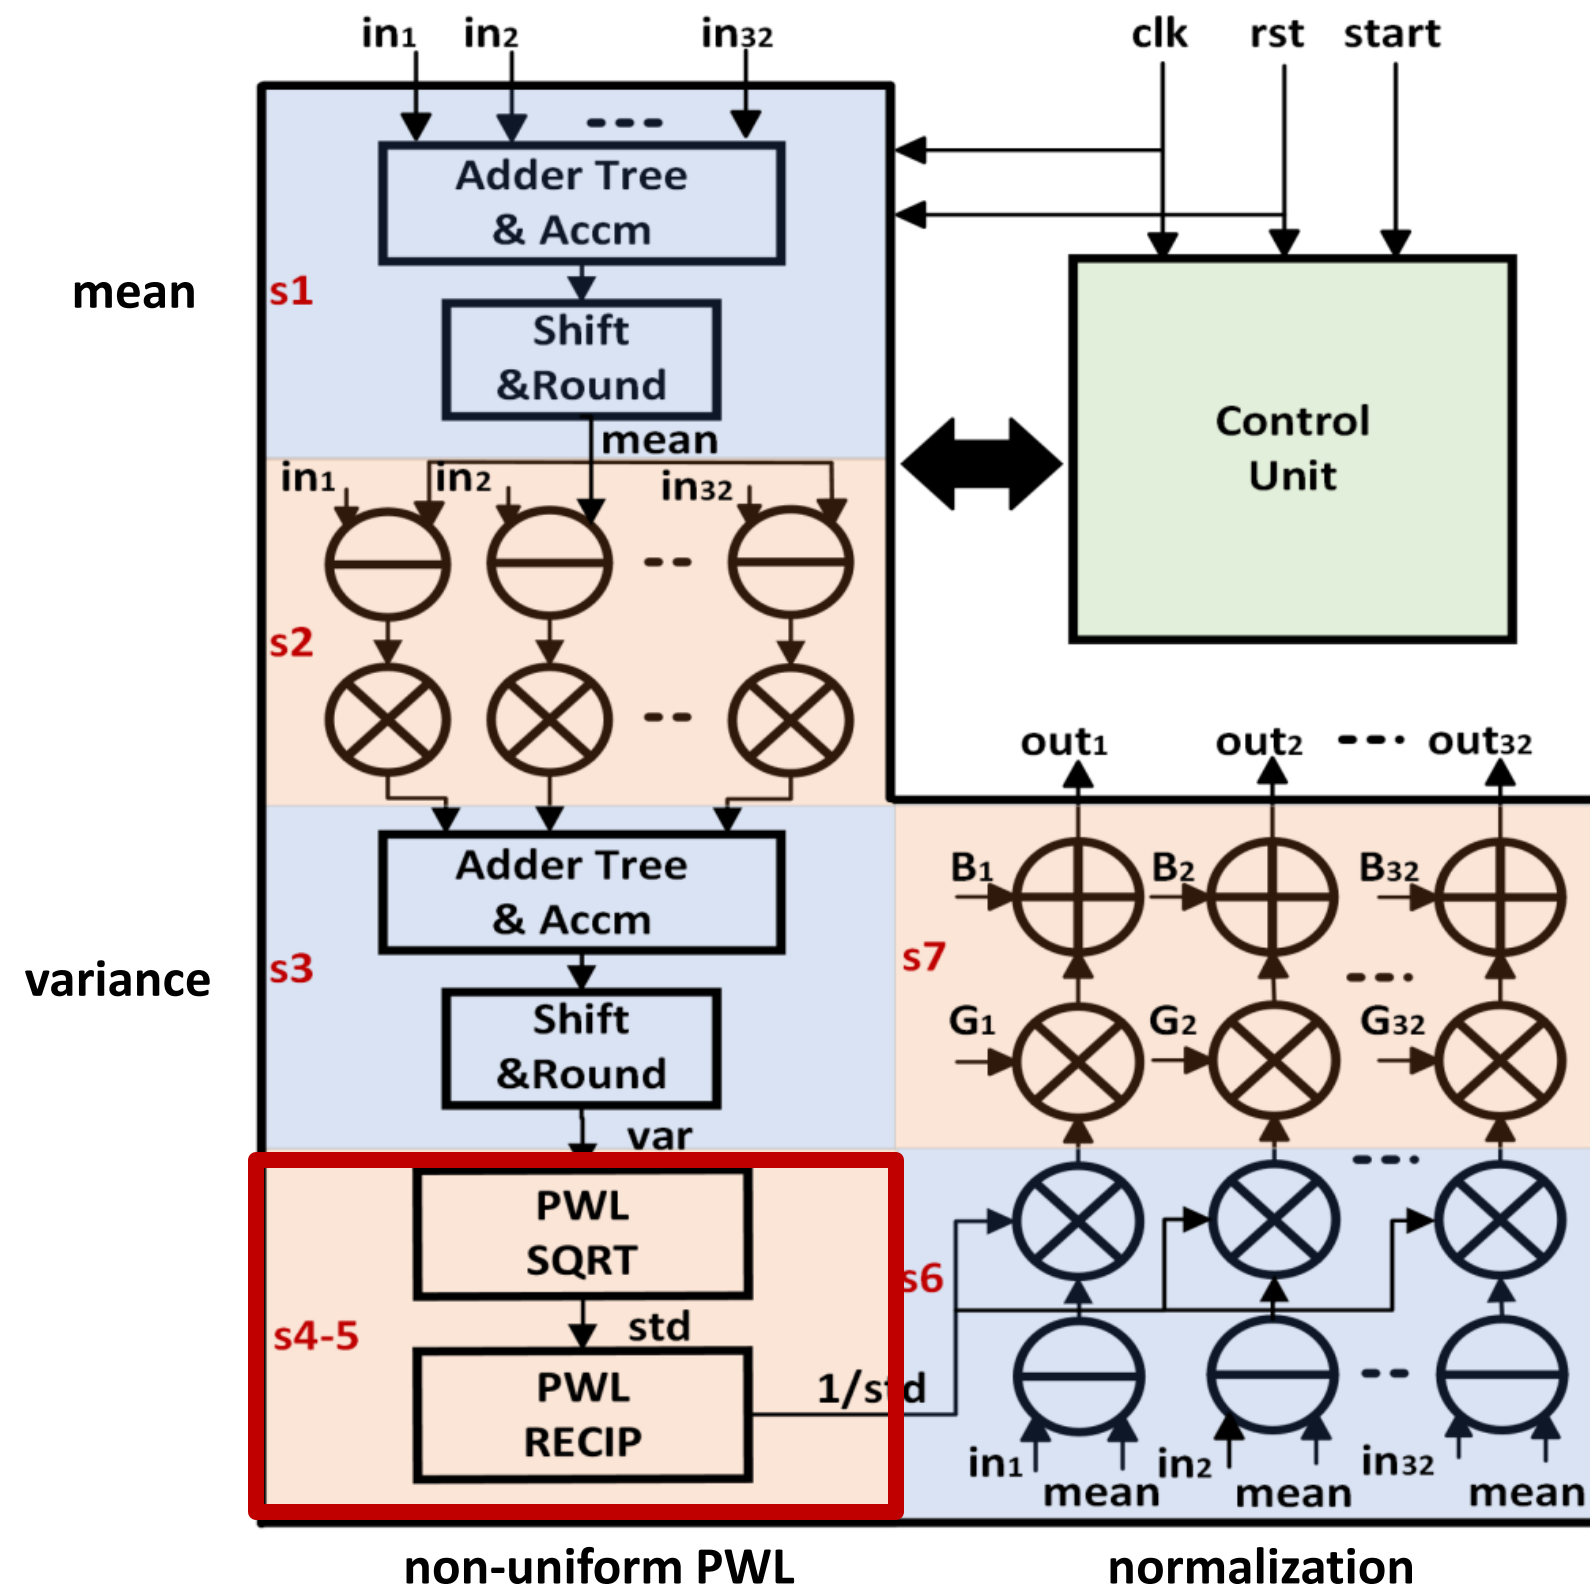
Latency increased due to **software**
(not FPGA) implementation.

Accuracy

```
===== Accuracy (% Error) =====  
[Sqrt PWL]      Mean Accuracy: 99.1760%  
[Recip PWL]     Mean Accuracy: 97.9223%
```

sqrt and reciprocal PWL approximations
achieved over **95% accuracy**.

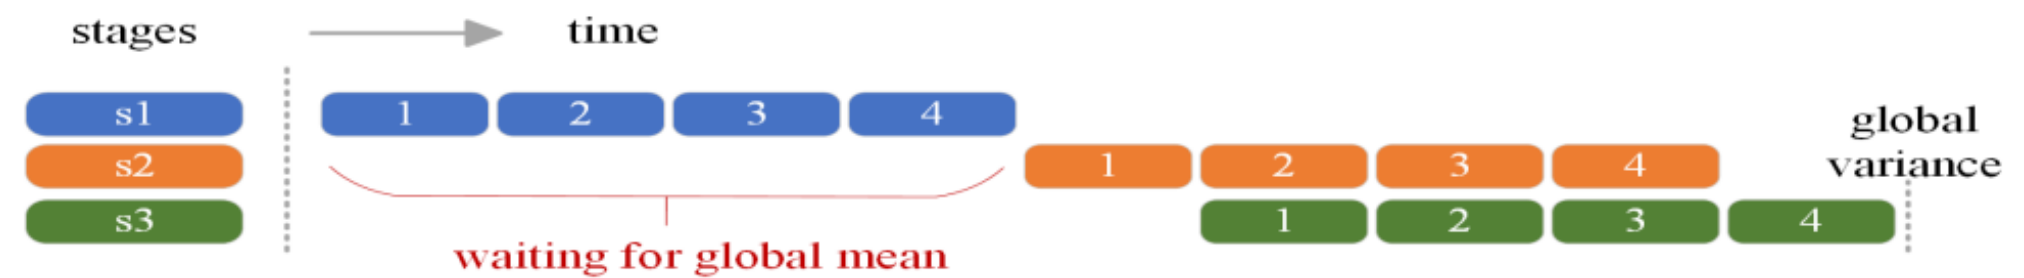
Standard Layer Normalization(PWL)



Total number of input: 512

Each batch: 32 elements

Each element: 16-bit fixed point



$s2$ must wait for the result of $s1$ until all inputs are presented

Pairwise Variance Algorithm

Algorithm:

the input data can be split into groups

$$\rightarrow \delta_{1(2)} = \mu_{1(3)} - \mu_{2(4)} \quad \delta = \frac{\mu_1 + \mu_2}{2} - \frac{\mu_3 + \mu_4}{2}$$

the variance of each group can be calculated separately

$$\rightarrow \text{intVar}_{1(2)} = \sigma_{1(3)}^2 + \sigma_{2(4)}^2 + \delta_{1(2)}^2 \times \frac{n_{1(3)} \times n_{2(4)}}{n_{1(3)} + n_{2(4)}}$$

intermediate values are used to calculate the global variance $\rightarrow \sigma^2 = \text{intVar}_1 + \text{intVar}_2 + \delta^2 \times \frac{(n_1 + n_2) \times (n_3 + n_4)}{n_1 + n_2 + n_3 + n_4}$

Pairwise Variance Algorithm

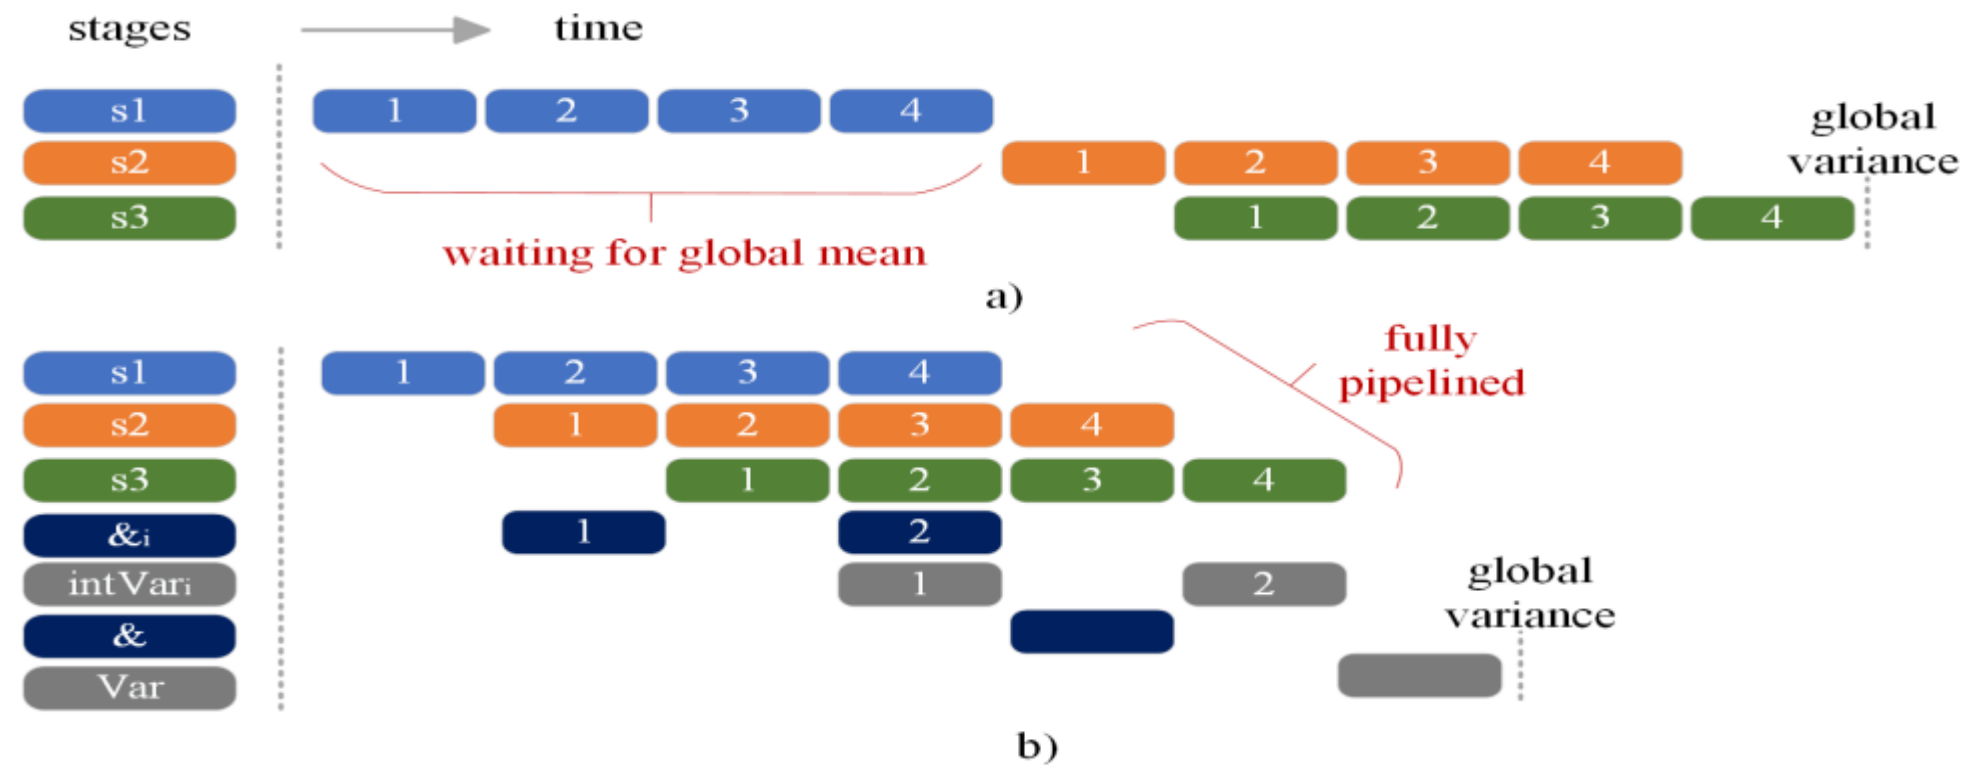


Fig. 5. Exemplified timing diagrams for a) standard variance b) pairwise variance calculation.

- Avoids waiting for global mean → reduces iteration count from 3 to 2
 - Enables **fully pipelined architecture** for s1–s3 stages
- Only one more clock cycle to calculate the global variance
- * No change to the remaining hardware components of s4-7

Pairwise Variance Algorithm

Standard variance calculation

```
# Variance methods
def true_variance(x):
    mean = x.mean(axis=-1, keepdims=True)
    return ((x - mean) ** 2).mean(axis=-1, keepdims=True) * (1)
```

: Two sequential divisions

→ major reason of latency

$$* (1) \sigma_i^2 = \frac{1}{d_{model}} \sum_{j=1}^{d_{model}} (x_{i,j} - \mu_i)^2$$

One-pass variance algorithm

```
def one_pass_variance(x):
    mean = x.mean(axis=-1, keepdims=True)
    mean_sq = (x**2).mean(axis=-1, keepdims=True)
    return mean_sq - mean**2 * (2)
```

: Reduces the computation to a single pass

→ lowers the latency

→ Increase the risk of floating-point errors

$$* (2) \sigma_i^2 = \frac{1}{d_{model}} \sum_{j=1}^{d_{model}} x_{i,j}^2 - \mu_i^2$$

Pairwise Variance Algorithm:

- Input is split into 16 groups
 - Variance is calculated in a hierarchical fashion: $16 \rightarrow 8 \rightarrow 4 \rightarrow 2 \rightarrow 1$
- Fully pipelined architecture leads to reduced overall latency

Pairwise Variance Algorithm

Pairwise Variance Algorithm

```
def pairwise_variance(x):  
    # Number of groups to split into (must be a power of two)  
    G = 16  
    D = x.shape[-1]  
    # Ensure the last dimension is divisible by G  
    assert D % G == 0, f"Last dim {D} must be divisible by {G}"  
  
    # Split the input tensor into G equal parts along the last axis  
    splits = np.split(x, G, axis=-1)  
  
    # Compute per-group counts, means, and sum of squared deviations:  
    #  $M_i = \sum_j (x_{ij} - \mu_i)^2 = n_i * \text{var}_i$   
    n_list = [s.shape[-1] for s in splits] # number of elements in each group  
    mu_list = [s.mean(axis=-1, keepdims=True) for s in splits] # group means ( $\mu_i$ )  
    M_list = [  
        np.var(s, axis=-1, keepdims=True, ddof=0)  
        * n # group sum of squared deviations ( $M_i$ )  
        for s, n in zip(splits, n_list)  
    ]
```

Number of groups: 16

N_list: Number of elements in each group

Mu_list: Group means

M_list: Variance * n

Pairwise Variance Algorithm

Pairwise Variance Algorithm

```
# Iteratively merge groups in pairs: 16 -> 8 -> 4 -> 2 -> 1
while len(mu_list) > 1:
    next_mu, next_M, next_n = [], [], []
    for i in range(0, len(mu_list), 2):
        # Grab two adjacent groups
        mu1, mu2 = mu_list[i], mu_list[i + 1]
        M1, M2 = M_list[i], M_list[i + 1]
        n1, n2 = n_list[i], n_list[i + 1]

        # Compute merged sum of squared deviations:
        # M_12 = M1 + M2 + (mu1 - mu2)^2 * (n1 * n2) / (n1 + n2)
        delta = mu1 - mu2 * (1)
        M12 = M1 + M2 + delta**2 * (n1 * n2) / (n1 + n2) * (2)
        next_M.append(M12)

        # Compute merged mean and count:
        # mu_12 = (n1 * mu1 + n2 * mu2) / (n1 + n2)
        next_mu.append((mu1 * n1 + mu2 * n2) / (n1 + n2))
        next_n.append(n1 + n2)

    # Prepare for next iteration
    mu_list, M_list, n_list = next_mu, next_M, next_n
```

- Start with 16 groups and iteratively merge them until only 1 group remains
- At each step, take adjacent values of μ , n , and M to compute δ_i and $intVar_i$
- Then, **update μ and n** and store them in the list for the next iteration.

* (1) $\delta_{1(2)} = \mu_{1(3)} - \mu_{2(4)}$

* (2) $intVar_{1(2)} = \sigma_{1(3)}^2 + \sigma_{2(4)}^2 + \delta_{1(2)}^2 \times \frac{n_{1(3)} \times n_{2(4)}}{n_{1(3)} + n_{2(4)}}$

Pairwise Variance Algorithm

Pairwise Variance Algorithm

```
# After merging, compute final biased variance:  
M_total = M_list[0] # total sum of squared deviations  
n_total = n_list[0] # total number of elements  
var_total = M_total / n_total # biased variance = M_total / n_total  
return var_total
```

$M_list = \text{variance} * n$

→ **Total variance** = $\frac{M_{total}}{n_{total}}$

Pairwise Variance Algorithm

Result Evaluation

Latency

```
===== Timing (ms) =====  
[True Var]      1.1087 ms  
[One-Pass Var]  0.1928 ms  
[Pairwise Var]  2.2948 ms
```

One-Pass < True < Pairwise

Accuracy

```
===== Accuracy (% Error vs PyTorch) =====  
[True Var]      100.0000%  
[One-Pass Var]  100.0000%  
[Pairwise Var]  100.0000%
```

One-Pass = True = Pairwise

Conclusion

- Proposed two LN hardware accelerators:
 - Standard Layer Normalization(PWL)
 - reduced the latency almost 3x
 - Pairwise variance algorithm
 - reduced the latency almost 4x
 - additional **27%** over standard LN (slightly more hardware resources)
- No fine-tuning required
- No retraining required