

LayerNorm FPGA

(Standard & Approximation)

2025.7.22.

Sangyun Kim
Eunju Kim

CONTENTS

1. LN Approximation Model

1. Pairwise module

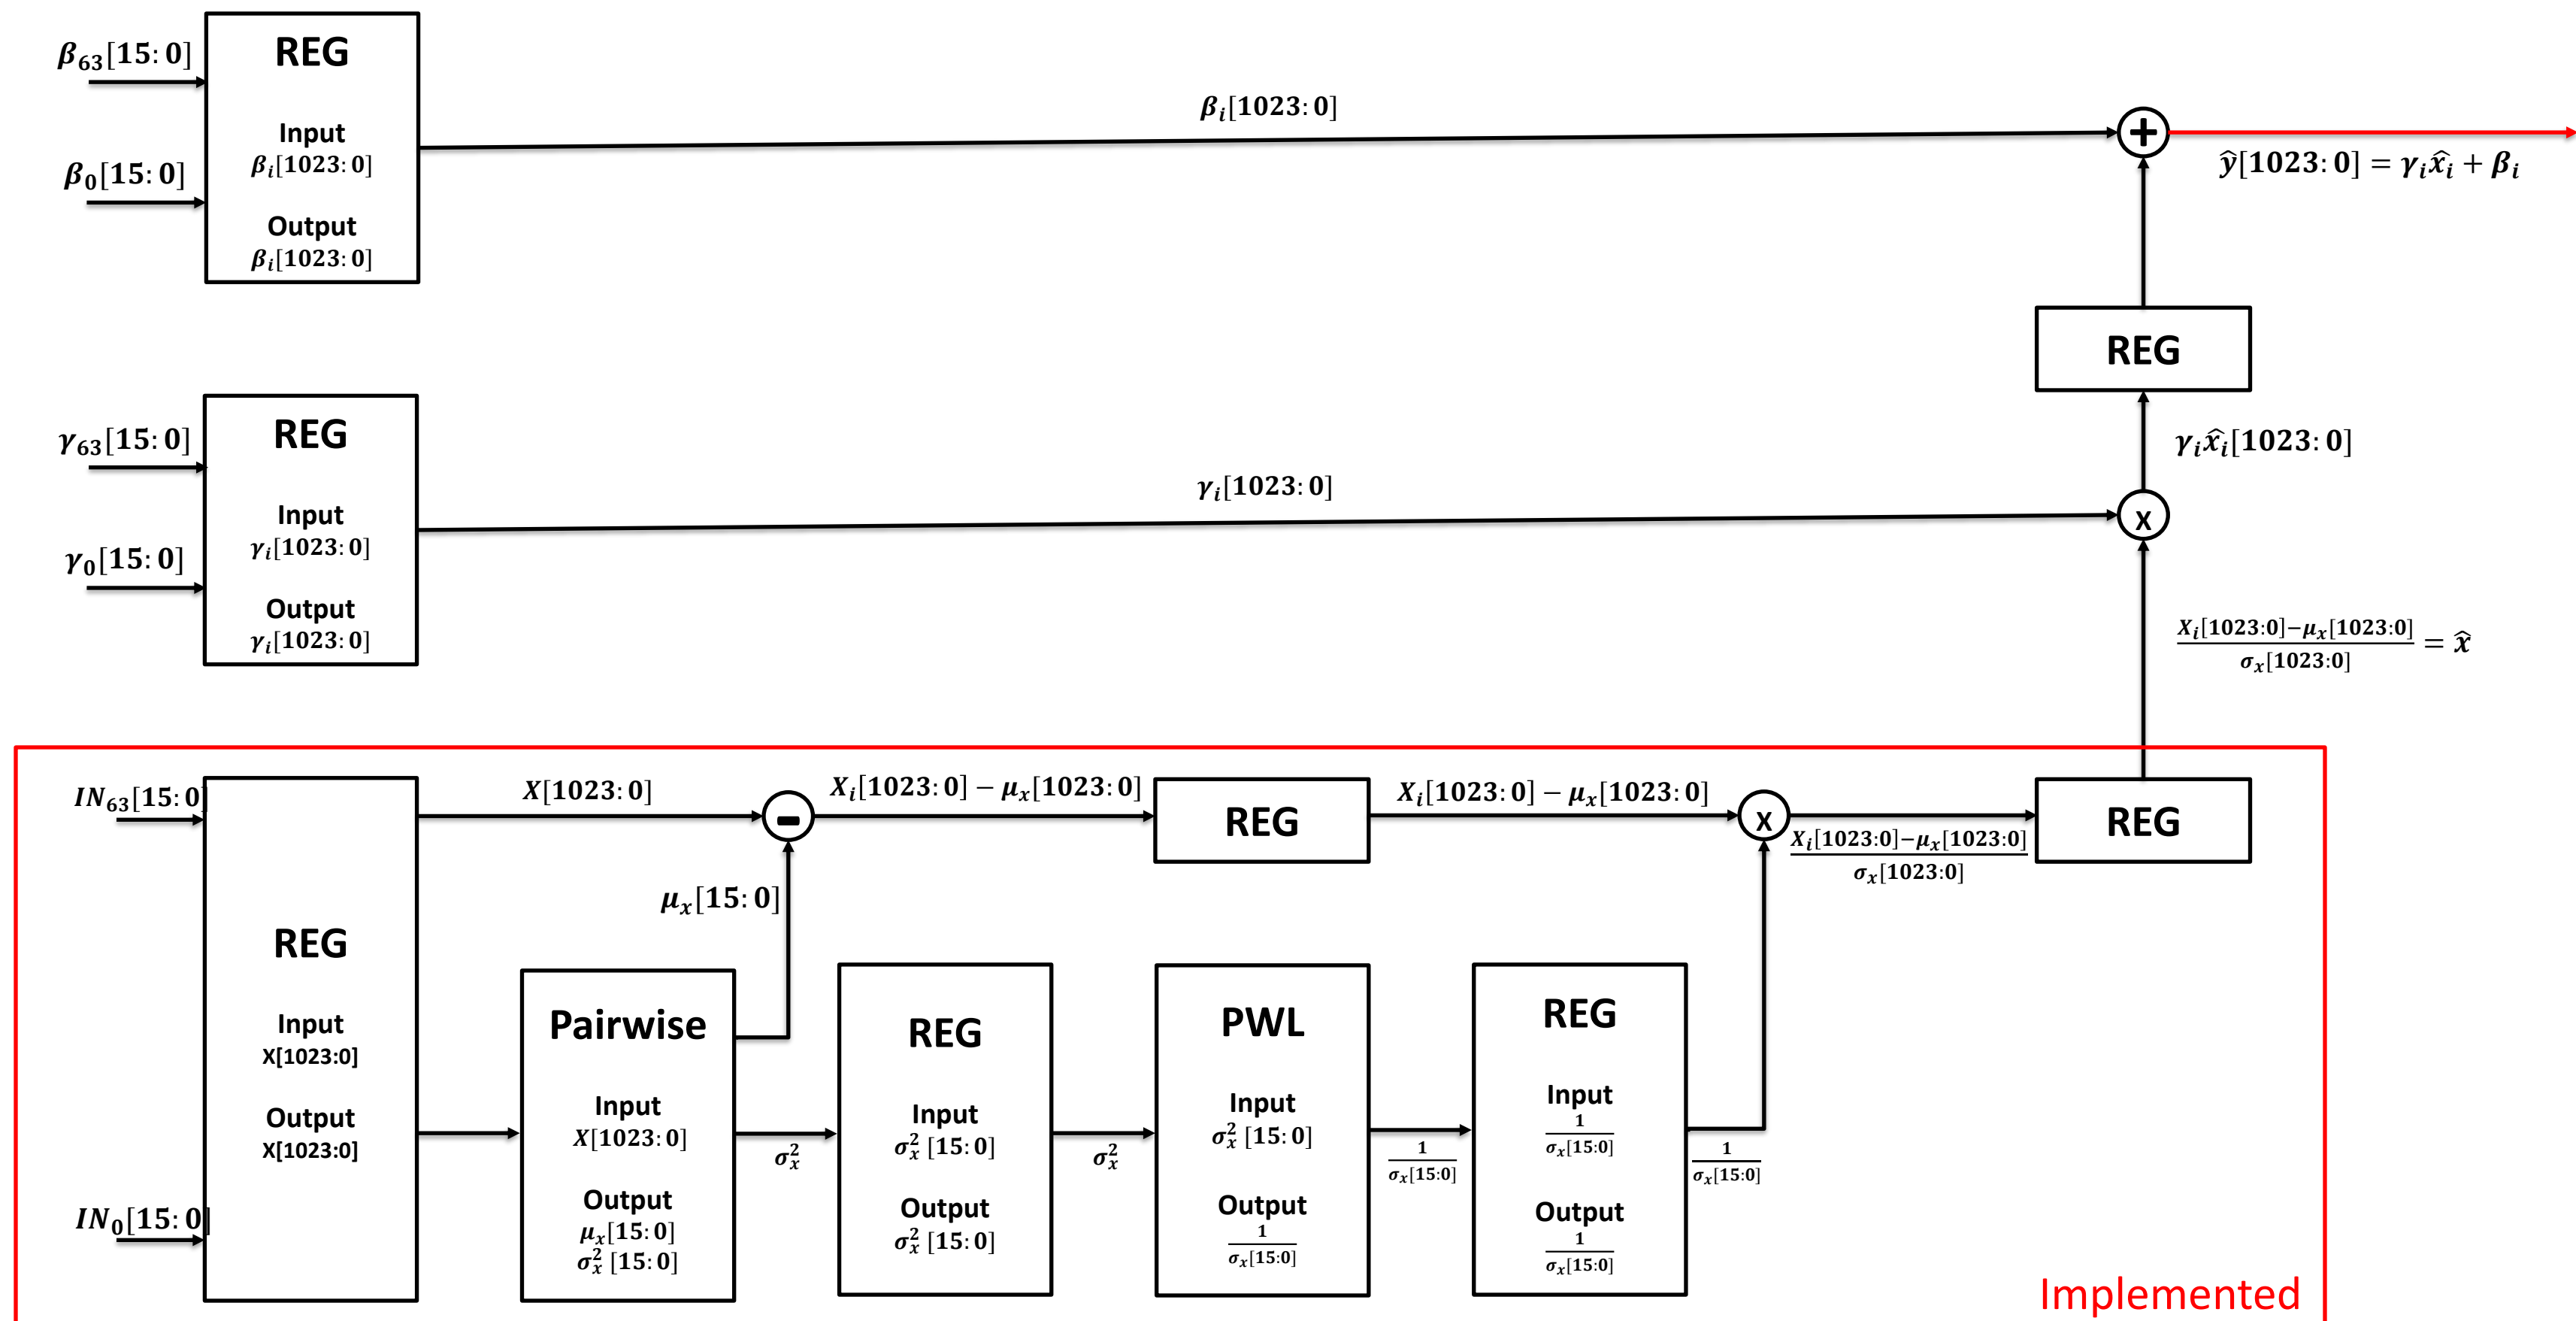
2. PWL module

3. Layer Normalization

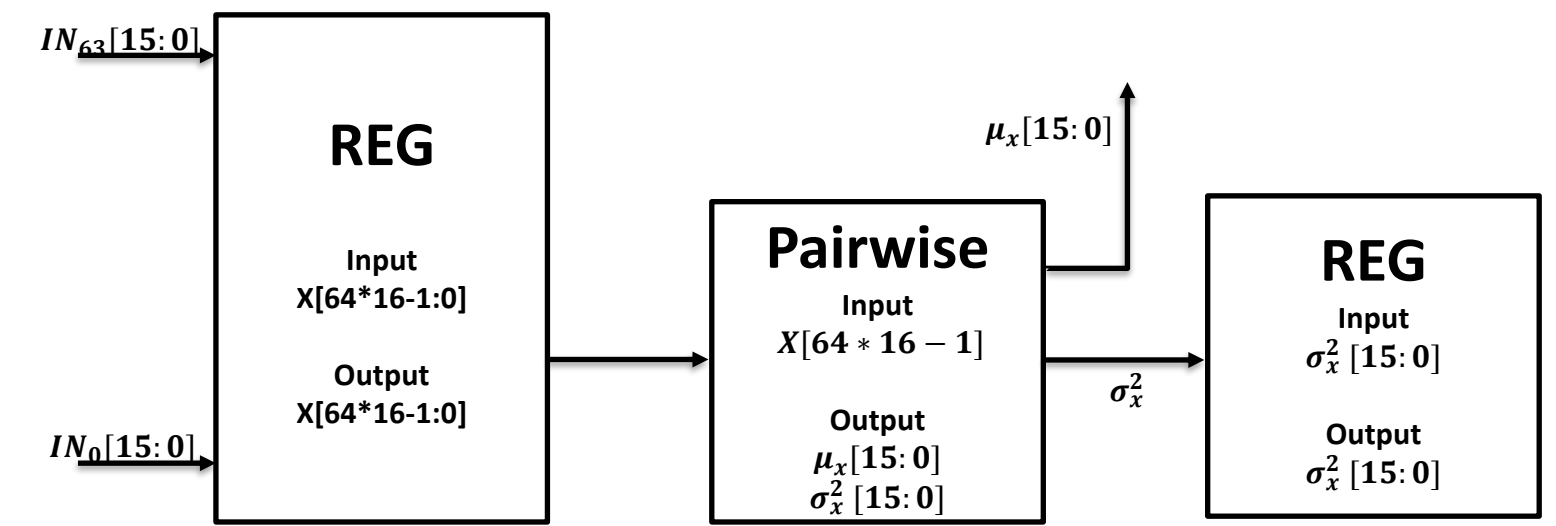
2. Apply FPGA

3. Plans for Next

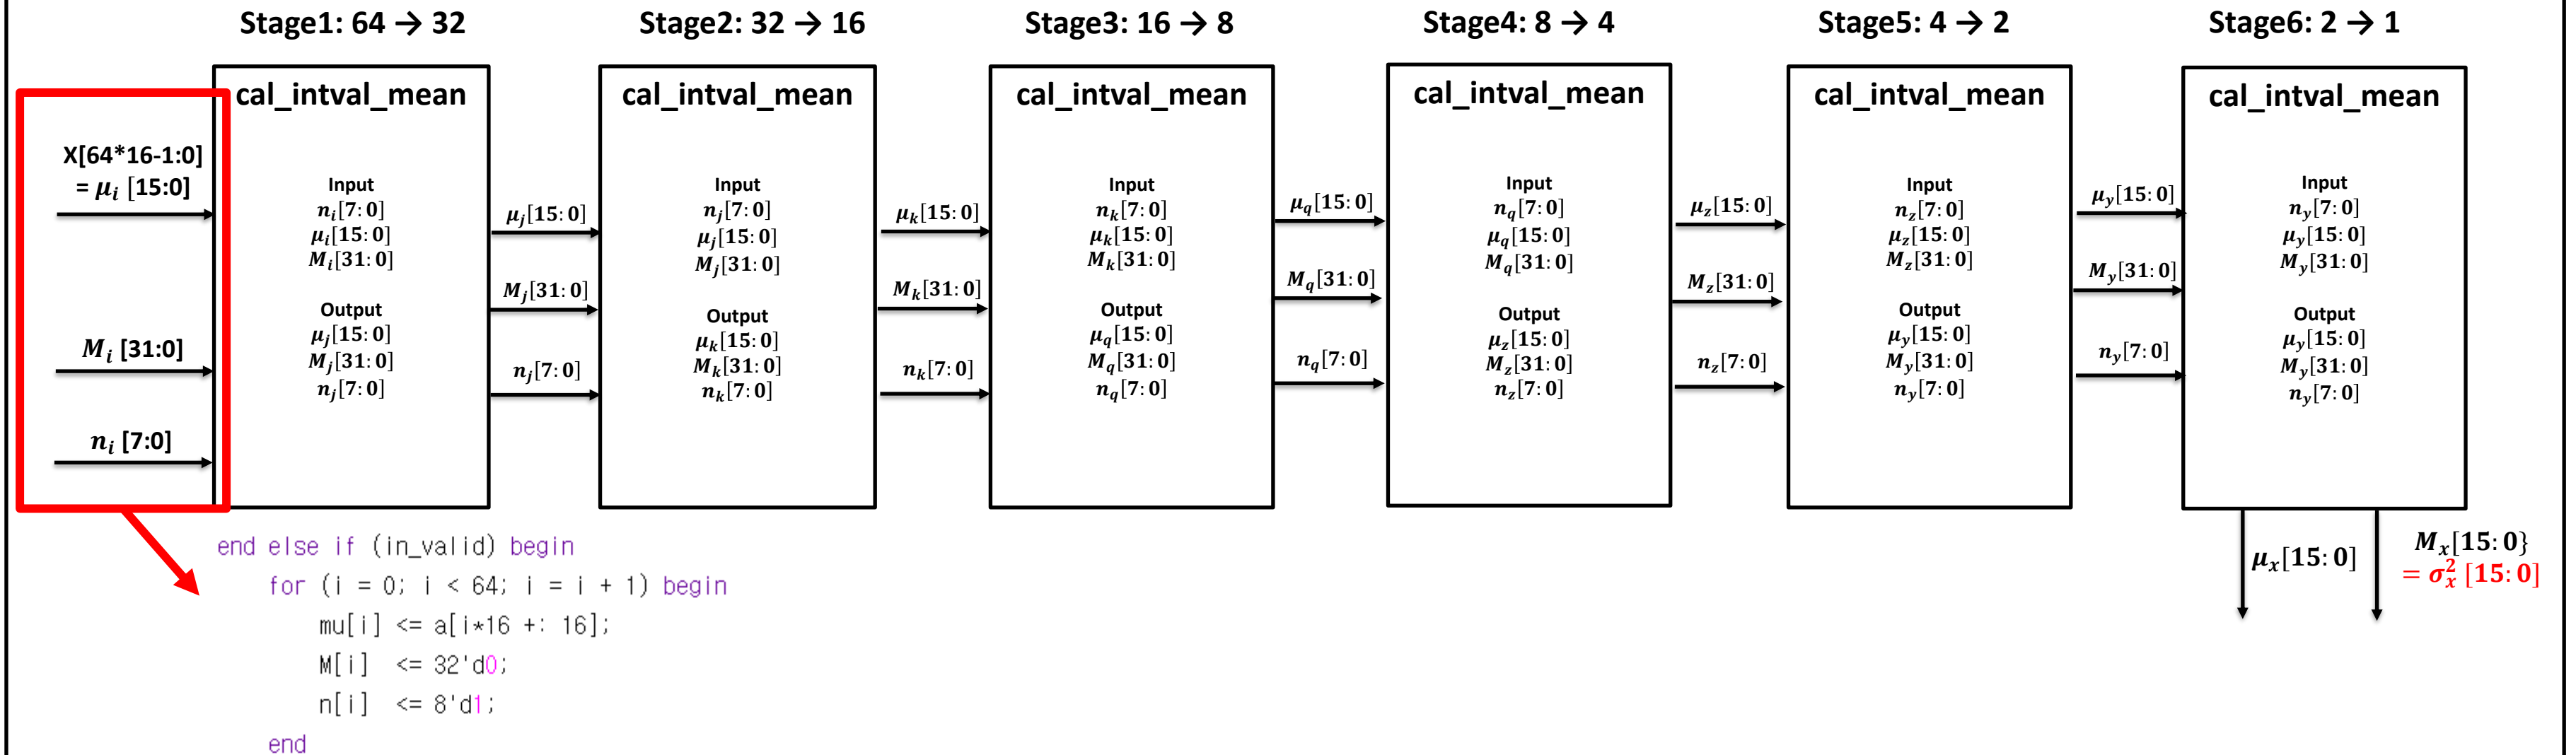
LN Approximation Model



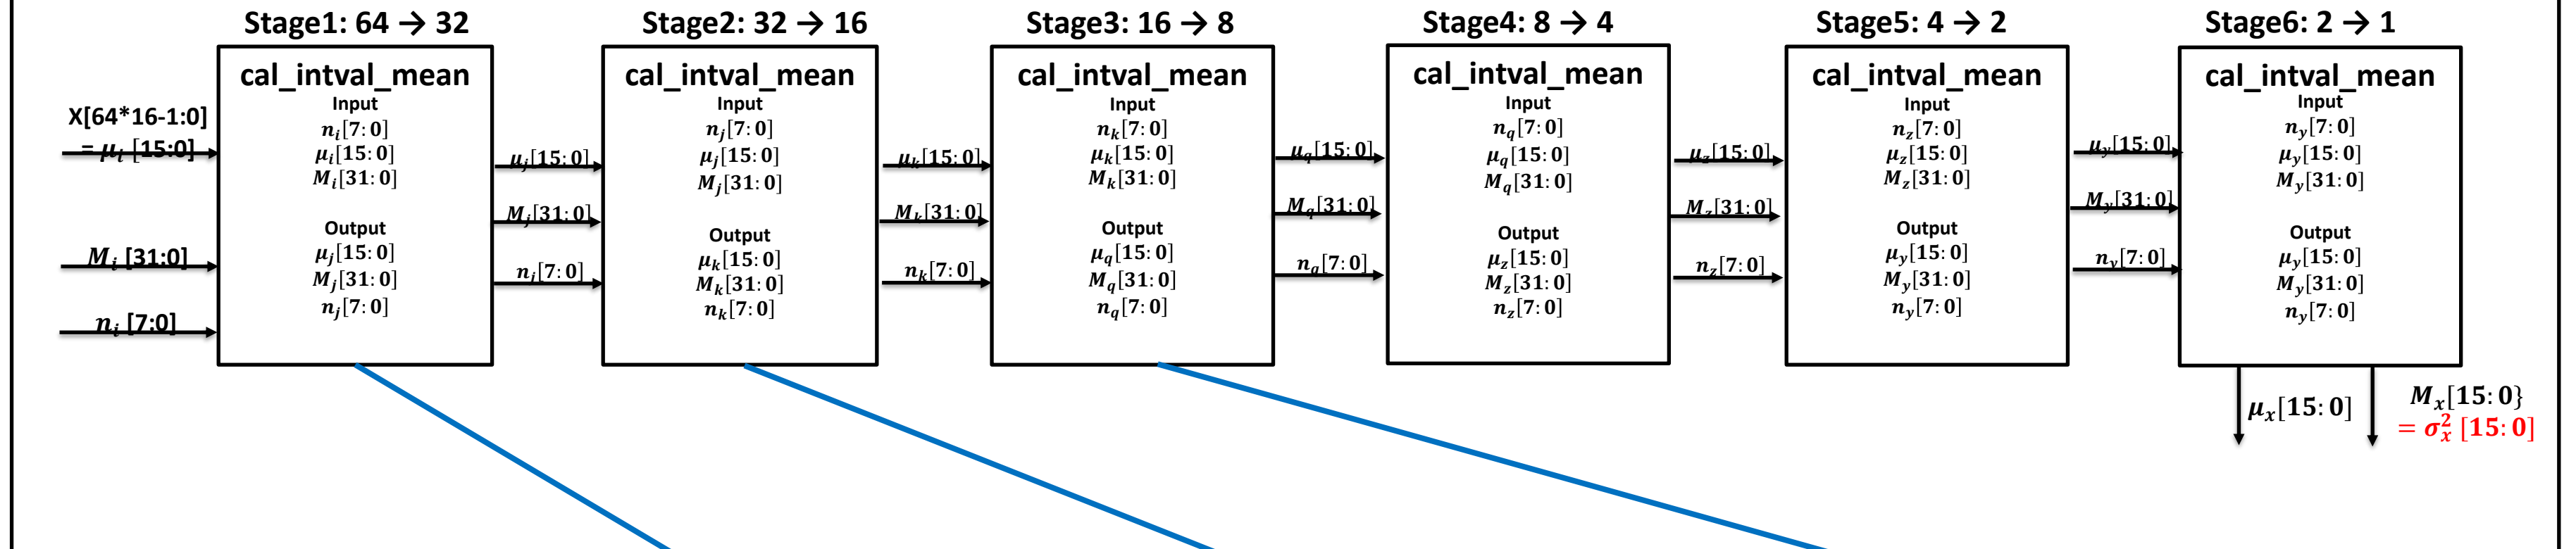
LN Approximation Model



Pairwise_module



Pairwise_module



```

module Pairwise_module (
    input wire clk,
    input wire rst_n,
    input wire in_valid,
    input wire [1023:0] a,
    output reg [15:0] mean_out,
    output reg [15:0] var_out,
    output reg out_valid
);

```

Stage1: 64 → 32

```

for (s = 0; s < 32; s = s + 1) begin : STAGE1
    cal_intval_mean u (
        .clk(clk), .rst_n(rst_n),
        .in_valid(in_valid),
        .out_valid(stage1_valid_bus[s]),
        .mu1(mu[2*s]), .mu2(mu[2*s+1]),
        .M1(M[2*s]), .M2(M[2*s+1]),
        .n1(n[2*s]), .n2(n[2*s+1]),
        .mu_out(mu_stage1[s]),
        .M_out(M_stage1[s]),
        .n_out(n_stage1[s])
    );
end

```

Stage2: 32 → 16

```

for (s = 0; s < 16; s = s + 1) begin : STAGE2
    cal_intval_mean u (
        .clk(clk), .rst_n(rst_n),
        .in_valid(stage1_valid),
        .out_valid(stage2_valid_bus[s]),
        .mu1(mu_stage1[2*s]), .mu2(mu_stage1[2*s+1]),
        .M1(M_stage1[2*s]), .M2(M_stage1[2*s+1]),
        .n1(n_stage1[2*s]), .n2(n_stage1[2*s+1]),
        .mu_out(mu_stage2[s]),
        .M_out(M_stage2[s]),
        .n_out(n_stage2[s])
    );
end

```

Stage3: 16 → 8

```

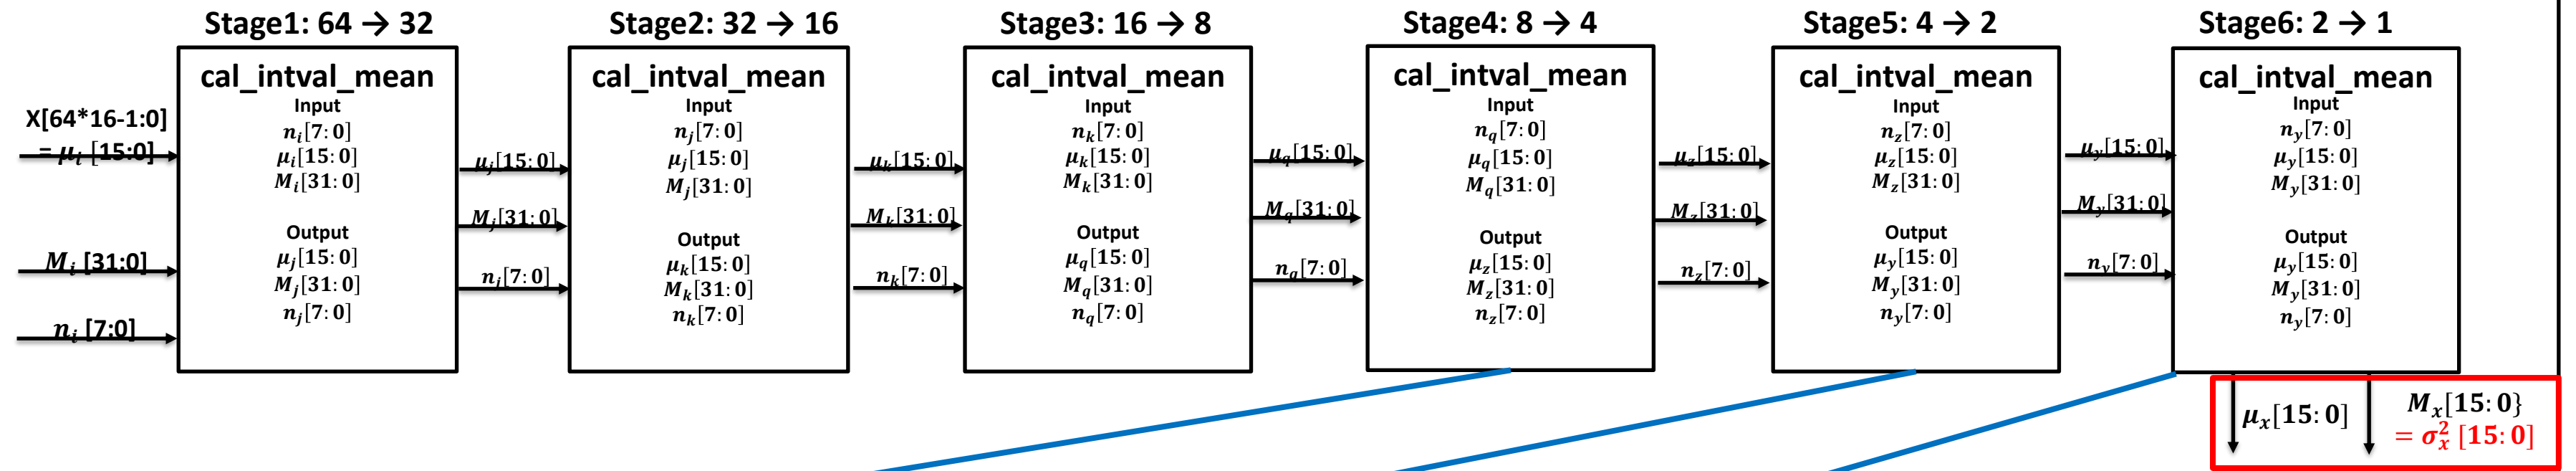
for (s = 0; s < 8; s = s + 1) begin : STAGE3
    cal_intval_mean u (
        .clk(clk), .rst_n(rst_n),
        .in_valid(stage2_valid),
        .out_valid(stage3_valid_bus[s]),
        .mu1(mu_stage2[2*s]), .mu2(mu_stage2[2*s+1]),
        .M1(M_stage2[2*s]), .M2(M_stage2[2*s+1]),
        .n1(n_stage2[2*s]), .n2(n_stage2[2*s+1]),
        .mu_out(mu_stage3[s]),
        .M_out(M_stage3[s]),
        .n_out(n_stage3[s])
    );
end

```

Stage4

end

Pairwise_module



Stage4: 8 → 4

```
for (s = 0; s < 4; s = s + 1) begin : STAGE4
    cal_intval_mean u (
        .clk(clk), .rst_n(rst_n),
        .in_valid(stage3_valid),
        .out_valid(stage4_valid_bus[s]),
        .mu1(mu_stage3[2*s]), .mu2(mu_stage3[2*s+1]),
        .M1(M_stage3[2*s]), .M2(M_stage3[2*s+1]),
        .n1(n_stage3[2*s]), .n2(n_stage3[2*s+1]),
        .mu_out(mu_stage4[s]),
        .M_out(M_stage4[s]),
        .n_out(n_stage4[s])
    );
end
```

Stage5: 4 → 2

```
for (s = 0; s < 2; s = s + 1) begin : STAGE5
    cal_intval_mean u (
        .clk(clk), .rst_n(rst_n),
        .in_valid(stage4_valid),
        .out_valid(stage5_valid_bus[s]),
        .mu1(mu_stage4[2*s]), .mu2(mu_stage4[2*s+1]),
        .M1(M_stage4[2*s]), .M2(M_stage4[2*s+1]),
        .n1(n_stage4[2*s]), .n2(n_stage4[2*s+1]),
        .mu_out(mu_stage5[s]),
        .M_out(M_stage5[s]),
        .n_out(n_stage5[s])
    );
end
```

Stage6: 2 → 1

```
cal_intval_mean u6 (
    .clk(clk), .rst_n(rst_n),
    .in_valid(stage5_valid),
    .out_valid(stage6_valid),
    .mu1(mu_stage5[0]), .mu2(mu_stage5[1]),
    .M1(M_stage5[0]), .M2(M_stage5[1]),
    .n1(n_stage5[0]), .n2(n_stage5[1]),
    .mu_out(mu_stage6),
    .M_out(M_stage6),
    .n_out()
);
```

```
// Output logic
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        mean_out <= 0;
        var_out <= 0;
        out_valid <= 0;
    end else if (stage6_valid_d1) begin
        mean_out <= mu_stage6;
        var_out <= M_stage6[23:8];
        out_valid <= 1'b1;
    end else begin
        out_valid <= 1'b0;
    end
end
```

LN Approximation Model

cal_intval_mean

- (1)

delta = mu1 - mu2

ADD
- (2)

delta2 = delta * delta

MULT
- (3)

cross = (n1 * n2 * delta2) / (n1 + n2)

ADD/MULT/SHIFT/LOG2
- (4)

mu_acc = (mu1 * n1 + mu2 * n2)

ADD/MULT
- (5)

M_out = M1 + M2 + cross

ADD
- (6)

mu_out = mu_acc / (n1 + n2)

ADD/SHIFT/LOG2
- (7)

n_out = n1 + n2

ADD

ADD(SUB)	c_addsub_add
MULT	mult_gen_mult
LOG2	Using LUT

```
module cal_intval_mean(  
    input wire clk,  
    input wire rst_n,  
    input wire in_valid,  
    output reg out_valid,  
  
    input wire [15:0] mu1,  
    input wire [15:0] mu2,  
    input wire [31:0] M1,  
    input wire [31:0] M2,  
    input wire [7:0] n1,  
    input wire [7:0] n2,  
  
    output reg [15:0] mu_out,  
    output reg [31:0] M_out,  
    output reg [7:0] n_out  
);  
  
    wire [31:0] n1_q16 = {n1, 16'b0};  
    wire [31:0] n2_q16 = {n2, 16'b0};  
    wire [31:0] mu1_q16 = {mu1, 8'b0};  
    wire [31:0] mu2_q16 = {mu2, 8'b0};  
  
    wire signed [31:0] delta;  
    wire [31:0] delta2;  
    wire [31:0] n1n2;  
    wire [31:0] cross_numerator;  
    wire [31:0] cross;  
    wire [31:0] mu1_n1, mu2_n2;  
    wire [31:0] mu_acc;  
    wire [31:0] M_tmp;  
    wire [31:0] M_sum;  
    wire [7:0] n_sum;  
    wire [31:0] mu_total;  
    wire [15:0] mu_final;  
    wire [3:0] shift_n;
```

- (1)

c_addsub_sub delta_sub (
 .A(mu1_q16), .B(mu2_q16), .CLK(clk), .CE(1'b1), .S(delta)
);

c_addsub_add muacc_add (
 .A(mu1_n1), .B(mu2_n2), .CLK(clk), .CE(1'b1), .S(mu_acc)
);
- (2)

mult_gen_mult delta_square_mul (
 .A(delta), .B(delta), .CLK(clk), .P(delta2)
);
- (3)

mult_gen_mult n1n2_mul (
 .A(n1_q16), .B(n2_q16), .CLK(clk), .P(n1n2)
);

 mult_gen_mult cross_num_mul (
 .A(n1n2), .B(delta2), .CLK(clk), .P(cross_numerator)
);

(4)

mult_gen_mult mu1n1_mul (
 .A(mu1_q16), .B(n1_q16), .CLK(clk), .P(mu1_n1)
);

 mult_gen_mult mu2n2_mul (
 .A(mu2_q16), .B(n2_q16), .CLK(clk), .P(mu2_n2)
);

(5)

c_addsub_add M_add (
 .A(M1), .B(M2), .CLK(clk), .CE(1'b1), .S(M_tmp)
);

 c_addsub_add M_cross_add (
 .A(M_tmp), .B(cross), .CLK(clk), .CE(1'b1), .S(M_sum)
);

(7)

c_addsub_add8 n_sum_add8 (
 .A(n1), .B(n2), .CLK(clk), .CE(1'b1), .S(n_sum)
);

 log2_lut u_log2 (
 .n_sum(n_sum), .log2_val(shift_n)
);

(3)

assign cross = cross_numerator >> shift_n;


(6)

assign mu_total = mu_acc >> shift_n;
assign mu_final = mu_total[23:8];


LN Approximation Model

(1) Input: 0x0080 0x0100 0x0180 ... 0x2000 (0.5 1.0 1.5 ... 32.0)

LATECNCY : 36




(2) Input: 32 values of 0x0000 + 32 values of 0x1fff (31.966)

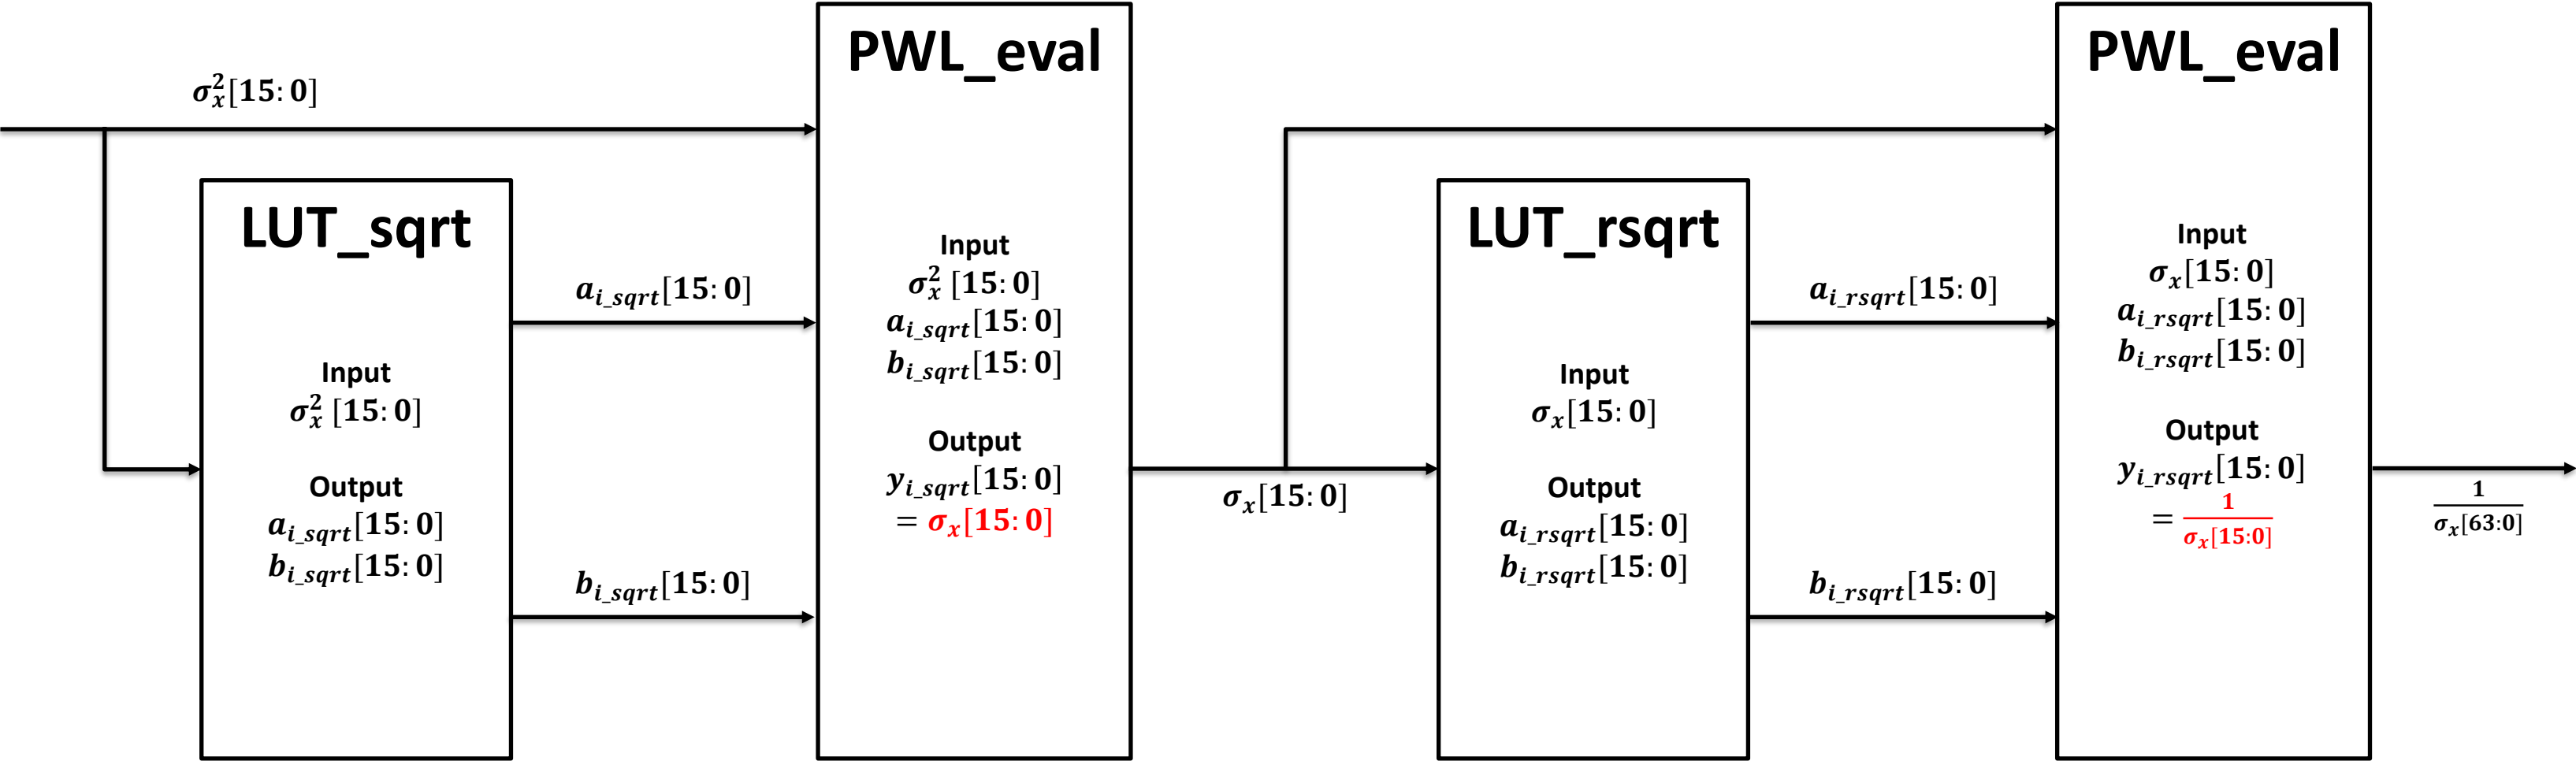


Q8.8 input hex (Dec)	(1)	(2)	Accuracy
Mean	0x1040 (16.25)	0x0fff (15.996)	100% (error=0%)
Mean expected	0x1040 (16.25)	0x0fff (15.996)	
Variance	0x5400 (84.0)	0xfc00 (252.0)	98.46% (error=1.54%)
Variance expected	0x5550 (85.313)	0xffb1 (255.246)	

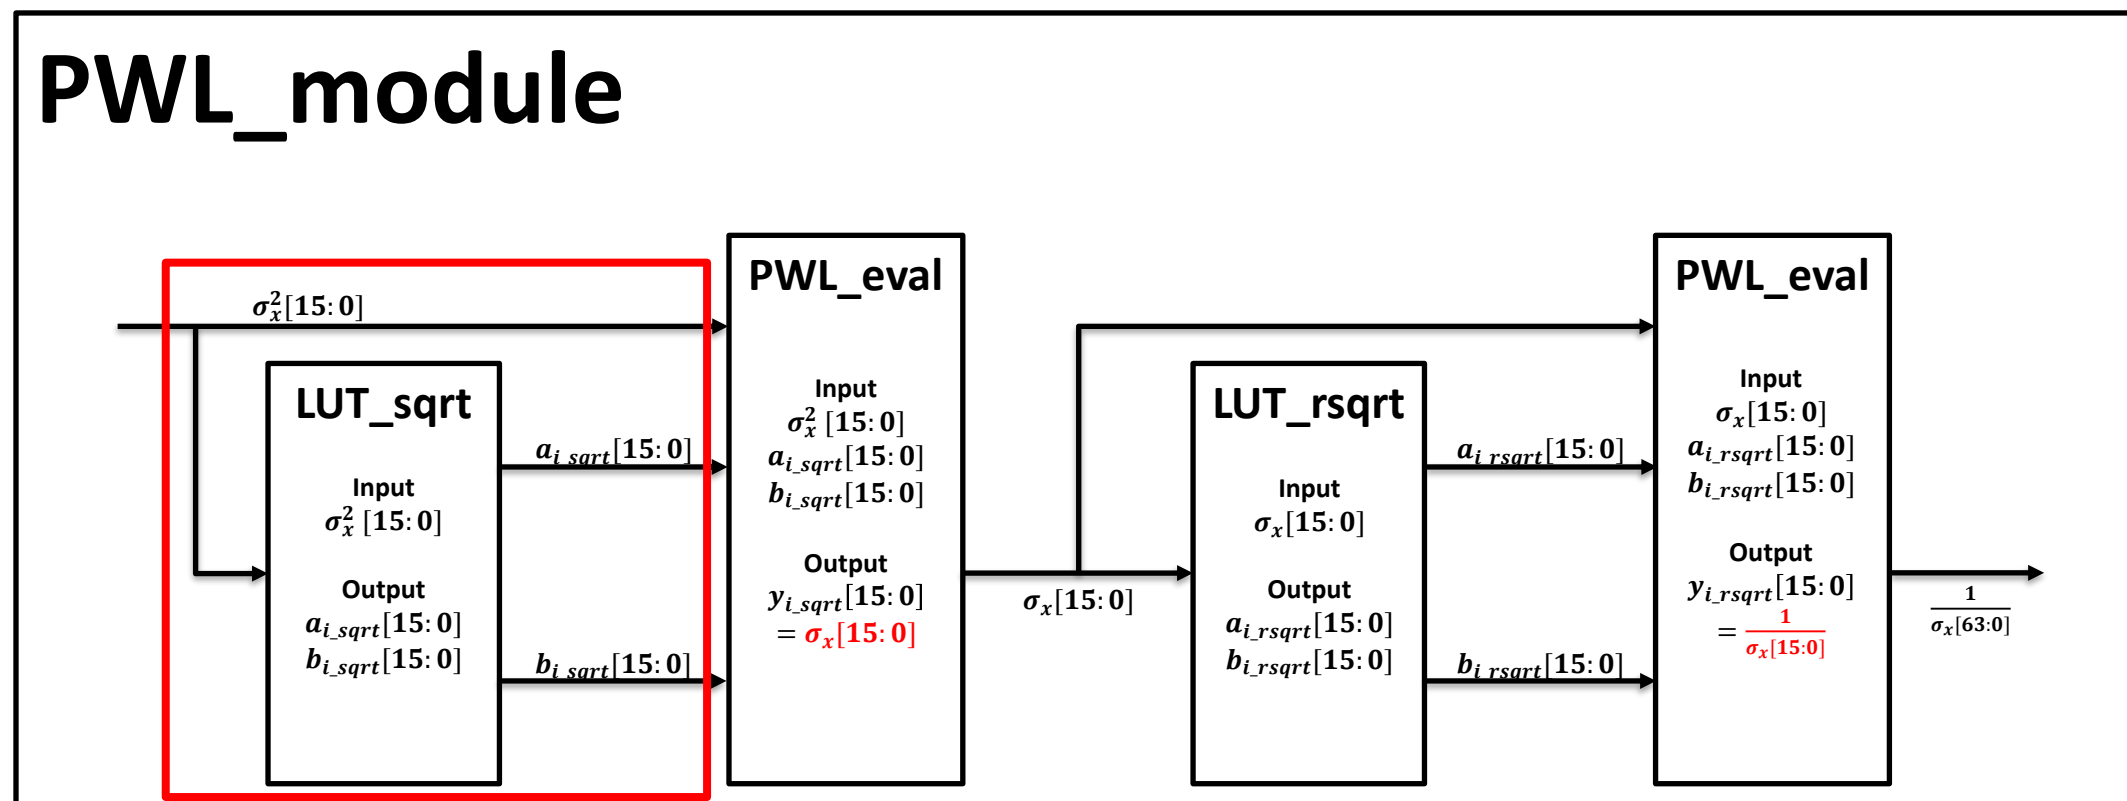
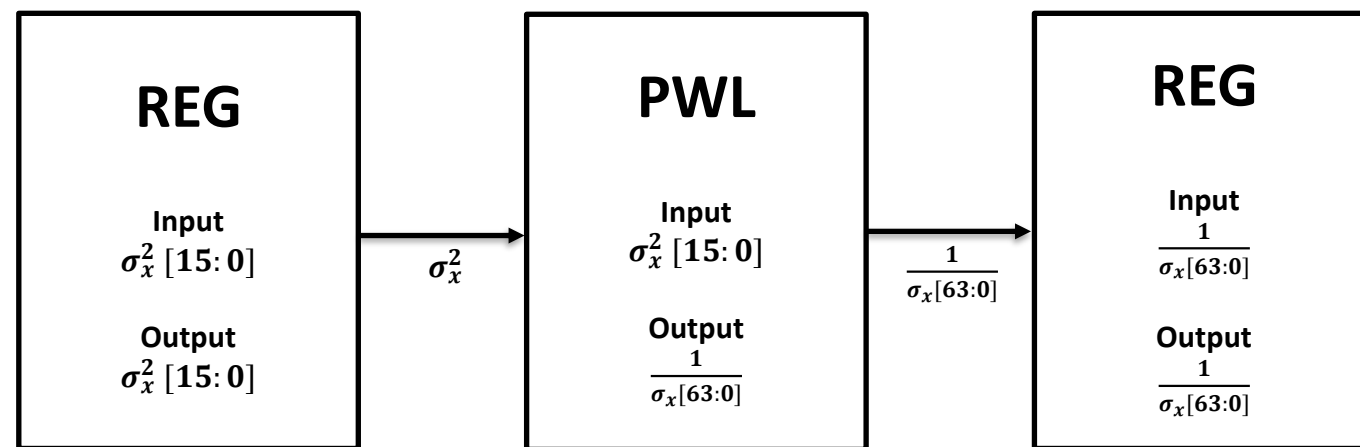
LN Approximation Model



PWL_module



LN Approximation Model



```

module LUT_sqrt (
    input wire clk,
    input wire rst_n,
    input wire in_valid,
    input wire [15:0] x_in,
    output reg out_valid,
    output reg [15:0] slope_out,
    output reg [15:0] intercept_out
);

localparam integer N_SEGMENTS = 8;

reg [15:0] breakpoints [0:N_SEGMENTS];
reg [15:0] slopes [0:N_SEGMENTS-1];
reg [15:0] intercepts [0:N_SEGMENTS-1];

reg [2:0] region;

// internal wire
reg [2:0] next_region;
reg region_valid;
  
```

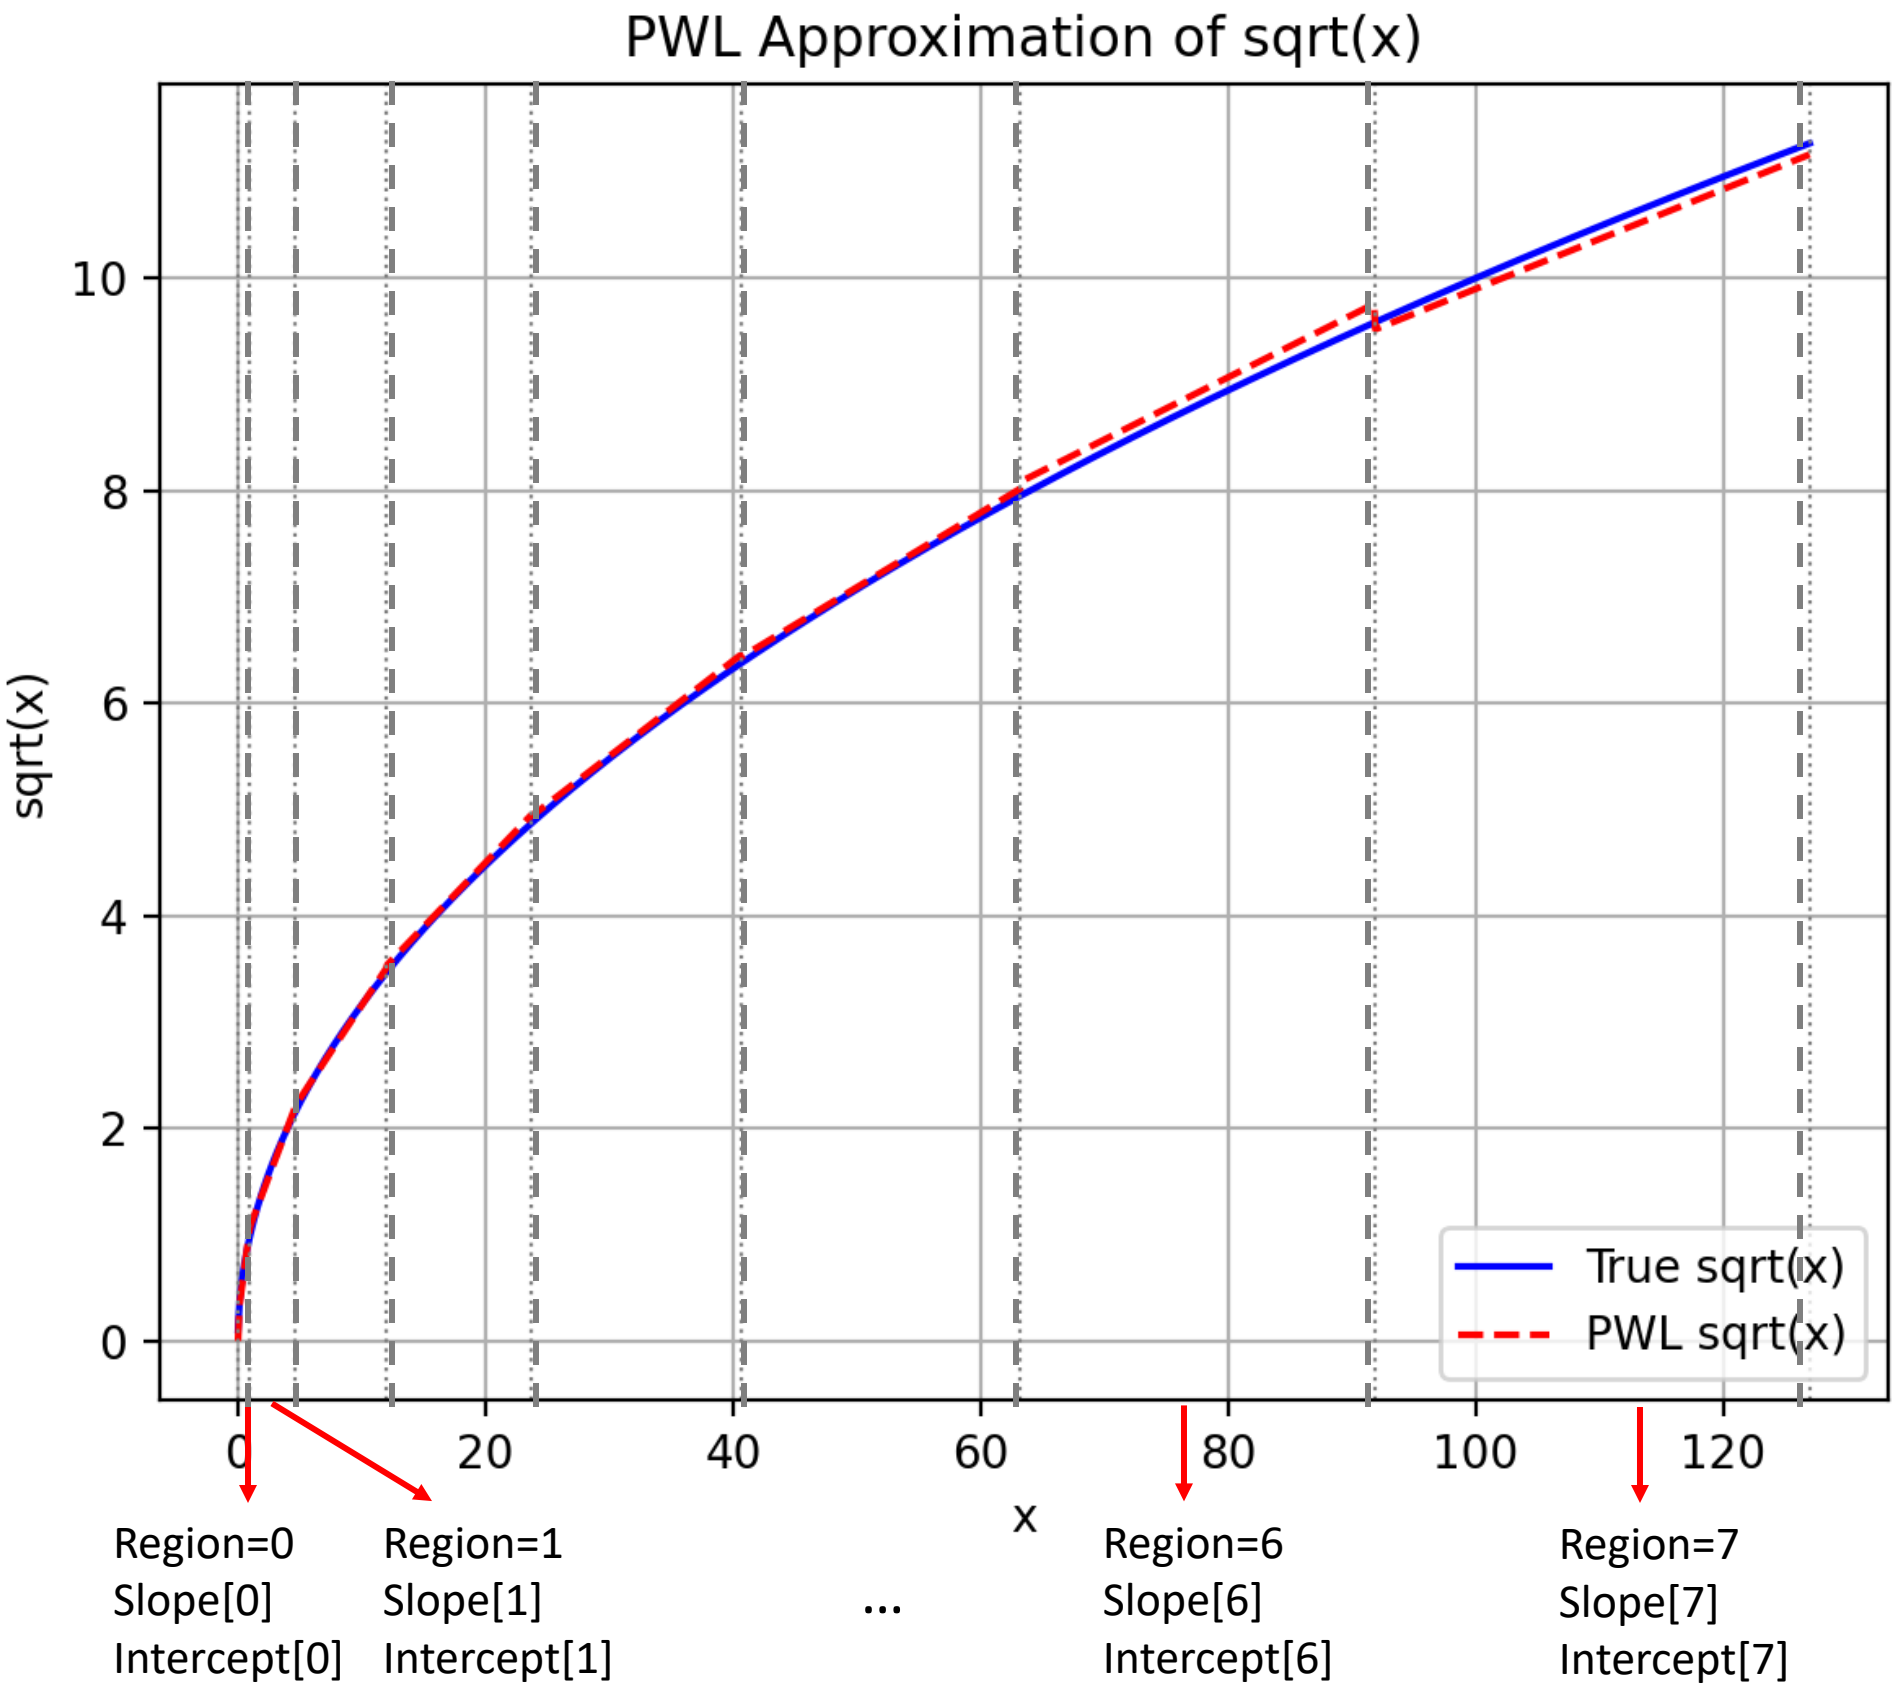
```

// Region decision
always @(*) begin
    if (x_in < breakpoints[1]) next_region = 0;
    else if (x_in < breakpoints[2]) next_region = 1;
    else if (x_in < breakpoints[3]) next_region = 2;
    else if (x_in < breakpoints[4]) next_region = 3;
    else if (x_in < breakpoints[5]) next_region = 4;
    else if (x_in < breakpoints[6]) next_region = 5;
    else if (x_in < breakpoints[7]) next_region = 6;
    else next_region = 7;
end

// Save Region
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        region <= 0;
        region_valid <= 0;
    end else begin
        if (in_valid) begin
            region <= next_region;
            region_valid <= 1;
        end else begin
            region_valid <= 0;
        end
    end
end

// Output
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        slope_out <= 0;
        intercept_out <= 0;
        out_valid <= 0;
    end else begin
        slope_out <= slopes[region];
        intercept_out <= intercepts[region];
        out_valid <= region_valid;
    end
end
  
```

LN Approximation Model



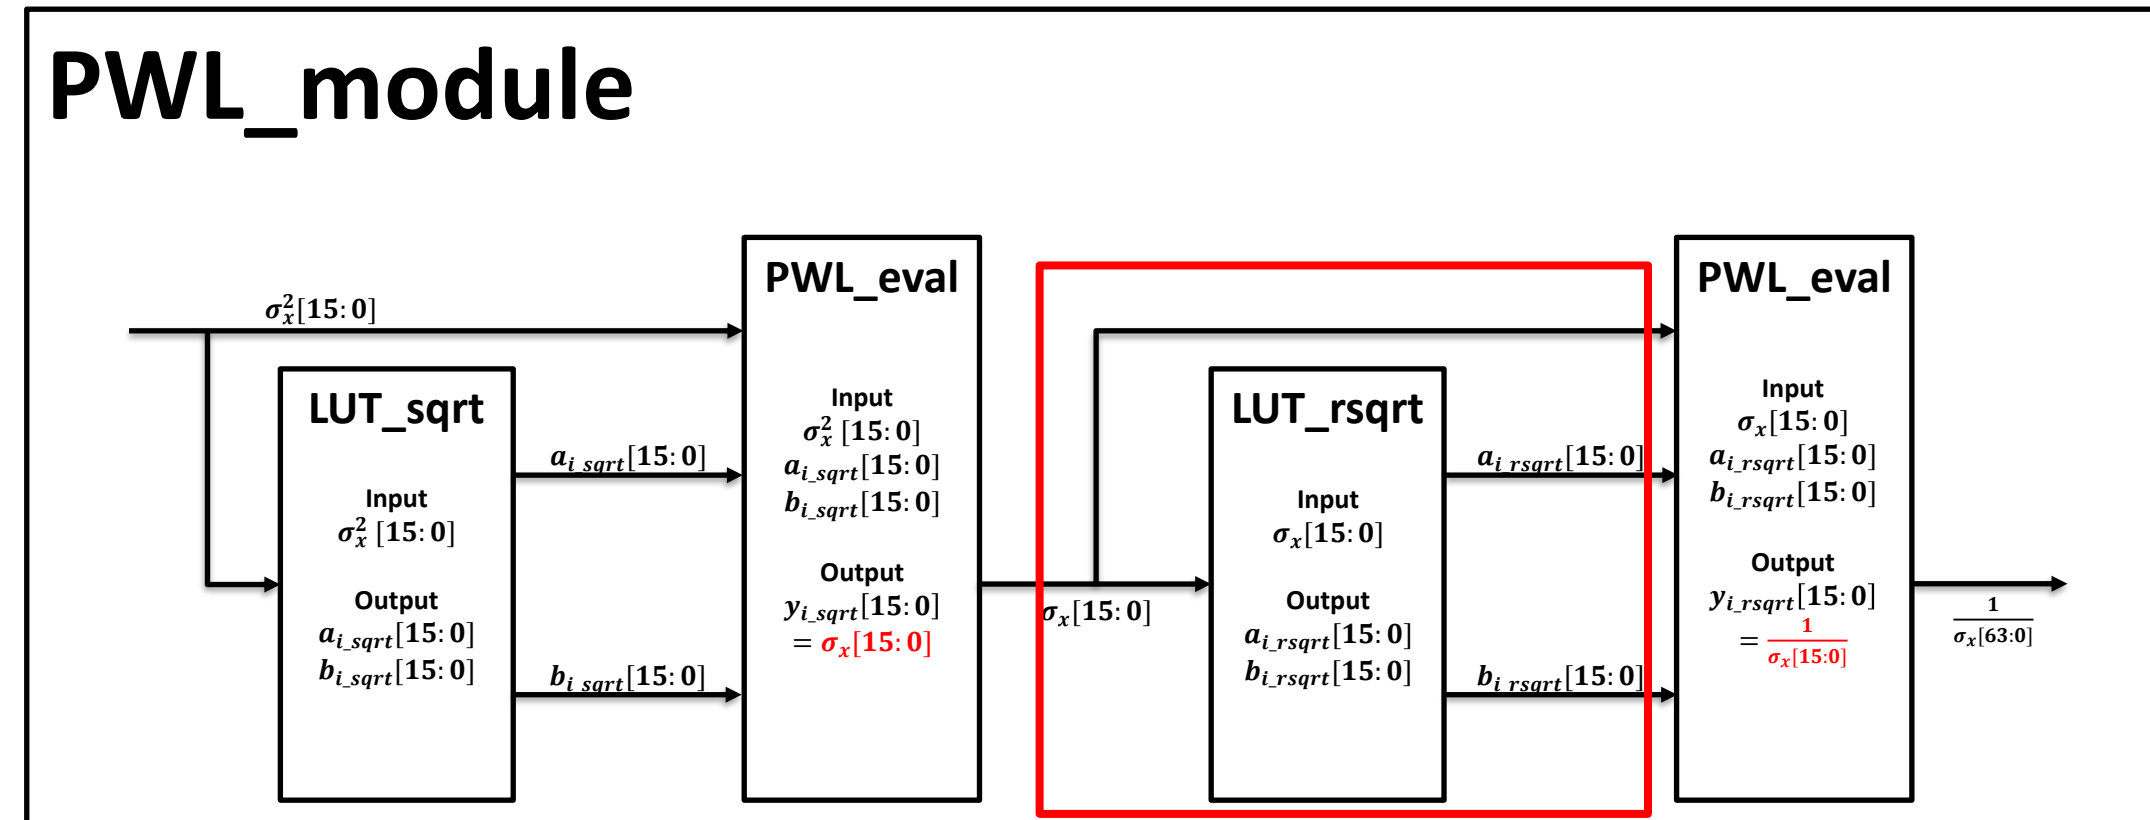
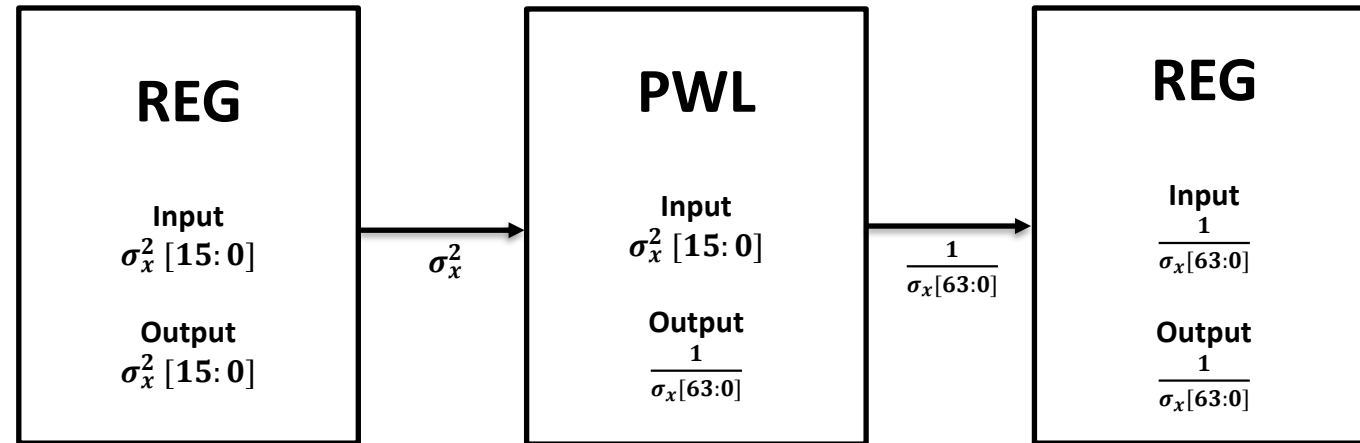
$$\text{sqrt}(x) = x * \text{slopes}[i] + \text{intercepts}[i]$$

```
initial begin
    breakpoints[0] = 16'h0003; // 0.0117
    breakpoints[1] = 16'h00F7; // 0.9648
    breakpoints[2] = 16'h04A9; // 4.6602
    breakpoints[3] = 16'h0C03; // 12.0117
    breakpoints[4] = 16'h17BF; // 23.7461
    breakpoints[5] = 16'h28A1; // 40.6328
    breakpoints[6] = 16'h3F38; // 63.2188
    breakpoints[7] = 16'h5BDB; // 91.8555
    breakpoints[8] = 16'h7F00; // 127.0000

    slopes[0] = 16'h00E2; // 0.8828
    slopes[1] = 16'h004F; // 0.3086
    slopes[2] = 16'h002D; // 0.1758
    slopes[3] = 16'h001F; // 0.1211
    slopes[4] = 16'h0017; // 0.0898
    slopes[5] = 16'h0012; // 0.0703
    slopes[6] = 16'h000F; // 0.0586
    slopes[7] = 16'h000C; // 0.0469

    intercepts[0] = 16'h0038; // 0.2188
    intercepts[1] = 16'h00C6; // 0.7734
    intercepts[2] = 16'h0166; // 1.3984
    intercepts[3] = 16'h0214; // 2.0781
    intercepts[4] = 16'h02CE; // 2.8086
    intercepts[5] = 16'h0393; // 3.5742
    intercepts[6] = 16'h0461; // 4.3789
    intercepts[7] = 16'h0536; // 5.2109
end
```

LN Approximation Model



```
module LUT_rsqr (
    input wire clk,
    input wire rst_n,
    input wire in_valid,
    input wire [15:0] x_in,
    output reg out_valid,
    output reg [15:0] slope_out,
    output reg [15:0] intercept_out
);

    localparam integer N_SEGMENTS = 8;

    reg [15:0] breakpoints [0:N_SEGMENTS];
    reg [15:0] slopes [0:N_SEGMENTS-1];
    reg [15:0] intercepts [0:N_SEGMENTS-1];

    reg [2:0] region;
    reg region_valid;
    reg [2:0] next_region;
```

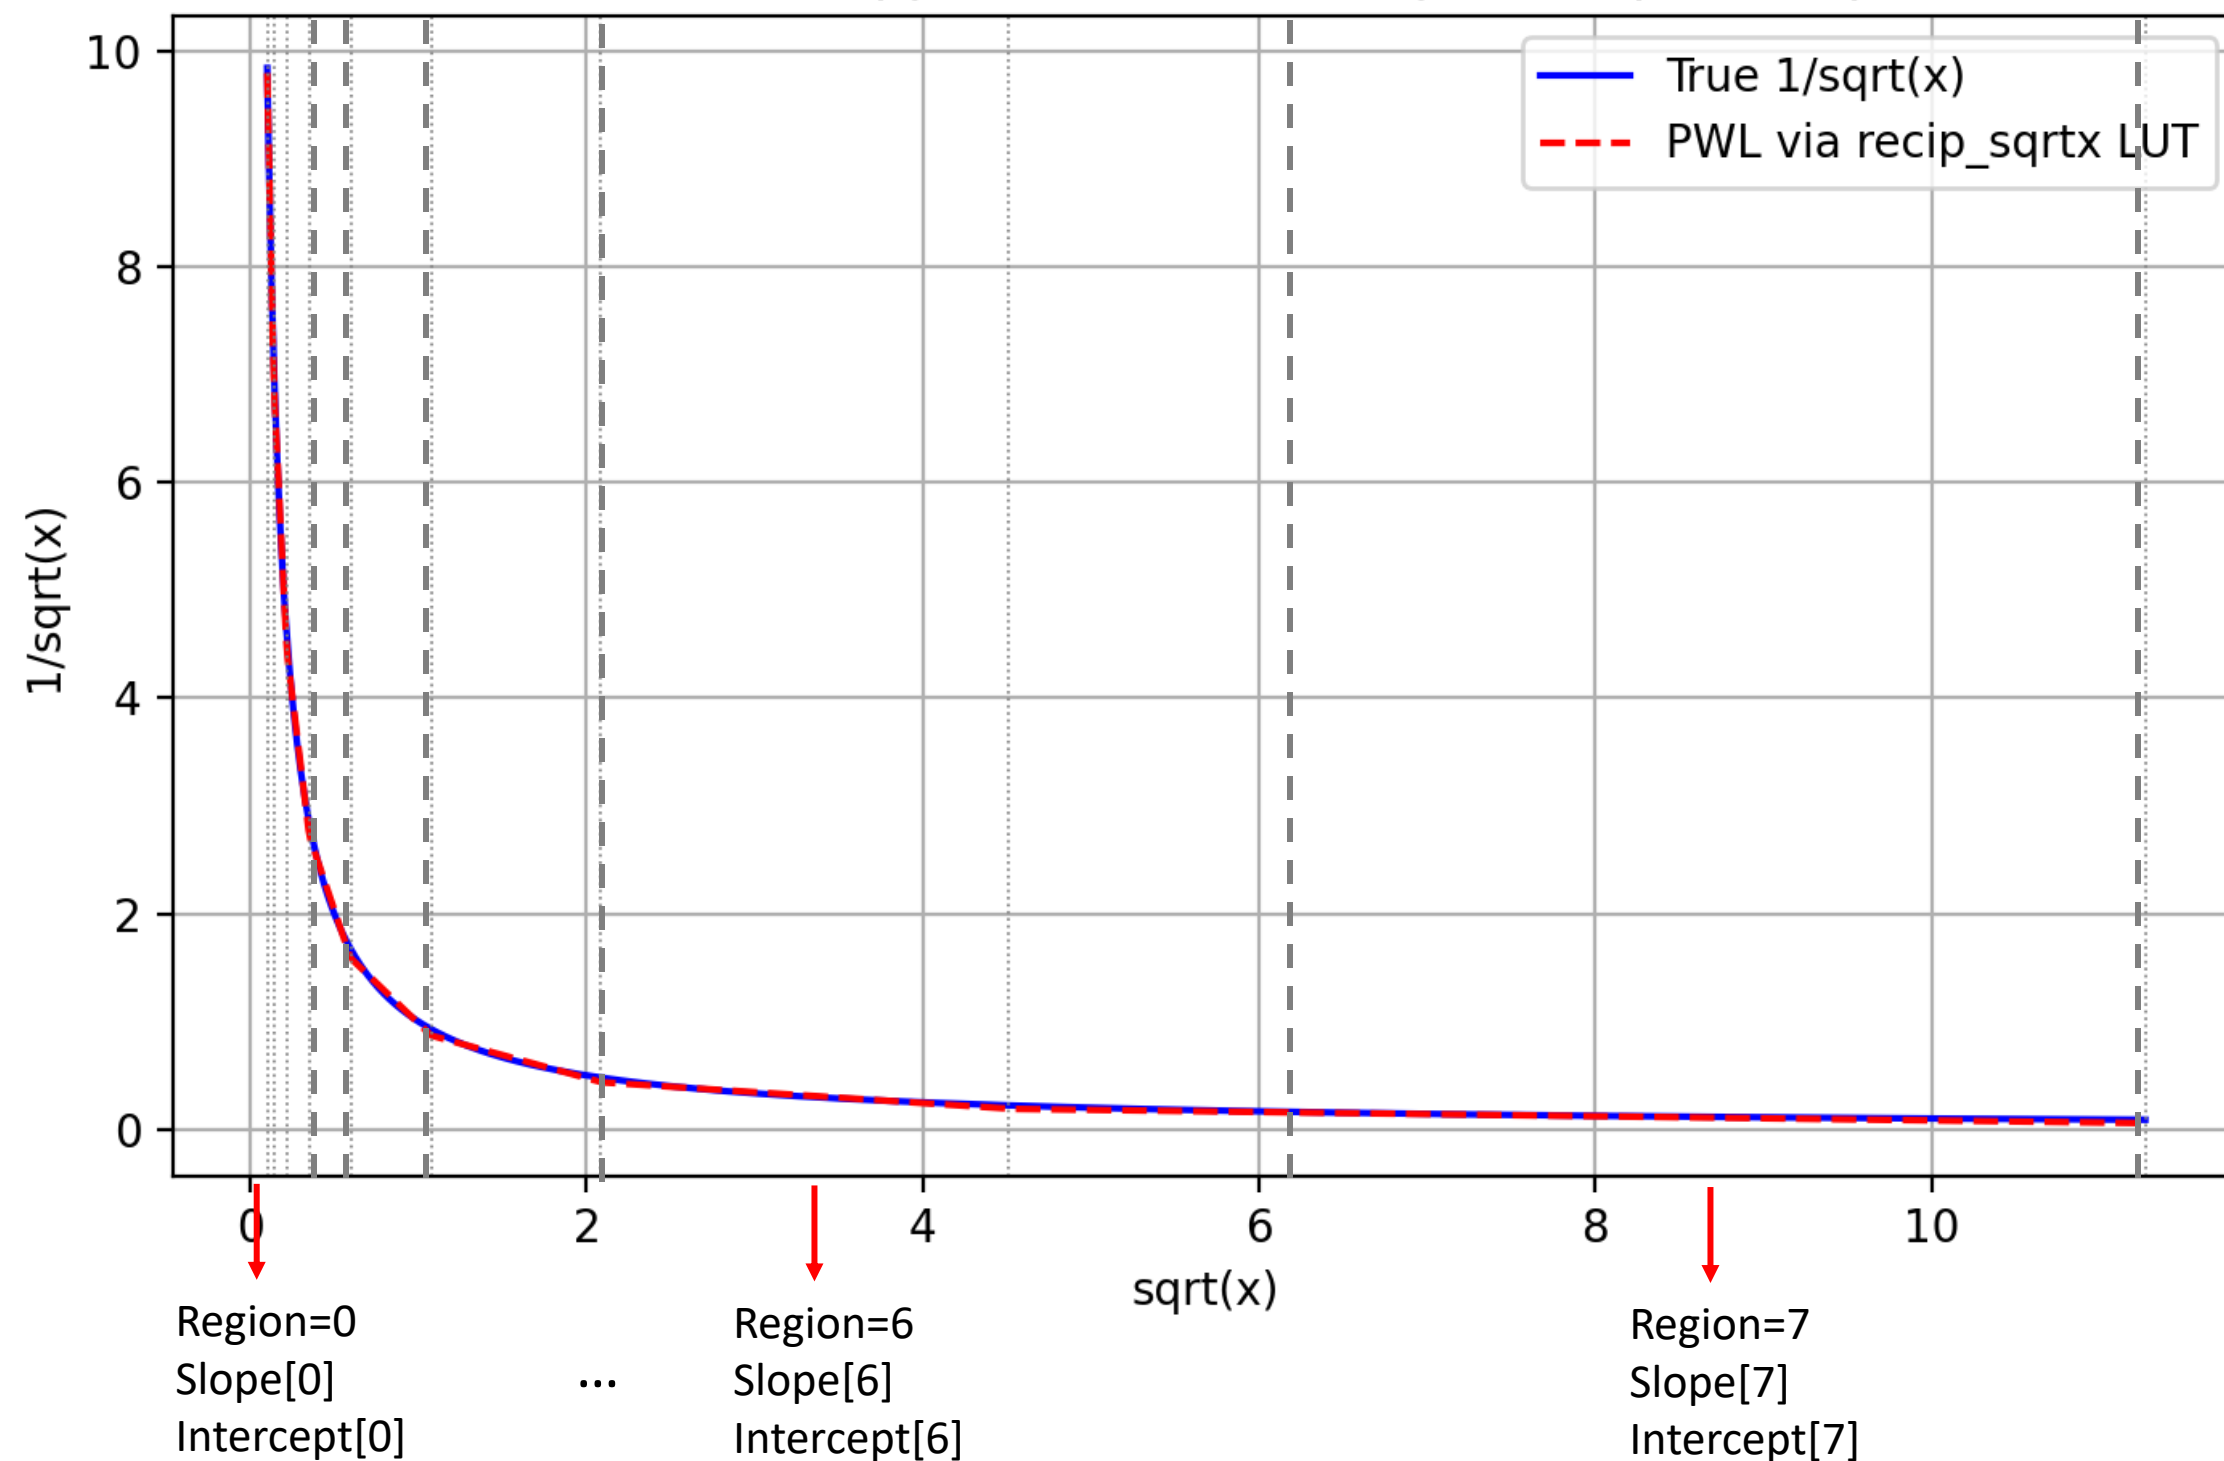
```
// Region decision (combinational)
always @(*) begin
    if (x_in < breakpoints[1]) next_region = 0;
    else if (x_in < breakpoints[2]) next_region = 1;
    else if (x_in < breakpoints[3]) next_region = 2;
    else if (x_in < breakpoints[4]) next_region = 3;
    else if (x_in < breakpoints[5]) next_region = 4;
    else if (x_in < breakpoints[6]) next_region = 5;
    else if (x_in < breakpoints[7]) next_region = 6;
    else next_region = 7;
end

// Save region (1 cycle)
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        region <= 0;
        region_valid <= 0;
    end else if (in_valid) begin
        region <= next_region;
        region_valid <= 1;
    end else begin
        region_valid <= 0;
    end
end

// Output (1 cycle)
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        slope_out <= 0;
        intercept_out <= 0;
        out_valid <= 0;
    end else begin
        slope_out <= slopes[region];
        intercept_out <= intercepts[region];
        out_valid <= region_valid;
    end
end
```

LN Approximation Model

LUT-based PWL Approximation of $1/\sqrt{x}$ (input = \sqrt{x})



$$1/\sqrt{x} = \sqrt{x} * \text{slopes}[i] + \text{intercepts}[i]$$

initial begin

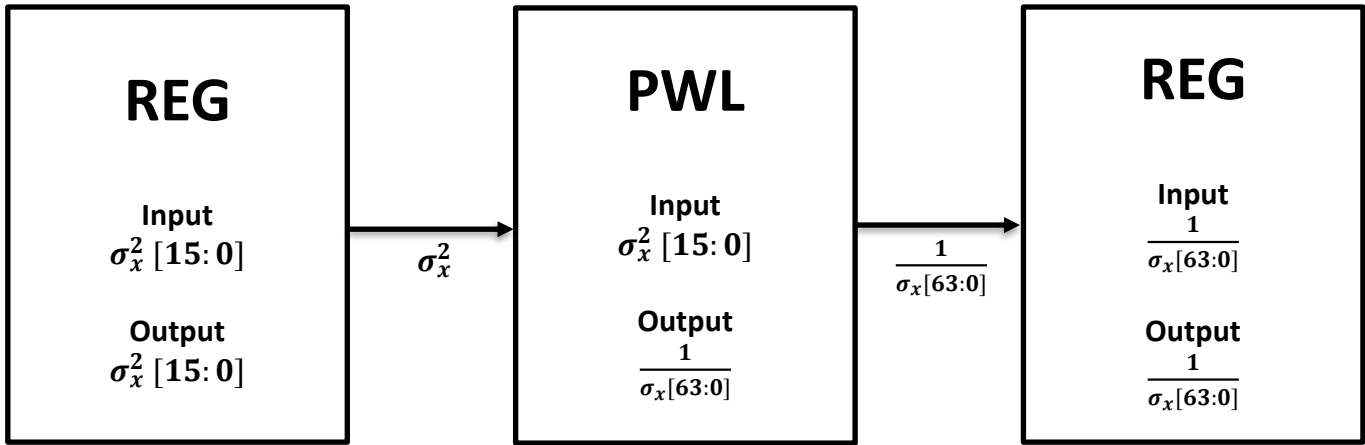
```
breakpoints[0] = 16'h001A; // 0.1016  
breakpoints[1] = 16'h0024; // 0.1406  
breakpoints[2] = 16'h0038; // 0.2188  
breakpoints[3] = 16'h005A; // 0.3516  
breakpoints[4] = 16'h0099; // 0.5977  
breakpoints[5] = 16'h0114; // 1.0703  
breakpoints[6] = 16'h0216; // 2.0898  
breakpoints[7] = 16'h0482; // 4.5078  
breakpoints[8] = 16'h0B45; // 11.2695
```

```
slopes[0] = 16'hB541; // -74.7461 → signed  
slopes[1] = 16'hDFA3; // -32.3633  
slopes[2] = 16'hF339; // -12.7773  
slopes[3] = 16'hFB5F; // -4.6289  
slopes[4] = 16'hFE80; // -1.5  
slopes[5] = 16'hFF93; // -0.4258  
slopes[6] = 16'hFFE6; // -0.1016  
slopes[7] = 16'hFFFB; // -0.0195
```

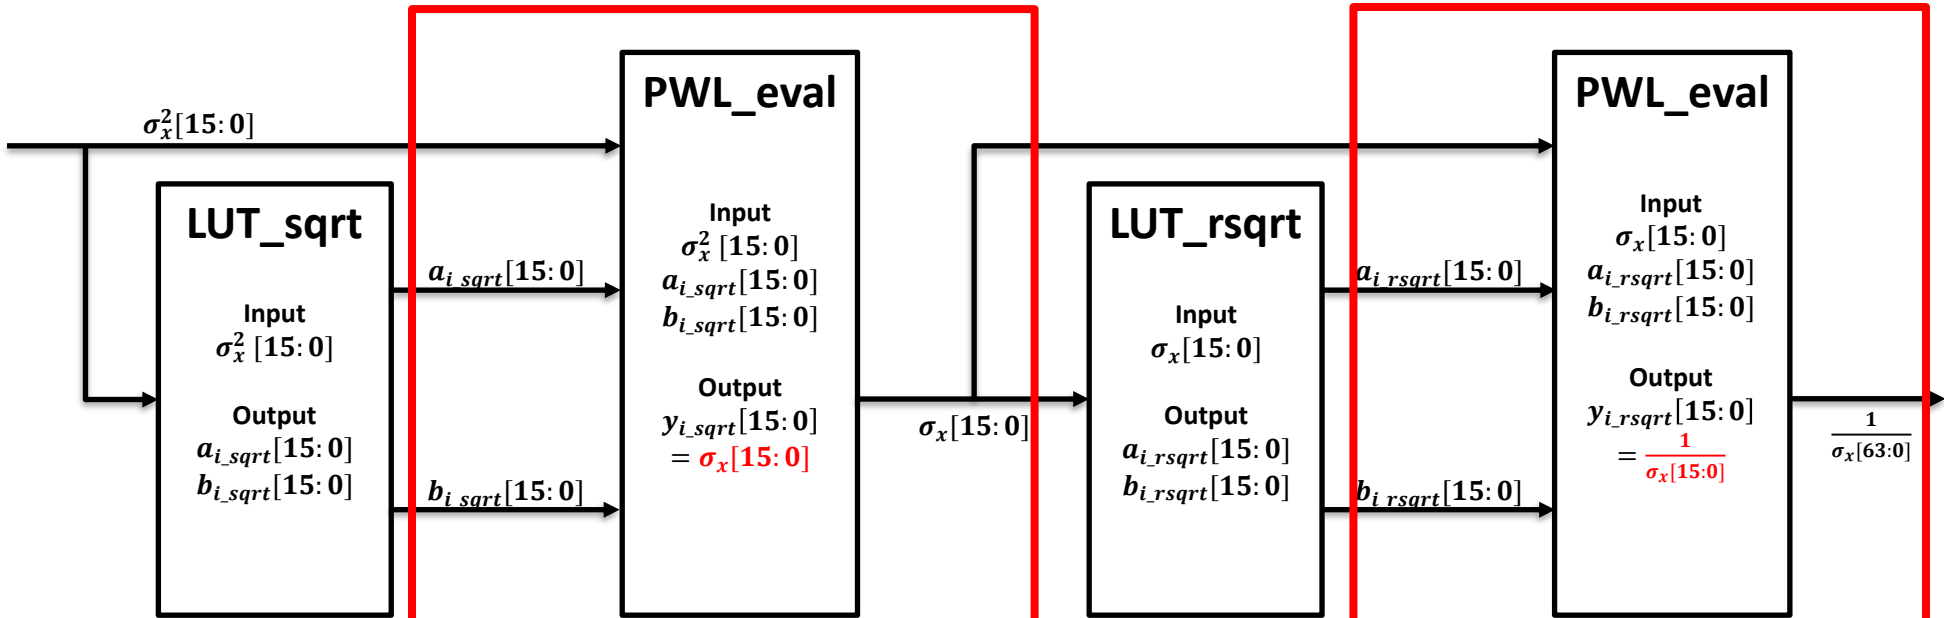
```
intercepts[0] = 16'h1164; // 17.3906  
intercepts[1] = 16'h0B78; // 11.4688  
intercepts[2] = 16'h0738; // 7.2188  
intercepts[3] = 16'h045A; // 4.3398  
intercepts[4] = 16'h027C; // 2.4844  
intercepts[5] = 16'h0154; // 1.3281  
intercepts[6] = 16'h00A6; // 0.6484  
intercepts[7] = 16'h0047; // 0.2773
```

end

LN Approximation Model



PWL_module



```
module PWL_eval #(parameter N = 16)(
    input wire clk,
    input wire rst_n,
    input wire in_valid,
    input wire signed [N-1:0] x_in, // Q8.8 input value
    input wire signed [N-1:0] slope, // Slope from LUT
    input wire signed [N-1:0] intercept, // Intercept from LUT
    output reg signed [N-1:0] y_out, // Result: y = slope * x + intercept
    output reg out_valid
);
```

```
// Multiplication result (Q8.8 * Q8.8 = Q16.16 → convert back to Q8.8)
```

```
wire signed [N-1:0] mult_result;
```

```
mult_gen_mult mul (
    .CLK(clk), .A(slope), .B(x_in), .P(mult_result)
);
```

```
// Addition using CLA (Carry Lookahead Adder)
```

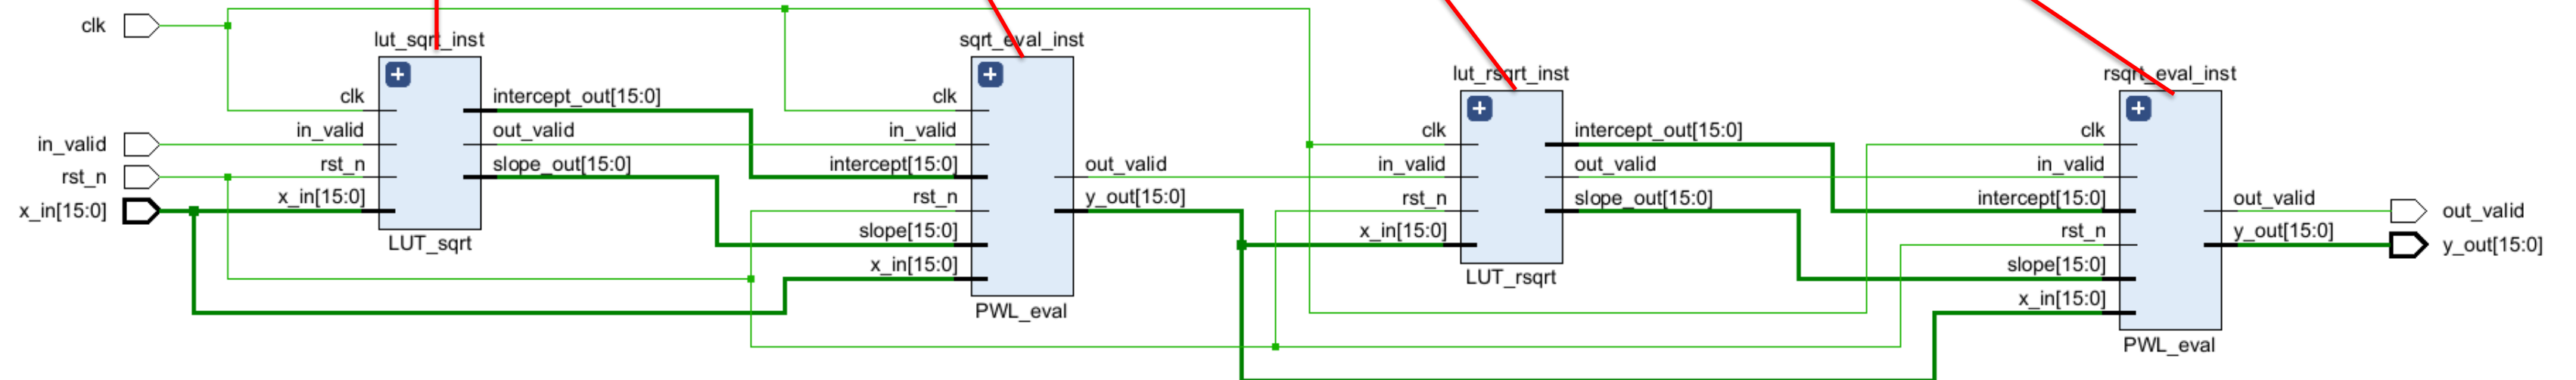
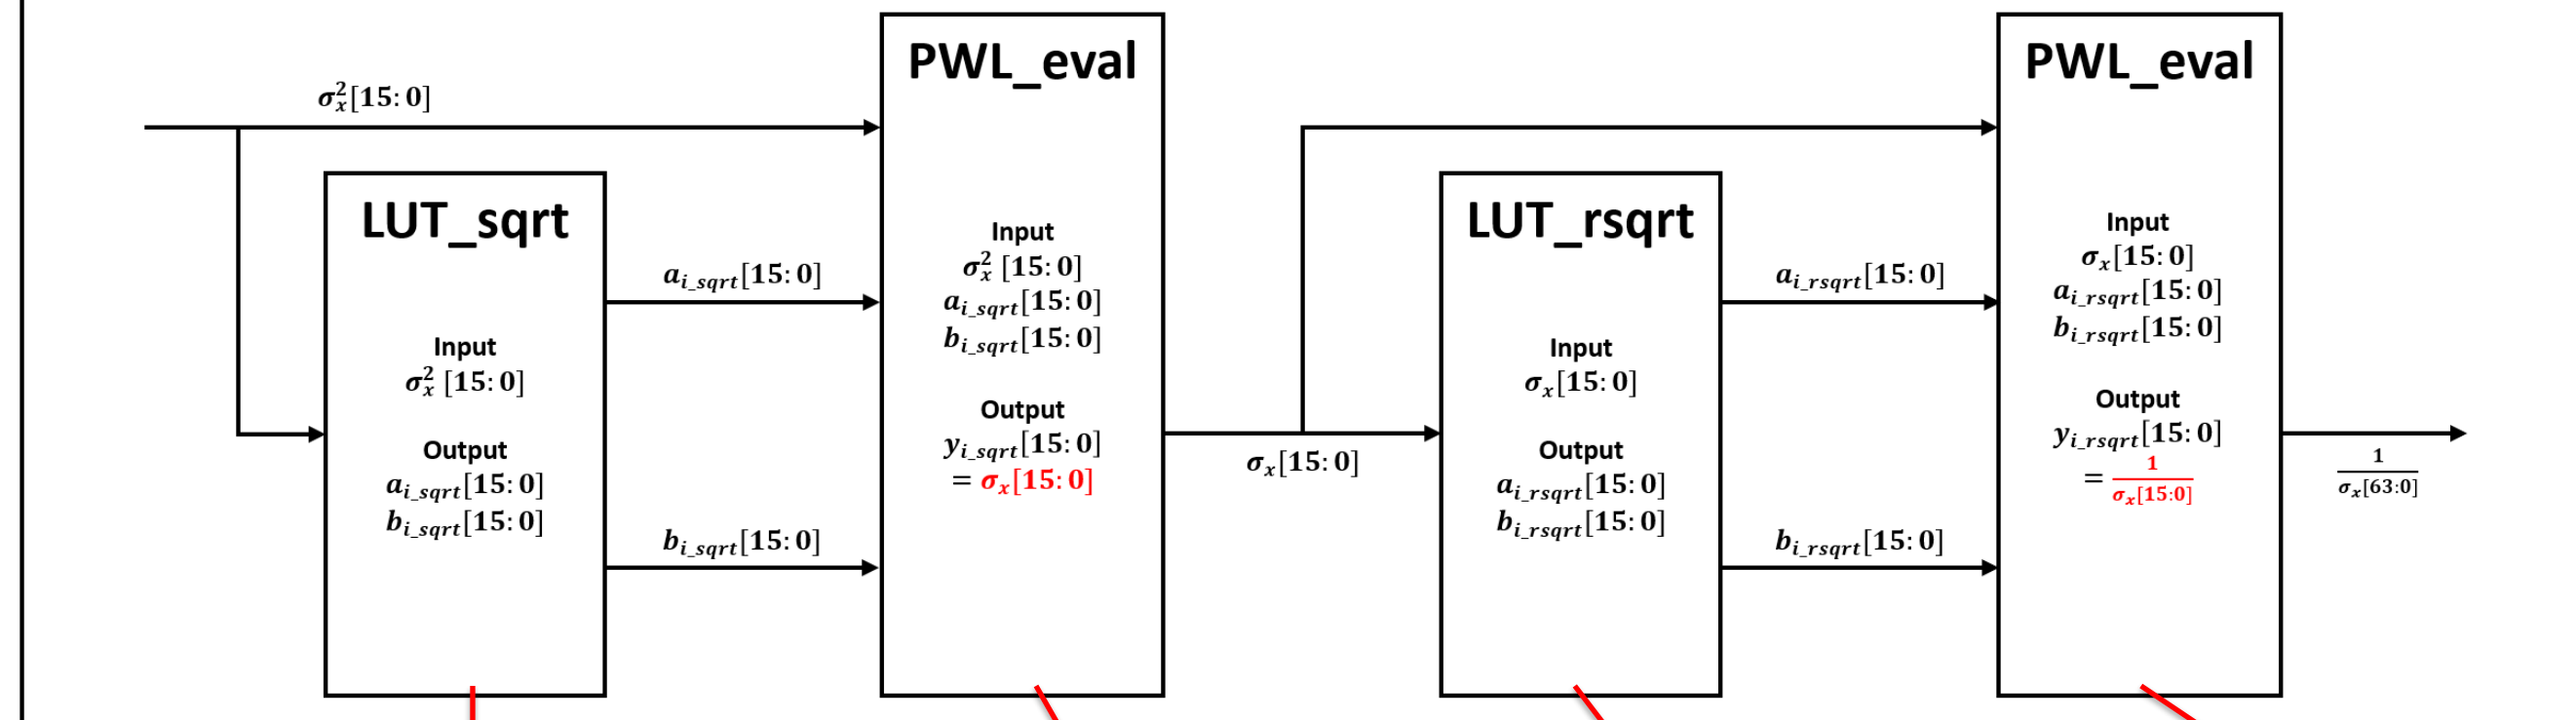
```
wire signed [N-1:0] sum_result;
```

```
c_addsub_add adder (
    .A(mult_result), .B(intercept), .CLK(clk), .CE(1'b1), .S(sum_result)
);
```


```
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        y_out <= 0;
        out_valid <= 0;
    end else begin
        if (valid_pipe[1]) begin
            y_out <= sum_result;
            out_valid <= 1'b1;
        end else begin
            out_valid <= 1'b0;
        end
    end
end
```

LN Approximation Model

PWL_module

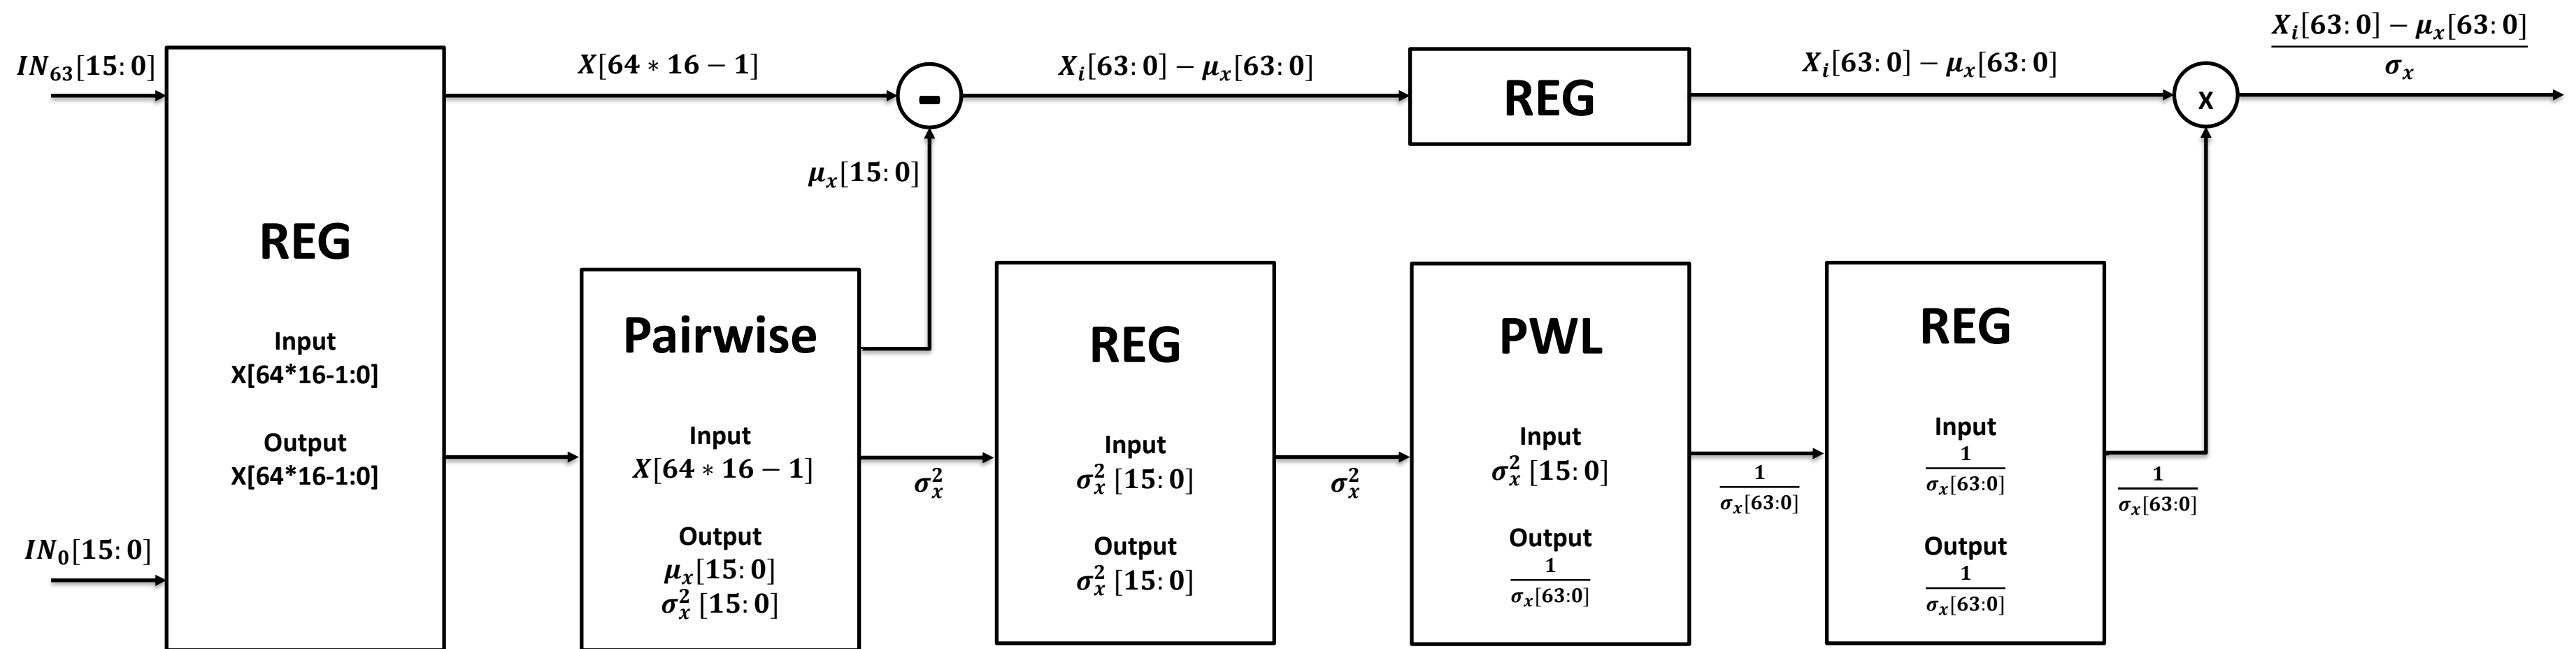


LN Approximation Model

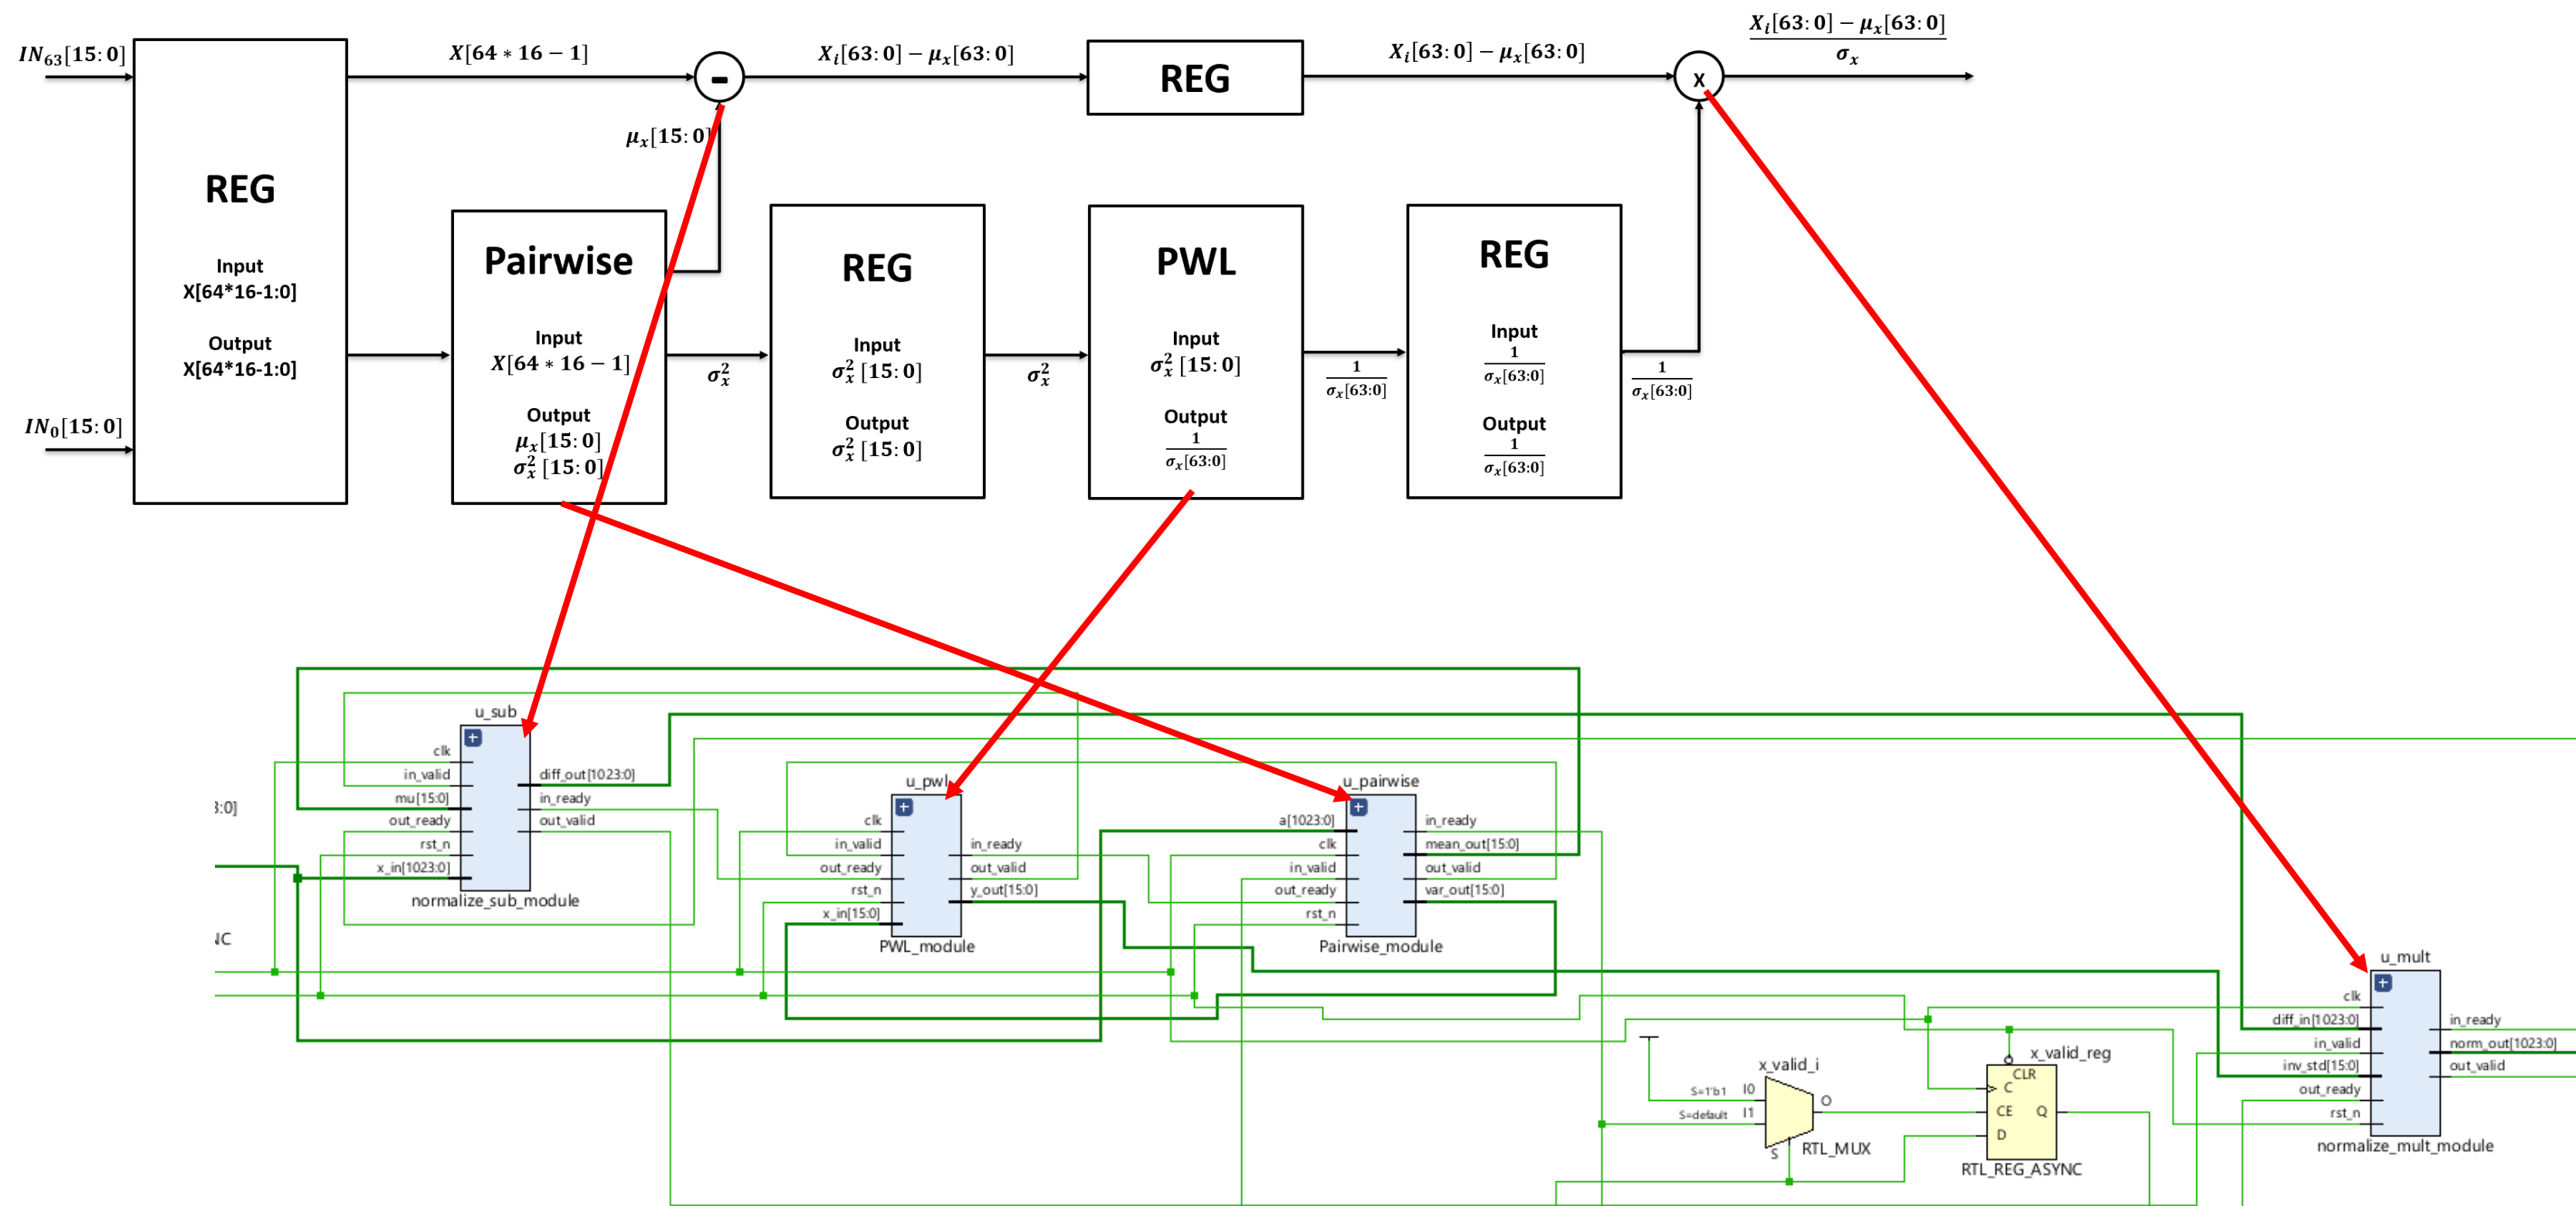


Q8.8 input hex (Dec)	X_in=0080 (0.5)	X_in=0200 (2.0)	X_in=0a00 (10.0)	X_in=1400 (20.0)	X_in=1e00 (30.0)	Accuracy
Sqrt	0x00a9 (0.660)	0x0164 (1.391)	0x0328 (3.156)	0480 (4.5)	0580 (5.5)	
Sqrt expected	0x00b5 (0.707)	0x016a (1.414)	0x032a (3.164)	0479 (4.473)	057a (5.477)	
Sqrt accuracy	93.37%	98.34%	0x99.75%	99.39%	99.57%	98.08% (error=1.92%)
reciprocal	0x017e (1.492)	0x00bc (0.734)	0x0053 (0.324)	0031 (0.191)	002b (0.168)	
reciprocal expected	0x016a (1.414)	0x00b5 (0.707)	0x0051 (0.316)	0039 (0.223)	002f (0.184)	
Reciprocal accuracy	94.48%	96.13%	97.53%	85.96%	91.49%	93.10% (error=6.90%)

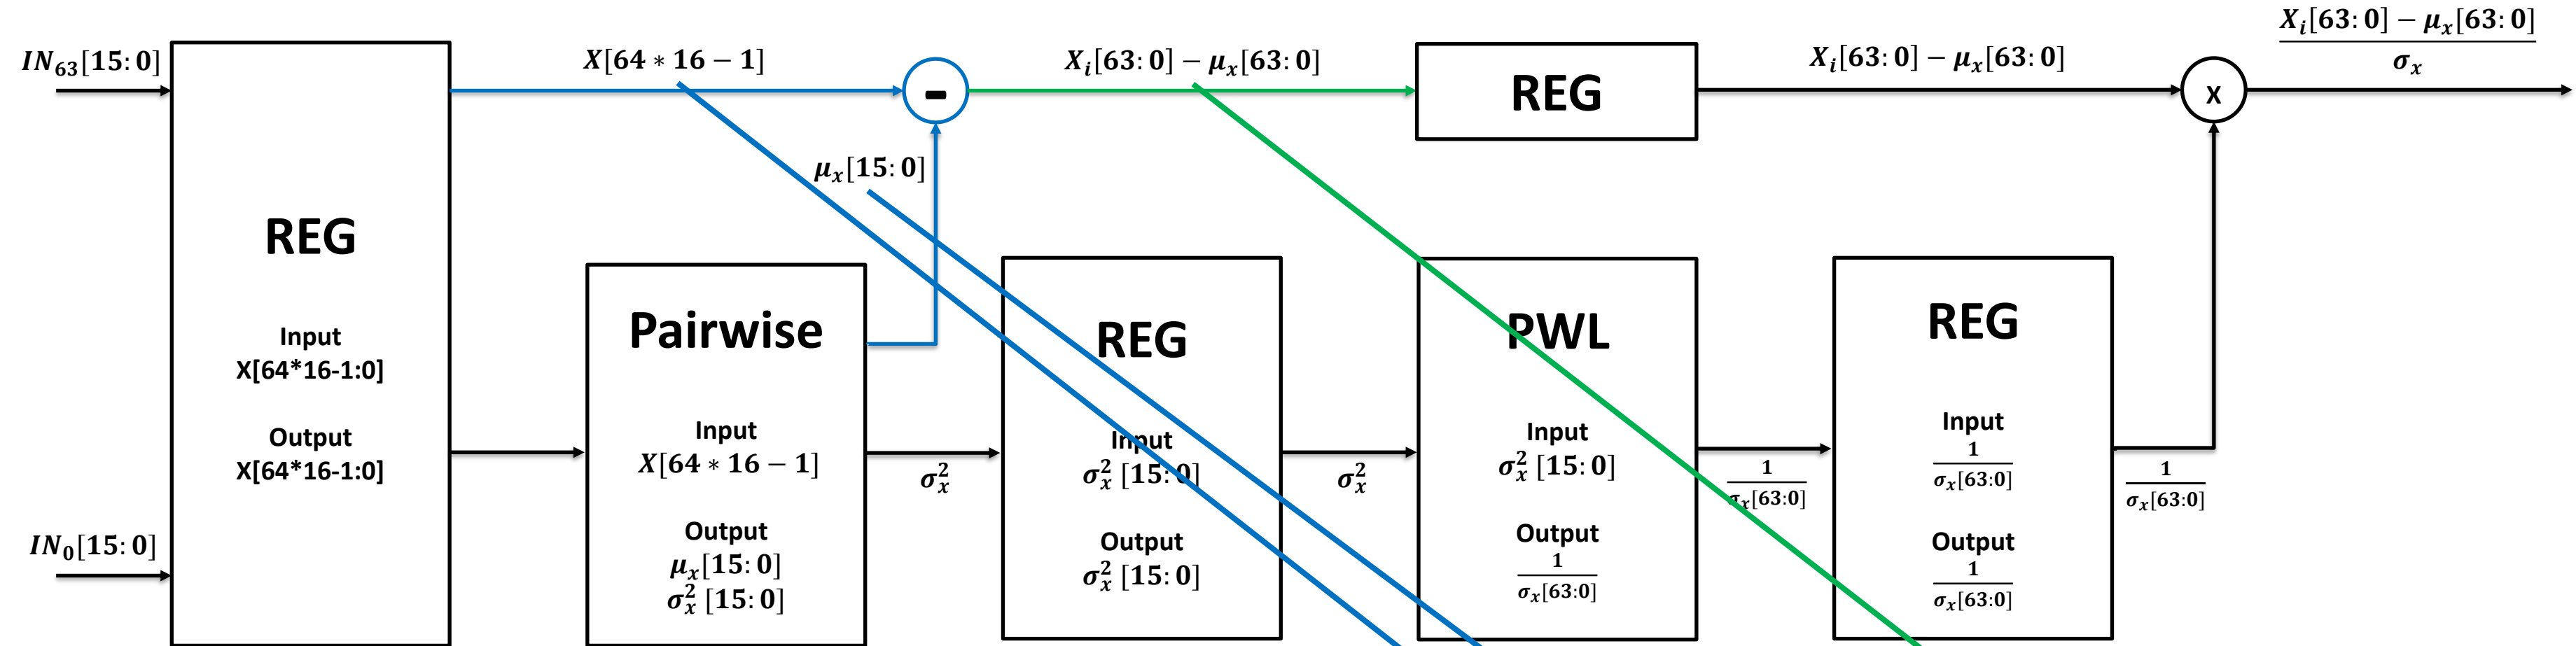
LN Approximation Model



LN Approximation Model



LN Approximation Model



```
genvar i;
generate
  for (i = 0; i < N; i = i + 1) begin : SUB_GEN
    assign x[i] = x_in[i*16 +: 16];

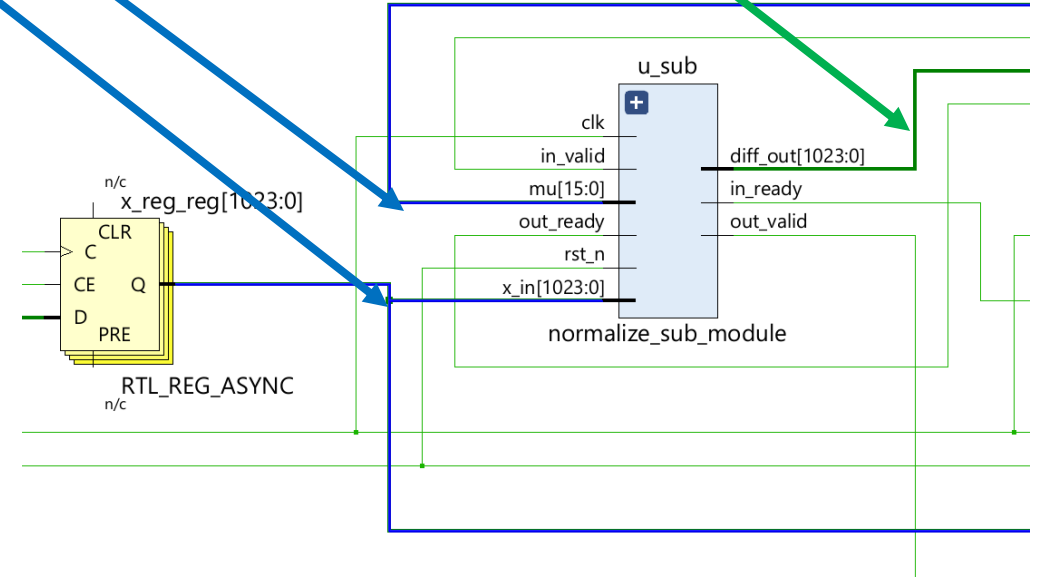
    c_addsub_sub sub_inst (
      .A(x[i]),
      .B(mu),
      .CLK(clk),
      .CE (1'b1),
      .S(diff[i])
    );

    assign diff_out[i*16 +: 16] = diff[i];
  end
endgenerate
```

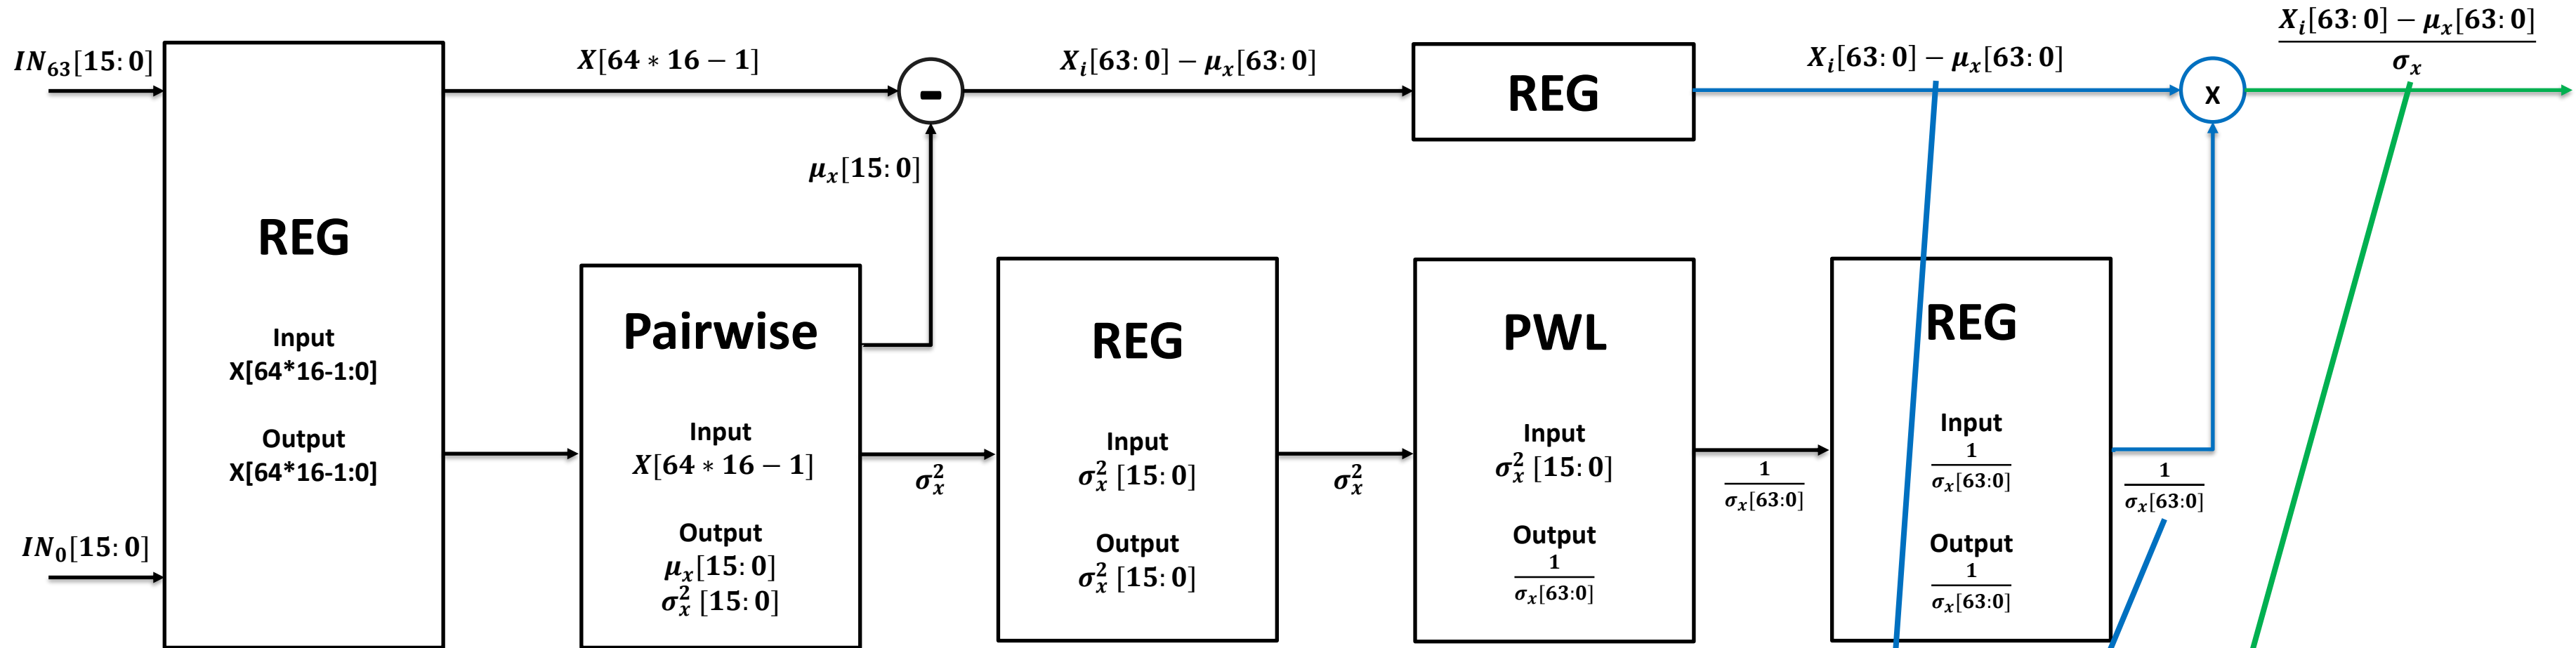
data

Input A : $X_i[63:0]$
Input B : $\mu_x[63:0]$
Output : $X_i[63:0] - \mu_x[63:0]$

Using addsub IP \rightarrow 1 Latency



LN Approximation Model



```
genvar i;
generate
  for (i = 0; i < N; i = i + 1) begin : MULT_GEN
    assign diff[i] = diff_in[i*16 +: 16];

    mult_gen_mult mult_inst (
      .A(diff[i]),
      .B(inv_std),
      .CLK(clk),
      .P(prod[i])
    );

    assign result[i] = prod[i][23:8]; // Q16.16 → Q8.8
    assign norm_out[i*16 +: 16] = result[i];
  end
endgenerate
```

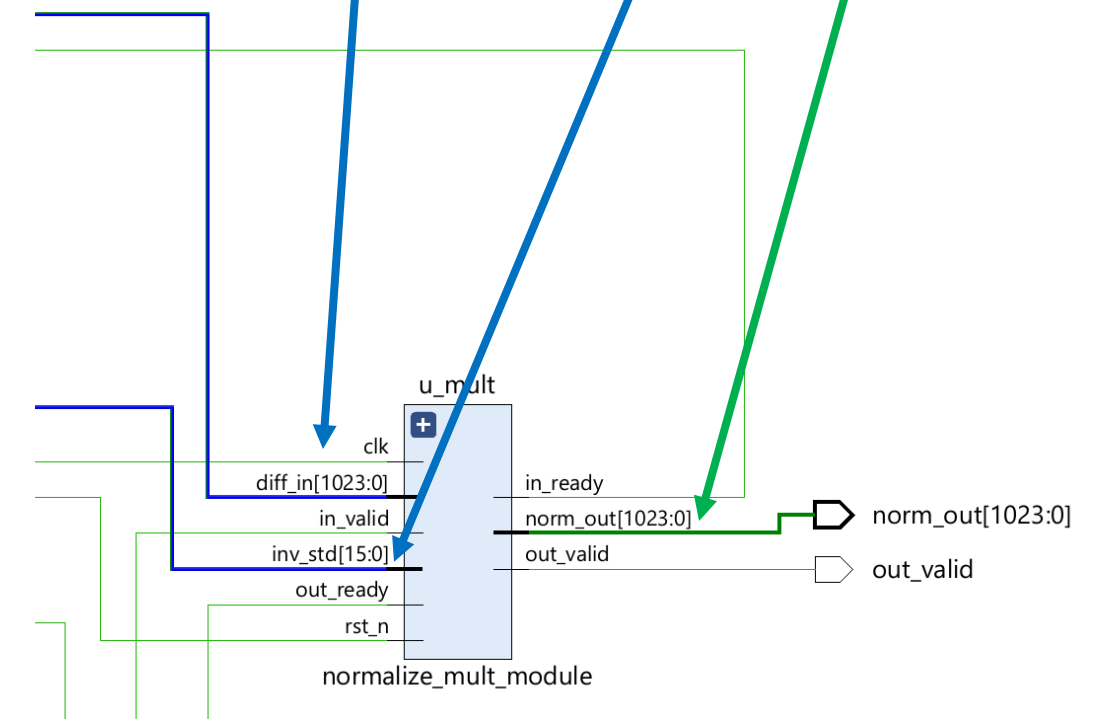
data

Input A : $X_i[63:0] - \mu_x[63:0]$

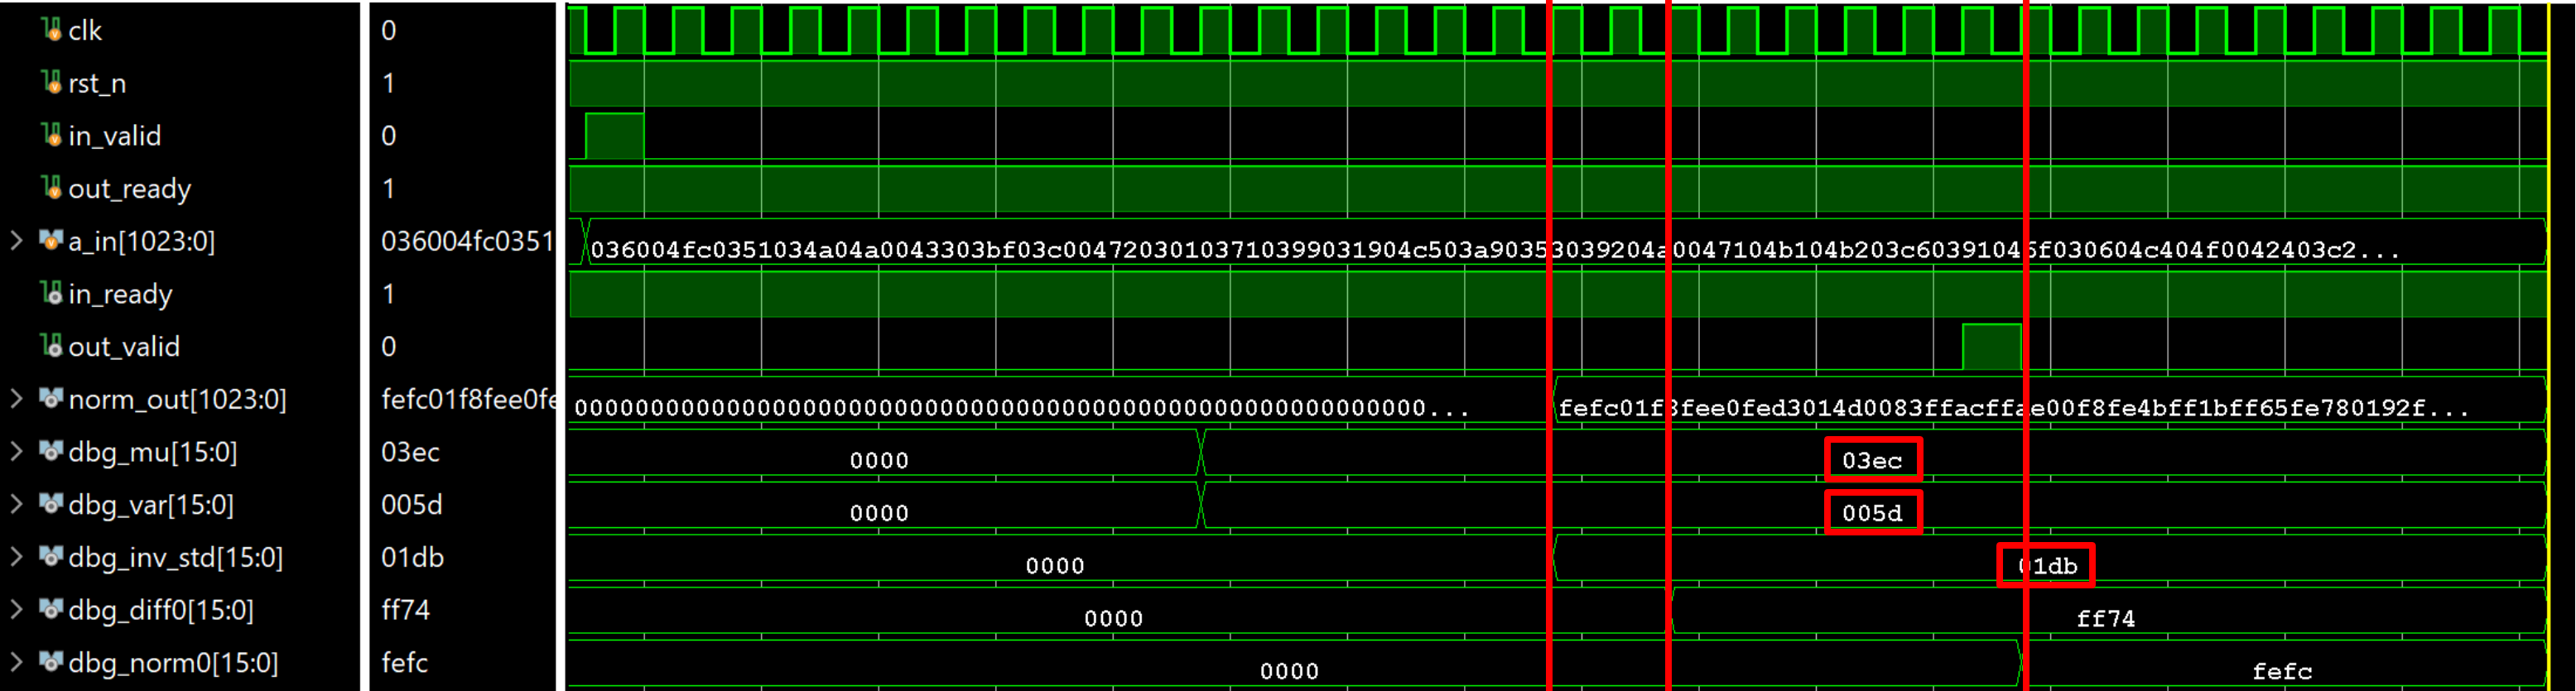
Input B : $\frac{1}{\sigma_x[63:0]}$

Output : $\frac{X_i[63:0] - \mu_x[63:0]}{\sigma_x}$

Using mult IP → 4 Latency



LN Approximation Model



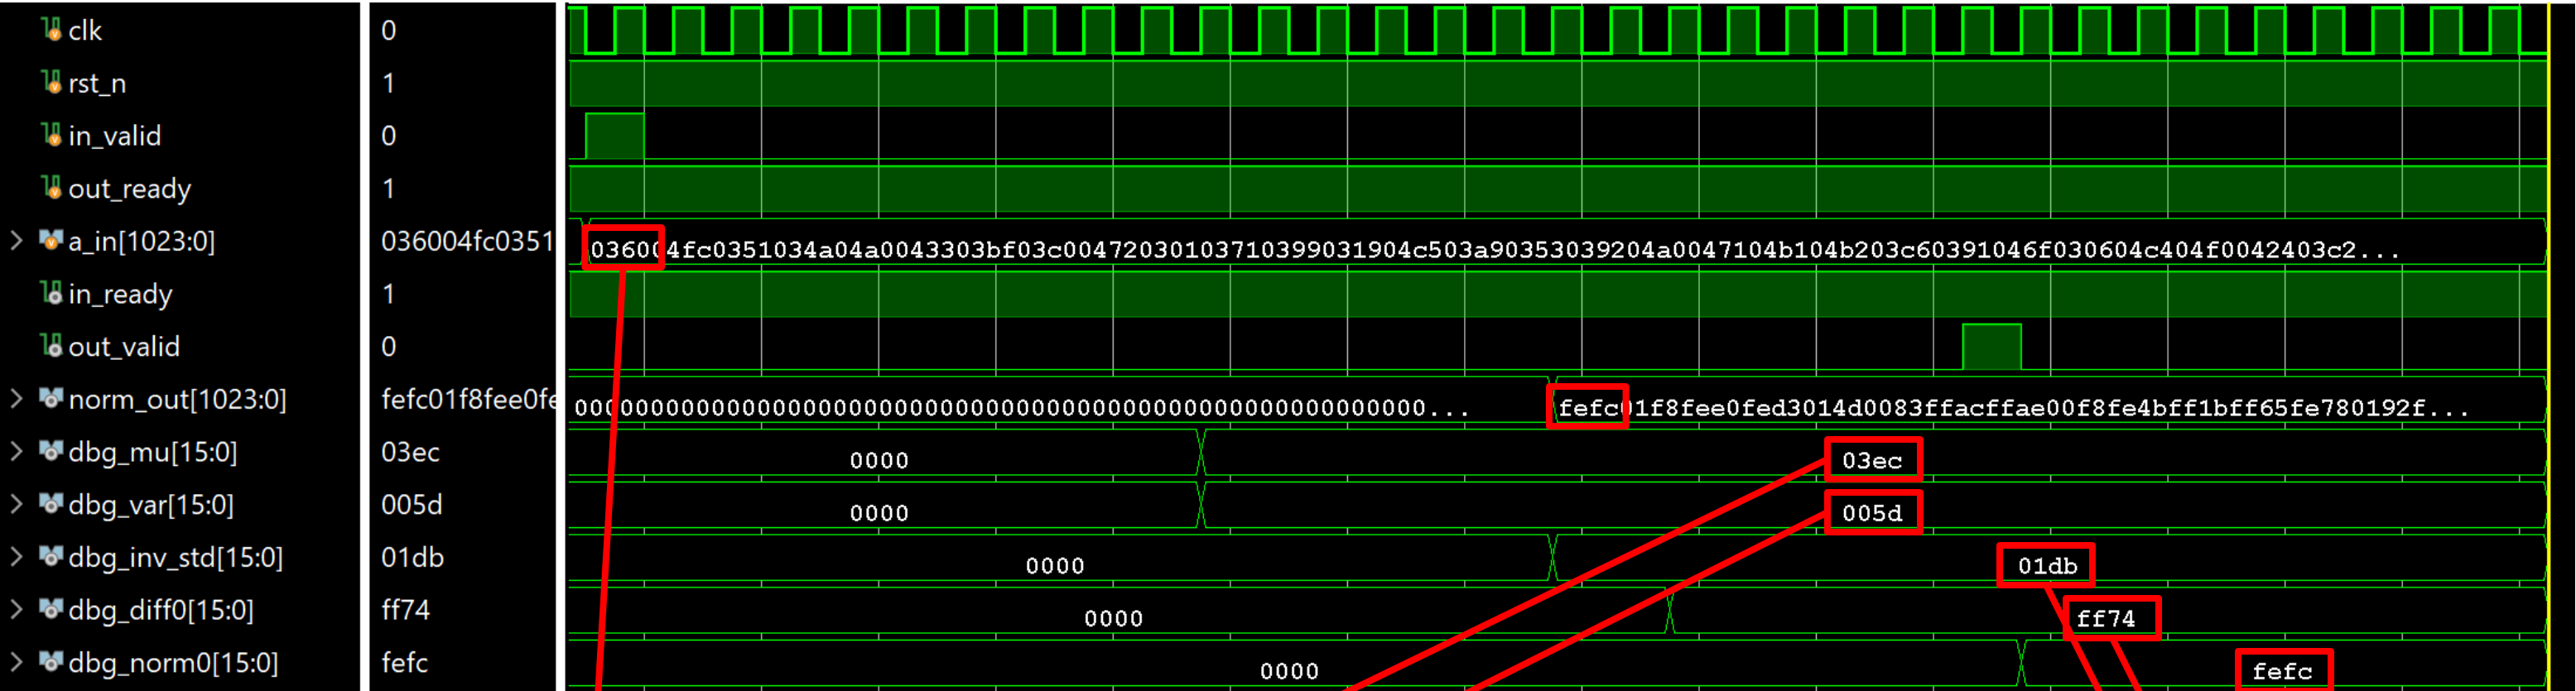
```
// 1 clock delay cycle to store reg
reg [0:0] delay_cnt;
always @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        delay_cnt <= 0;
        out_valid <= 0;
    end else if (in_valid && in_ready) begin
        delay_cnt <= 1;
        out_valid <= 0;
    end else if (delay_cnt != 0) begin
        delay_cnt <= delay_cnt - 1;
        out_valid <= 1; // out_valid is 1 after 1 clk delayed
    end else if (out_valid && out_ready) begin
        out_valid <= 0;
    end else begin
        out_valid <= 0;
    end
end
end
```

Using addsub IP → 2 Latency = 1 Latency(add) + 1 Latency(reg)
Using mult IP → 6 Latency = 4 Latency(mult) + 1 Latency(reg) + 1 Latency(dbg)

During debugging, extra 1 cycles are added for safe result verification, resulting in more delay than the actual waveform.

Thus will reduce the Latency.

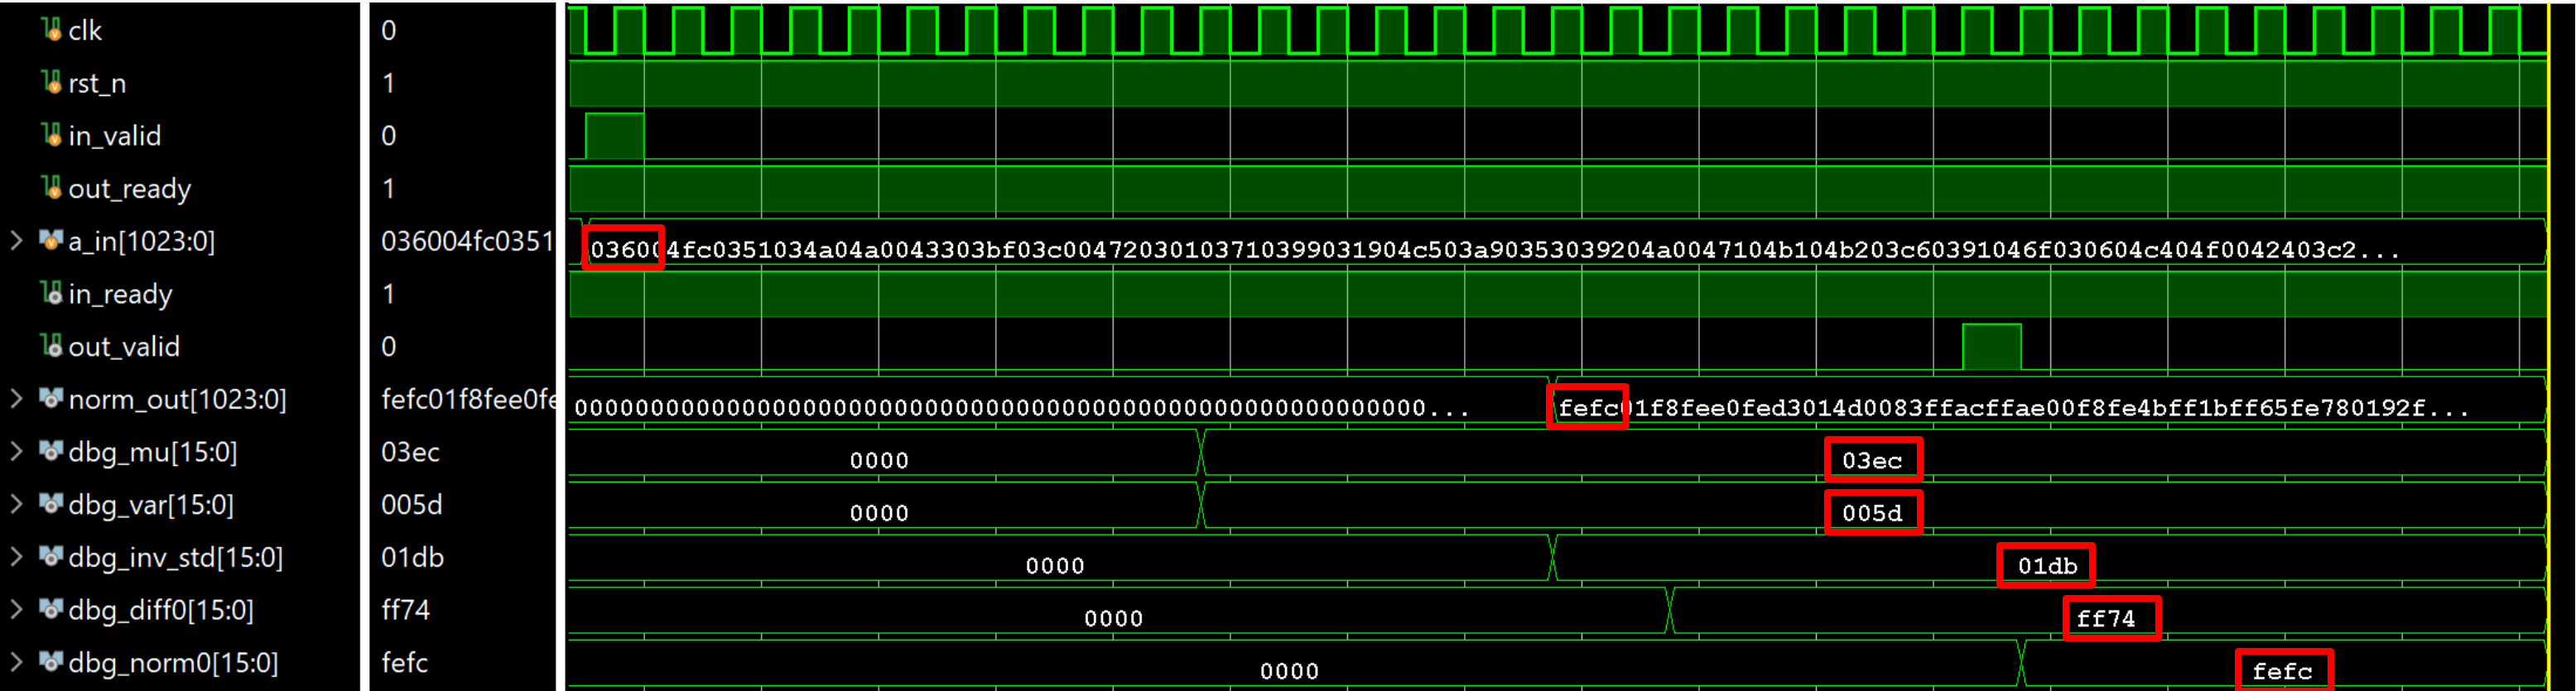
LN Approximation Model



64 random input values ranging from 3.00 to 5.00

Input data	0x0360(3.375)				
Result_Mean	0x03ec(3.922)				
Result_Var	0x005d(0.363)	→	0x005d(Var) * 0x00E2(a) + 0x0038(b)	→	\sqrt{Var} : 0x008A(0.539)
\sqrt{Var}	0x008A(0.539)	→	0x0x008A(\sqrt{Var}) * 0xFB5F(a) + 0x045A(b)	→	$1/\sqrt{Var}$: 0x01DB(1.855)
Result_Diff	0x0360(3.375) / 0x03ec(3.922)	→	0x0360(3.375) - 0x03ec(3.922)	→	Diff : 0xFF74(-0.547)
Result_Norm	0xFF74(-0.547) / 0x01DB(1.855)	→	0xFF74(-0.547) * 0x01DB(1.855)	→	Norm : 0xFEFC(-1.016)

LN Approximation Model



64 random input values ranging from 3.00 to 5.00

	Real	Approx	Accuracy(%)
Input Value	3.37	0x0360(3.375)	99.85
Mean	3.937	0x03ec(3.922)	99.61
Variance	0.368	0x005d(0.363)	98.64
1/Variance	1.649	0x01db(1.855)	87.51
Result(Norm)	-0.935	0xfefc(-1.016)	91.34
Norm mean	0	0.004	-

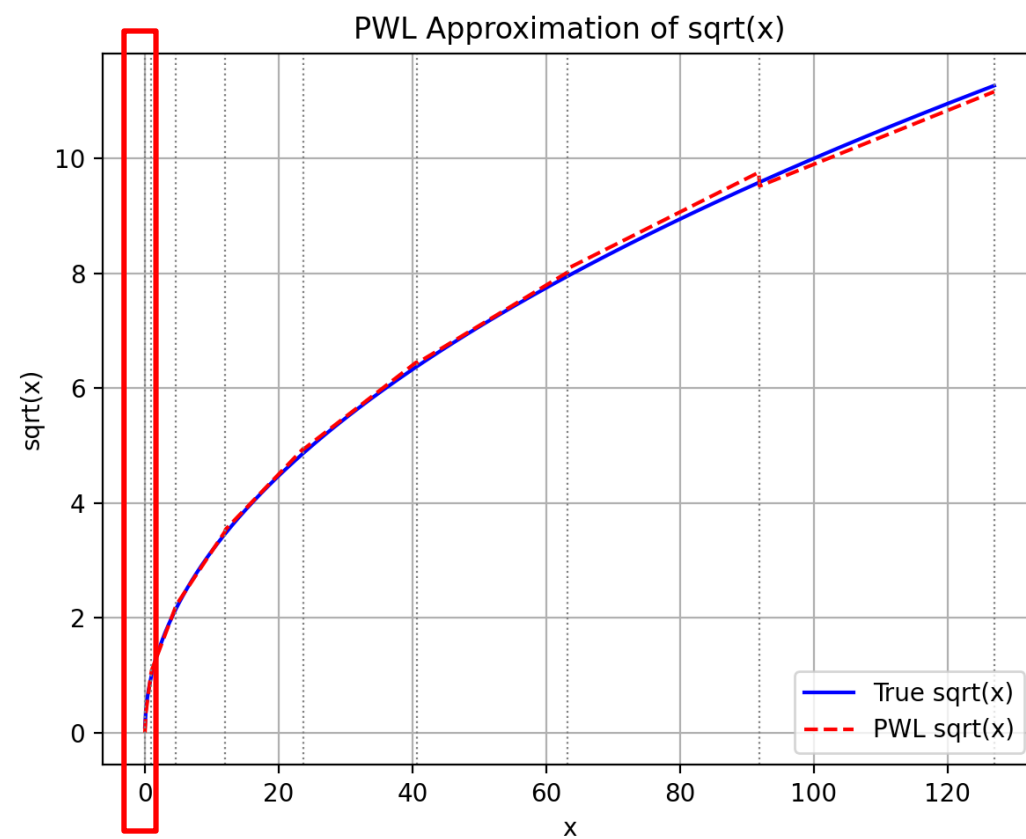
Accuracy(%)	data count
90 ~ 100	31
80 ~ 90	26
Other	7

LN Approximation Model

64 random input values ranging from 3.00 to 5.00

Accuracy dropped due to the narrow input range ($\text{var} \approx 0$), which is much **smaller than** the **expected real-case variance** ($20 \sim 30$).

Therefore, **LUT range will be redefined** based on the **actual input distribution**.




Apply FPGA

Input : 64 * 16bit → [1023:0] : 1024EA

Output 64 * 16bit → [1023:0] : 1024EA

Other port(clk, rst...) : 6EA




Need 2054 I/O port, But only exist 210 I/O port

→ Using FSM(Finite State Machine)

Apply FPGA

When Calculate add, sub, mult.

When not using IP



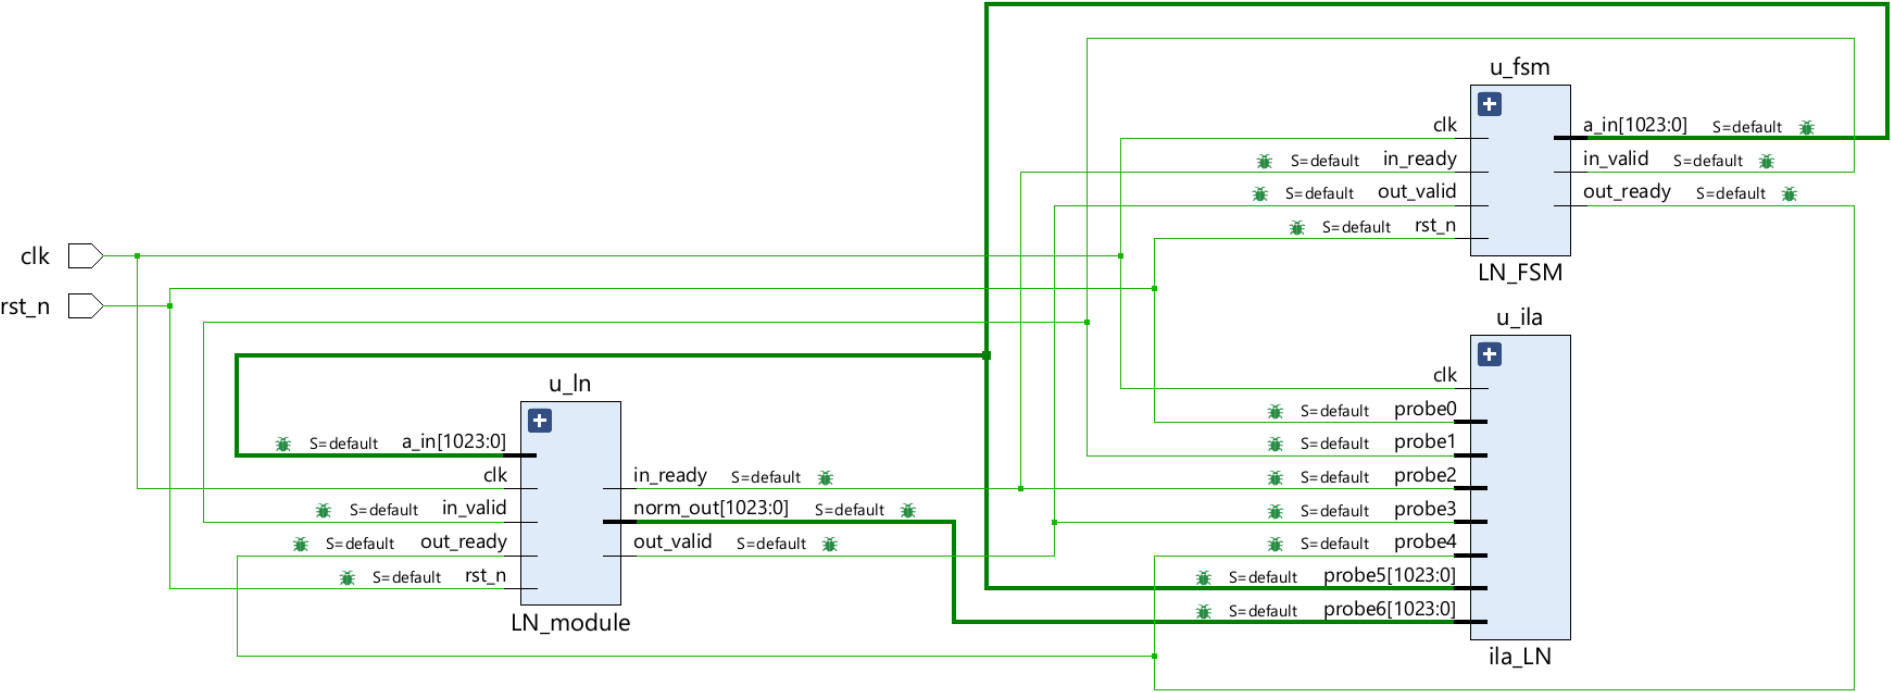
Too many used LUT, DSP → Consider IP

Using IP

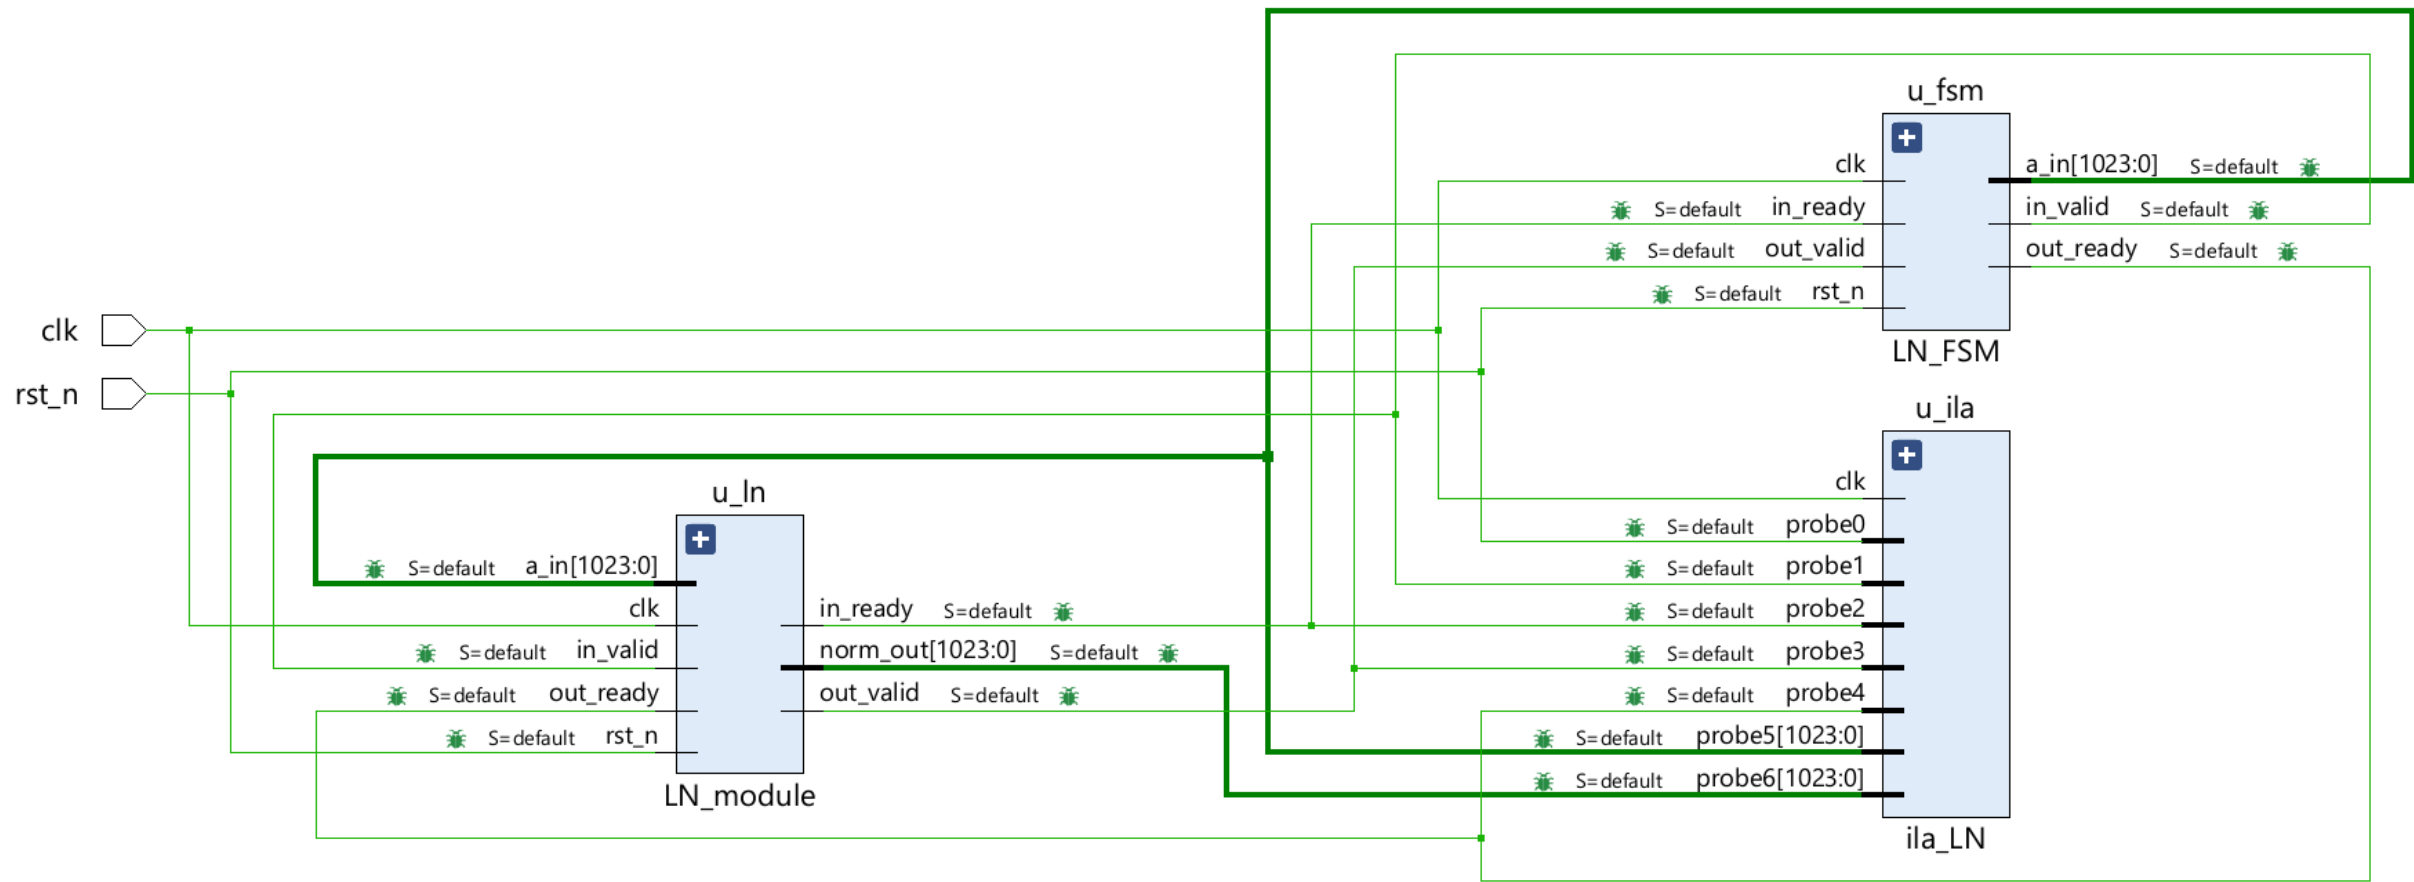


Logic Overview

- Top_module** : include LN_module, FSM, ILA.
- LN_module** : Performs normalization on input data
- FSM** : Used to verify correct operation on FPGA before **applying BRAM**
- ILA** : Used to observe signal values and **confirm correct operation on FPGA**



Logic Overview

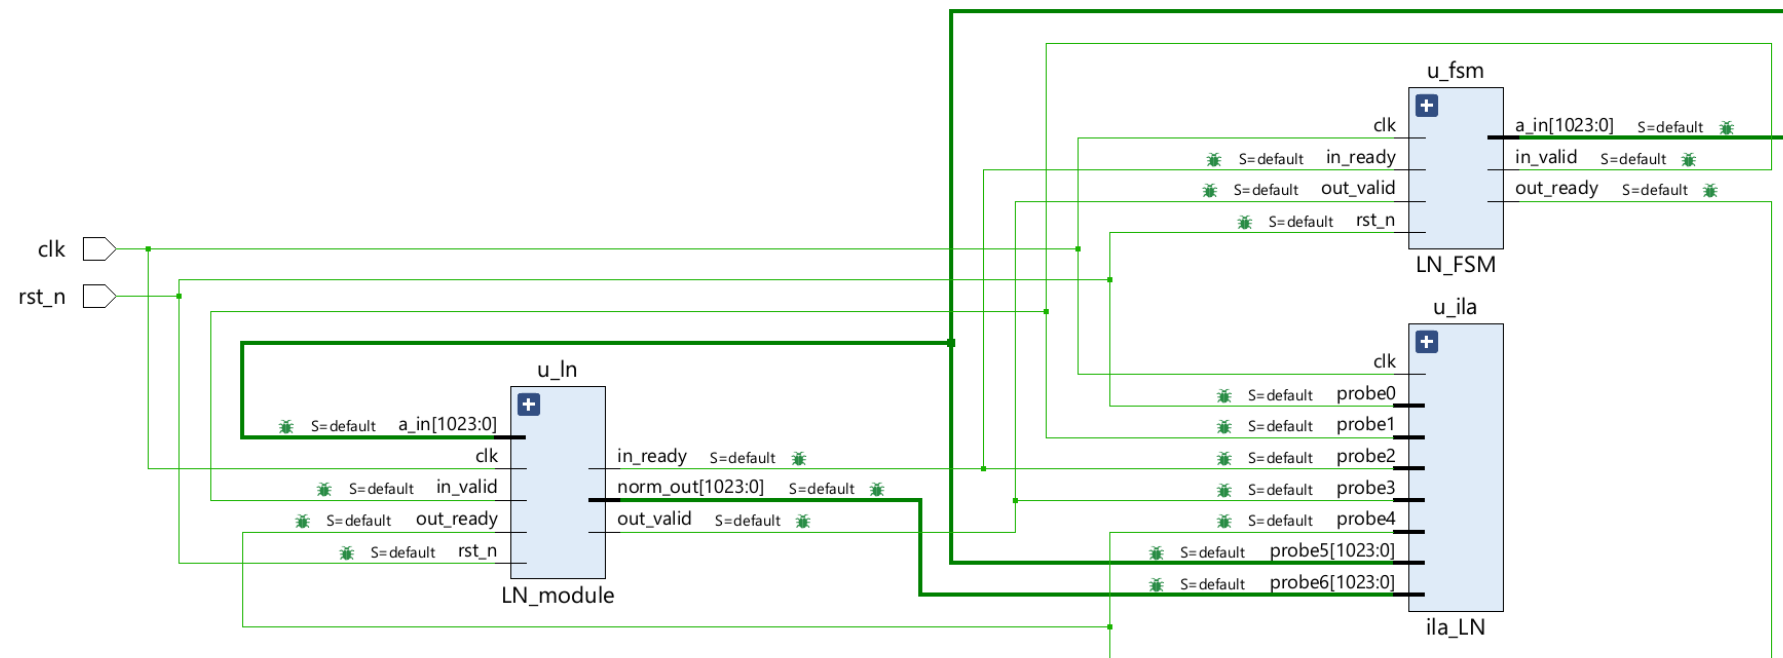


```
LN_module u_ln (  
    .clk(clk),           // 1bit, input  
    .rst_n(rst_n),       // 1bit, input  
    .in_valid(in_valid), // 1bit, input  
    .in_ready(in_ready), // 1bit, output  
    .out_valid(out_valid), // 1bit, output  
    .out_ready(out_ready), // 1bit, input  
    .a_in(a_in),         // 64*16bit => 1024, input  
    .norm_out(norm_out)  // 64*16bit => 1024, output  
);
```

```
LN_FSM u_fsm (  
    .clk(clk),           // 1bit, input  
    .rst_n(rst_n),       // 1bit, input  
    .in_valid(in_valid), // 1bit, input  
    .in_ready(in_ready), // 1bit, input  
    .a_in(a_in),         // 64*16bit => 1024, output  
    .out_valid(out_valid), // 1bit, output  
    .out_ready(out_ready) // 1bit, output
```

```
ila_LN u_ila (  
    .clk(clk),           // 1bit, input  
    .probe0(rst_n),      // 1bit, input  
    .probe1(in_valid),   // 1bit, input  
    .probe2(in_ready),   // 1bit, input  
    .probe3(out_valid),  // 1bit, input  
    .probe4(out_ready),  // 1bit, input  
    .probe5(a_in),       // 64*16bit => 1024, input  
    .probe6(norm_out)    // 64*16bit => 1024, input  
);
```

Apply FPGA



However, **Error occurred** due to excessive **resource usage** during the **pairwise** computation of the **64 input values**.

Therefore, the data input will be changed to **32 elements**, as proposed in the paper.

Apply FPGA

Error about 64 Inputs.

Error due to **excessive resource usage**.

Occurred during pairwise computation **(64 inputs)**.

Plan to Solve : change 32 Inputs

Change input size : 64 → 32 elements

Mult IP : 1-cycle → 4-cycle (allows higher freq.)

Latency Change

Pairwise stage : 6 cycles → 15 cycles

GPT-2 model(768) : +12 cycles

Removed 64→32 conversion stage : -15 cycles

Conclusion

Higher frequency possible

Net latency reduced by 3 cycles

Reduces resource consumption compared to 64-input design

Plans for Next

- **Change the number of input data from 64 to 32.**
- **Redefine the LUT based on the updated input data.**
- **Reduce unnecessary latency (originally added for debugging purposes).**
- **Enable continuous data input through pipelining. +) also using BRAM**
- **Upload the design to the FPGA and compare it with the baseline model.**